



# RISC-V Supervisor Binary Interface Specification

RISC-V Platform Runtime Services Task Group

Version -, 2023-06-02: Draft

# Table of Contents

|   |    |
|---|----|
| Preamble  | 1  |
| Copyright and license information                       | 2  |
| Change Log  | 3  |
| Version 2.0-rc1   | 3  |
| Version 1.0.0   | 3  |
| Version 1.0-rc3   | 3  |
| Version 1.0-rc2   | 3  |
| Version 1.0-rc1   | 4  |
| Version 0.3.0   | 4  |
| Version 0.3-rc1   | 4  |
| Version 0.2   | 4  |
| 1. Introduction   | 5  |
| 2. Terms and Abbreviations                              | 7  |
| 3. Binary Encoding                                      | 8  |
| 3.1. HART list parameter                                | 9  |
| 3.2. Shared memory physical address range parameter     | 9  |
| 4. Base Extension (EID #0x10)                           | 11 |
| 4.1. Function: Get SBI specification version (FID #0)   | 11 |
| 4.2. Function: Get SBI implementation ID (FID #1)       | 11 |
| 4.3. Function: Get SBI implementation version (FID #2)  | 11 |
| 4.4. Function: Probe SBI extension (FID #3)             | 11 |
| 4.5. Function: Get machine vendor ID (FID #4)           | 11 |
| 4.6. Function: Get machine architecture ID (FID #5)     | 12 |
| 4.7. Function: Get machine implementation ID (FID #6)   | 12 |
| 4.8. Function Listing                                   | 12 |
| 4.9. SBI Implementation IDs                             | 12 |
| 5. Legacy Extensions (EIDs #0x00 - #0x0F)               | 14 |
| 5.1. Extension: Set Timer (EID #0x00)                   | 14 |
| 5.2. Extension: Console Putchar (EID #0x01)             | 14 |
| 5.3. Extension: Console Getchar (EID #0x02)             | 15 |
| 5.4. Extension: Clear IPI (EID #0x03)                   | 15 |
| 5.5. Extension: Send IPI (EID #0x04)                    | 15 |
| 5.6. Extension: Remote FENCE.I (EID #0x05)              | 15 |
| 5.7. Extension: Remote SFENCE.VMA (EID #0x06)           | 16 |
| 5.8. Extension: Remote SFENCE.VMA with ASID (EID #0x07) | 16 |
| 5.9. Extension: System Shutdown (EID #0x08)             | 16 |
| 5.10. Function Listing                                  | 16 |
| 6. Timer Extension (EID #0x54494D45 "TIME")             | 18 |
| 6.1. Function: Set Timer (FID #0)                       | 18 |
| 6.2. Function Listing                                   | 18 |

|   |    |
|---|----|
| 7. IPI Extension (EID #0x735049 "sPI: s-mode IPI")              | 19 |
| 7.1. Function: Send IPI (FID #0)                                | 19 |
| 7.2. Function Listing   | 19 |
| 8. RFENCE Extension (EID #0x52464E43 "RFNC")                    | 20 |
| 8.1. Function: Remote FENCE.I (FID #0)                          | 20 |
| 8.2. Function: Remote SFENCE.VMA (FID #1)                       | 20 |
| 8.3. Function: Remote SFENCE.VMA with ASID (FID #2)             | 21 |
| 8.4. Function: Remote HFENCE.GVMA with VMID (FID #3)            | 21 |
| 8.5. Function: Remote HFENCE.GVMA (FID #4)                      | 22 |
| 8.6. Function: Remote HFENCE.VVMA with ASID (FID #5)            | 22 |
| 8.7. Function: Remote HFENCE.VVMA (FID #6)                      | 23 |
| 8.8. Function Listing   | 23 |
| 9. Hart State Management Extension (EID #0x48534D "HSM")        | 24 |
| 9.1. Function: HART start (FID #0)                              | 25 |
| 9.2. Function: HART stop (FID #1)                               | 27 |
| 9.3. Function: HART get status (FID #2)                         | 27 |
| 9.4. Function: HART suspend (FID #3)                            | 27 |
| 9.5. Function Listing   | 29 |
| 10. System Reset Extension (EID #0x53525354 "SRST")             | 30 |
| 10.1. Function: System reset (FID #0)                           | 30 |
| 10.2. Function Listing  | 31 |
| 11. Performance Monitoring Unit Extension (EID #0x504D55 "PMU") | 32 |
| 11.1. Event: Hardware general events (Type #0)                  | 32 |
| 11.2. Event: Hardware cache events (Type #1)                    | 33 |
| 11.3. Event: Hardware raw events (Type #2)                      | 34 |
| 11.4. Event: Firmware events (Type #15)                         | 35 |
| 11.5. Function: Get number of counters (FID #0)                 | 36 |
| 11.6. Function: Get details of a counter (FID #1)               | 36 |
| 11.7. Function: Find and configure a matching counter (FID #2)  | 37 |
| 11.8. Function: Start a set of counters (FID #3)                | 38 |
| 11.9. Function: Stop a set of counters (FID #4)                 | 39 |
| 11.10. Function: Read a firmware counter (FID #5)               | 40 |
| 11.11. Function: Read a firmware counter high bits (FID #6)     | 41 |
| 11.12. Function: Enable PMU snapshot feature (FID #7)           | 41 |
| 11.13. Function Listing   | 42 |
| 12. Debug Console Extension (EID #0x4442434E "DBCN")            | 43 |
| 12.1. Function: Console Write (FID #0)                          | 43 |
| 12.2. Function: Console Read (FID #1)                           | 44 |
| 12.3. Function: Console Write Byte (FID #2)                     | 44 |
| 12.4. Function Listing  | 45 |
| 13. System Suspend Extension (EID #0x53555350 "SUSP")           | 46 |
| 13.1. Function: System Suspend (FID #0)                         | 46 |
| 13.2. Function Listing  | 47 |

|  |    |
|--|----|
| 14. CPPC Extension (EID #0x43505043 "CPPC")                            | 49 |
| 14.1. Function: Probe CPPC register (FID #0)                           | 50 |
| 14.2. Function: Read CPPC register (FID #1)                            | 51 |
| 14.3. Function: Read CPPC register high bits (FID #2)                  | 51 |
| 14.4. Function: Write to CPPC register (FID #3)                        | 52 |
| 14.5. Function Listing   | 52 |
| 15. Nested Acceleration Extension (EID #0x4E41434C "NACL")             | 53 |
| 15.1. Feature: Synchronize CSR (ID #0)                                 | 54 |
| 15.2. Feature: Synchronize HFENCE (ID #1)                              | 55 |
| 15.3. Feature: Synchronize SRET (ID #2)                                | 57 |
| 15.4. Feature: Autoswap CSR (ID #3)                                    | 59 |
| 15.5. Function: Probe nested acceleration feature (FID #0)             | 59 |
| 15.6. Function: Set nested acceleration shared memory (FID #1)         | 60 |
| 15.7. Function: Synchronize shared memory CSRs (FID #2)                | 60 |
| 15.8. Function: Synchronize shared memory HFENCES (FID #3)             | 61 |
| 15.9. Function: Synchronize shared memory and emulate SRET (FID #4)    | 62 |
| 15.10. Function Listing  | 62 |
| 16. Steal-time Accounting Extension (EID #0x535441 "STA")              | 63 |
| 16.1. Function: Set Steal-time Shared Memory Address (FID #0)          | 63 |
| 16.2. Function Listing   | 65 |
| 17. SBI Supervisor Software Events Extension (EID #0x535345 "SSE")     | 66 |
| 17.1. Software Event Identification                                    | 66 |
| 17.2. Software Event States  | 66 |
| 17.3. Software Event Priority  | 67 |
| 17.4. Software Event Attributes  | 68 |
| 17.5. Software Event Handler   | 68 |
| 17.6. Software Event Injection   | 71 |
| 17.7. Software Event Completion  | 71 |
| 17.8. Function: Get a software event attribute (FID #0)                | 72 |
| 17.9. Function: Set a software event attribute (FID #1)                | 72 |
| 17.10. Function: Register a software event (FID #2)                    | 73 |
| 17.11. Function: Unregister a software event (FID #3)                  | 74 |
| 17.12. Function: Enable a software event (FID #4)                      | 74 |
| 17.13. Function: Disable a software event (FID #5)                     | 75 |
| 17.14. Function: Complete software event handling (FID #6)             | 75 |
| 17.15. Function: Signal a software event (FID #7)                      | 76 |
| 18. Experimental SBI Extension Space (EIDs #0x08000000 - #0x08FFFFFF)  | 78 |
| 19. Vendor Specific Extension Space (EIDs #0x09000000 - #0x09FFFFFF)   | 79 |
| 20. Firmware Specific Extension Space (EIDs #0x0A000000 - #0x0AFFFFFF) | 80 |
| References   | 81 |

# Preamble



*This document is in the [Development state](#)*

*Assume everything can change but backward compatibility with perviously ratified specification will be maintained.*

# Copyright and license information

This RISC-V SBI specification is © 2022 RISC-V International.

It is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at [creativecommons.org/licenses/by/4.0/](https://creativecommons.org/licenses/by/4.0/).

# Change Log

## Version 2.0-rc1

- Added common description for shared memory physical address range parameter
- Added SBI debug console extension
- Relaxed the counter width requirement on SBI PMU firmware counters
- Added `sbi_pmu_counter_fw_read_hi()` in SBI PMU extension
- Reserved space for SBI implementation specific firmware events
- Added SBI system suspend extension
- Added SBI CPPC extension
- Clarified that an SBI extension can be partially implemented only if it defines a mechanism to discover implemented SBI functions
- Added error code `SBI_ERR_NO_SHMEM`
- Added SBI nested acceleration extension
- Added common description for a virtual HART
- Added SBI steal-time accounting extension
- Added SBI PMU snapshot extension
- Added error codes `SBI_ERR_INVALID_STATE`, `SBI_ERR_BAD_RANGE`, `SBI_ERR_BUSY`
- Added SBI supervisor software event extension

## Version 1.0.0

- Updated the version for ratification

## Version 1.0-rc3

- Updated the calling convention
- Fixed a typo in PMU extension
- Added a abbreviation table

## Version 1.0-rc2

- Update to RISC-V formatting
- Improved the introduction
- Removed all references to RV32

## Version 1.0-rc1

- A typo fix

## Version 0.3.0

- Few typo fixes
- Updated the LICENSE with detailed text instead of a hyperlink

## Version 0.3-rc1

- Improved document styling and naming conventions
- Added SBI system reset extension
- Improved SBI introduction section
- Improved documentation of SBI hart state management extension
- Added suspend function to SBI hart state management extension
- Added performance monitoring unit extension
- Clarified that an SBI extension shall not be partially implemented

## Version 0.2

- The entire v0.1 SBI has been moved to the legacy extension, which is now an optional extension. This is technically a backwards-incompatible change because the legacy extension is optional and v0.1 of the SBI doesn't allow probing, but it's as good as we can do.



# Chapter 1. Introduction

This specification describes the RISC-V Supervisor Binary Interface, known from here on as SBI. The SBI allows supervisor-mode (S-mode or VS-mode) software to be portable across all RISC-V implementations by defining an abstraction for platform (or hypervisor) specific functionality. The design of the SBI follows the general RISC-V philosophy of having a small core along with a set of optional modular extensions.

An SBI extension defines a set of SBI functions which provides a particular functionality to supervisor-mode software. SBI extensions as a whole are optional and cannot be partially implemented unless an SBI extension defines a mechanism to discover implemented SBI functions. If `sbi_probe_extension()` signals that an extension is available, all functions present in the SBI version reported by `sbi_get_spec_version()` must conform to that version of the SBI specification.

The higher privilege software providing SBI interface to the supervisor-mode software is referred as an SBI implementation or Supervisor Execution Environment (SEE). An SBI implementation (or SEE) can be platform runtime firmware executing in machine-mode (M-mode) (see below [Figure 1](#)) or it can be some hypervisor executing in hypervisor-mode (HS-mode) (see below [Figure 2](#)).

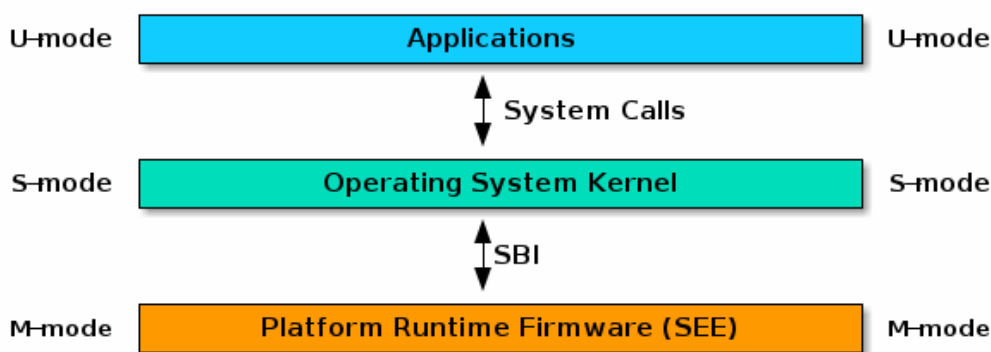


Figure 1. RISC-V System without H-extension




Figure 2. RISC-V System with H-extension

HARTs are provisioned by the SBI implementation for supervisor-mode software. Hence, from the perspective of the SBI implementation, the S-mode HART contexts are referred to as virtual HARTs. In the case that the implementation is a hypervisor, virtual HARTs represent the VS-mode guest contexts.

The SBI specification doesn't specify any method for hardware discovery. The supervisor software must rely on the other industry standard hardware discovery methods (i.e. Device Tree or ACPI) for that.

## Chapter 2. Terms and Abbreviations

This specification uses the following terms and abbreviations:

| Term | Meaning                                     |
|------|---|
| ACPI | Advanced Configuration and Power Interface  |
| ASID | Address Space Identifier                    |
| BMC  | Baseboard Management Controller             |
| CPPC | Collaborative Processor Performance Control |
| EID  | Extension ID                                |
| FID  | Function ID                                 |
| HSM  | Hart State Management                       |
| IPI  | Inter Processor Interrupt                   |
| PMU  | Performance Monitoring Unit                 |
| SBI  | Supervisor Binary Interface                 |
| SEE  | Supervisor Execution Environment            |
| VMID | Virtual Machine Identifier                  |

# Chapter 3. Binary Encoding

All SBI functions share a single binary encoding, which facilitates the mixing of SBI extensions. The SBI specification follows the below calling convention.

- An **ECALL** is used as the control transfer instruction between the supervisor and the SEE.
- **a7** encodes the SBI extension ID (**EID**),
- **a6** encodes the SBI function ID (**FID**) for a given extension ID encoded in **a7** for any SBI extension defined in or after SBI v0.2.
- All registers except **a0** & **a1** must be preserved across an SBI call by the callee.
- SBI functions must return a pair of values in **a0** and **a1**, with **a0** returning an error code. This is analogous to returning the C structure

```
struct sbiret {
    long error;
    long value;
};
```

In the name of compatibility, SBI extension IDs (**EIDs**) and SBI function IDs (**FIDs**) are encoded as signed 32-bit integers. When passed in registers these follow the standard above calling convention rules.

The [Table 1](#) below provides a list of Standard SBI error codes.

Table 1. Standard SBI Errors

| Error Type                | Value | Description                        |
|---------------------------|-------|------------------------------------|
| SBI_SUCCESS               | 0     | Completed successfully             |
| SBI_ERR_FAILED            | -1    | Failed                             |
| SBI_ERR_NOT_SUPPORTED     | -2    | Not supported                      |
| SBI_ERR_INVALID_PARAM     | -3    | Invalid parameter(s)               |
| SBI_ERR_DENIED            | -4    | Denied or not allowed              |
| SBI_ERR_INVALID_ADDRESS   | -5    | Invalid address(s)                 |
| SBI_ERR_ALREADY_AVAILABLE | -6    | Already available                  |
| SBI_ERR_ALREADY_STARTED   | -7    | Already started                    |
| SBI_ERR_ALREADY_STOPPED   | -8    | Already stopped                    |
| SBI_ERR_NO_SHMEM          | -9    | Shared memory not available        |
| SBI_ERR_INVALID_STATE     | -10   | Invalid operation on state machine |
| SBI_ERR_BAD_RANGE         | -11   | Value out of valid range           |
| SBI_ERR_BUSY              | -12   | Resource busy                      |

An **ECALL** with an unsupported SBI extension ID (**EID**) or an unsupported SBI function ID (**FID**) must return the error code **SBI\_ERR\_NOT\_SUPPORTED**.

Every SBI function should prefer **unsigned long** as the data type. It keeps the specification simple and easily adaptable for all RISC-V ISA types. In case the data is defined as 32bit wide, higher privilege software must ensure that it only uses 32 bit data only.

## 3.1. HART list parameter

If an SBI function needs to pass a list of harts to the higher privilege mode, it must use a hart mask as defined below. This is applicable to any extensions defined in or after v0.2.

Any function, requiring a hart mask, need to pass following two arguments.

- **unsigned long hart\_mask** is a scalar bit-vector containing hartids
- **unsigned long hart\_mask\_base** is the starting hartid from which bit-vector must be computed.

In a single SBI function call, maximum number harts that can be set is always XLEN. If a lower privilege mode needs to pass information about more than XLEN harts, it should invoke multiple instances of the SBI function call. **hart\_mask\_base** can be set to **-1** to indicate that **hart\_mask** can be ignored and all available harts must be considered.

Any function using hart mask may return error values listed in the [Table 2](#) below which are in addition to function specific error values.

Table 2. HART Mask Errors

| Error code            | Description  |
|-----------------------|--|
| SBI_ERR_INVALID_PARAM | Either <b>hart_mask_base</b> , or at least one hartid from <b>hart_mask</b> , is not valid, i.e. either the hartid is not enabled by the platform or is not available to the supervisor. |

## 3.2. Shared memory physical address range parameter

If an SBI function needs to pass a shared memory physical address range to the SBI implementation (or higher privilege mode), then this physical memory address range MUST satisfy the following requirements:

- The SBI implementation MUST check that the supervisor-mode software is allowed to access the specified physical memory range with the access type requested (read and/or write).
- The SBI implementation MUST access the specified physical memory range using the PMA attributes.



*If the supervisor-mode software accesses the same physical memory range using a memory type different than the PMA, then a loss of coherence or unexpected memory ordering may occur. The invoking software should follow the rules and sequences defined in the RISC-V Svpbmt specification to prevent the loss of coherence and memory ordering.*

- The data in the shared memory MUST follow little-endian byte ordering.

It is recommended that a memory physical address passed to an SBI function should use at least two

**unsigned long** parameters to support platforms which have memory physical addresses wider than XLEN bits.

## Chapter 4. Base Extension (EID #0x10)

The base extension is designed to be as small as possible. As such, it only contains functionality for probing which SBI extensions are available and for querying the version of the SBI. All functions in the base extension must be supported by all SBI implementations, so there are no error returns defined.

### 4.1. Function: Get SBI specification version (FID #0)

```
struct sbiret sbi_get_spec_version(void);
```

Returns the current SBI specification version. This function must always succeed. The minor number of the SBI specification is encoded in the low 24 bits, with the major number encoded in the next 7 bits. Bit 31 must be 0 and is reserved for future expansion.

### 4.2. Function: Get SBI implementation ID (FID #1)

```
struct sbiret sbi_get_impl_id(void);
```

Returns the current SBI implementation ID, which is different for every SBI implementation. It is intended that this implementation ID allows software to probe for SBI implementation quirks.

### 4.3. Function: Get SBI implementation version (FID #2)

```
struct sbiret sbi_get_impl_version(void);
```

Returns the current SBI implementation version. The encoding of this version number is specific to the SBI implementation.

### 4.4. Function: Probe SBI extension (FID #3)

```
struct sbiret sbi_probe_extension(long extension_id);
```

Returns 0 if the given SBI extension ID (EID) is not available, or 1 if it is available unless defined as any other non-zero value by the implementation.

### 4.5. Function: Get machine vendor ID (FID #4)

```
struct sbiret sbi_get_mvvendorid(void);
```

Return a value that is legal for the `mvvendorid` CSR and 0 is always a legal value for this CSR.

## 4.6. Function: Get machine architecture ID (FID #5)

```
struct sbiret sbi_get_marchid(void);
```

Return a value that is legal for the `marchid` CSR and 0 is always a legal value for this CSR.

## 4.7. Function: Get machine implementation ID (FID #6)

```
struct sbiret sbi_get_mimpid(void);
```

Return a value that is legal for the `mimpid` CSR and 0 is always a legal value for this CSR.

## 4.8. Function Listing

Table 3. Base Function List

| Function Name                         | SBI Version | FID | EID  |
|---------------------------------------|-------------|-----|------|
| <code>sbi_get_sbi_spec_version</code> | 0.2         | 0   | 0x10 |
| <code>sbi_get_sbi_impl_id</code>      | 0.2         | 1   | 0x10 |
| <code>sbi_get_sbi_impl_version</code> | 0.2         | 2   | 0x10 |
| <code>sbi_probe_extension</code>      | 0.2         | 3   | 0x10 |
| <code>sbi_get_mvvendorid</code>       | 0.2         | 4   | 0x10 |
| <code>sbi_get_marchid</code>          | 0.2         | 5   | 0x10 |
| <code>sbi_get_mimpid</code>           | 0.2         | 6   | 0x10 |

## 4.9. SBI Implementation IDs

Table 4. SBI Implementation IDs

| Implementation ID | Name                       |
|-------------------|----------------------------|
| 0                 | Berkeley Boot Loader (BBL) |
| 1                 | OpenSBI                    |
| 2                 | Xvisor                     |



---

| Implementation ID | Name        |
|-------------------|-------------|
| 3                 | KVM         |
| 4                 | RustSBI     |
| 5                 | Diosix      |
| 6                 | Coffer      |
| 7                 | Xen Project |

## Chapter 5. Legacy Extensions (EIDs #0x00 - #0x0F)

The legacy SBI extensions follow a slightly different calling convention as compared to the SBI v0.2 (or higher) specification where:

- The SBI function ID field in **a6** register is ignored because these are encoded as multiple SBI extension IDs.
- Nothing is returned in **a1** register.
- All registers except **a0** must be preserved across an SBI call by the callee.
- The value returned in **a0** register is SBI legacy extension specific.

The page and access faults taken by the SBI implementation while accessing memory on behalf of the supervisor are redirected back to the supervisor with **sepc** CSR pointing to the faulting **ECALL** instruction.

The legacy SBI extensions is deprecated in favor of the other extensions listed below. The legacy console SBI functions (**sbi\_console\_getchar()** and **sbi\_console\_putchar()**) are expected to be deprecated; they have no replacement.

### 5.1. Extension: Set Timer (EID #0x00)

```
long sbi_set_timer(uint64_t stime_value)
```

Programs the clock for next event after **stime\_value** time. This function also clears the pending timer interrupt bit.

If the supervisor wishes to clear the timer interrupt without scheduling the next timer event, it can either request a timer interrupt infinitely far into the future (i.e., (uint64\_t)-1), or it can instead mask the timer interrupt by clearing **sie.STIE** CSR bit.

This SBI call returns 0 upon success or an implementation specific negative error code.

### 5.2. Extension: Console Puchar (EID #0x01)

```
long sbi_console_putchar(int ch)
```

Write data present in **ch** to debug console.

Unlike **sbi\_console\_getchar()**, this SBI call **will block** if there remain any pending characters to be transmitted or if the receiving terminal is not yet ready to receive the byte. However, if the console doesn't exist at all, then the character is thrown away.

This SBI call returns 0 upon success or an implementation specific negative error code.

## 5.3. Extension: Console Getchar (EID #0x02)

```
long sbi_console_getchar(void)
```

Read a byte from debug console.

The SBI call returns the byte on success, or -1 for failure.

## 5.4. Extension: Clear IPI (EID #0x03)

```
long sbi_clear_ipi(void)
```

Clears the pending IPIs if any. The IPI is cleared only in the hart for which this SBI call is invoked. `sbi_clear_ipi()` is deprecated because S-mode code can clear `sip.SSIP` CSR bit directly.

This SBI call returns 0 if no IPI had been pending, or an implementation specific positive value if an IPI had been pending.

## 5.5. Extension: Send IPI (EID #0x04)

```
long sbi_send_ipi(const unsigned long *hart_mask)
```

Send an inter-processor interrupt to all the harts defined in `hart_mask`. Interprocessor interrupts manifest at the receiving harts as Supervisor Software Interrupts.

`hart_mask` is a virtual address that points to a bit-vector of harts. The bit vector is represented as a sequence of unsigned longs whose length equals the number of harts in the system divided by the number of bits in an unsigned long, rounded up to the next integer.

This SBI call returns 0 upon success or an implementation specific negative error code.

## 5.6. Extension: Remote FENCE.I (EID #0x05)

```
long sbi_remote_fence_i(const unsigned long *hart_mask)
```

Instructs remote harts to execute `FENCE.I` instruction. The `hart_mask` is same as described in `sbi_send_ipi()`.

This SBI call returns 0 upon success or an implementation specific negative error code.

## 5.7. Extension: Remote SFENCE.VMA (EID #0x06)

```
long sbi_remote_sfence_vma(const unsigned long *hart_mask,
                           unsigned long start,
                           unsigned long size)
```

Instructs the remote harts to execute one or more **SFENCE.VMA** instructions, covering the range of virtual addresses between start and size.

This SBI call returns 0 upon success or an implementation specific negative error code.

## 5.8. Extension: Remote SFENCE.VMA with ASID (EID #0x07)

```
long sbi_remote_sfence_vma_asid(const unsigned long *hart_mask,
                                unsigned long start,
                                unsigned long size,
                                unsigned long asid)
```

Instruct the remote harts to execute one or more **SFENCE.VMA** instructions, covering the range of virtual addresses between start and size. This covers only the given **ASID**.

This SBI call returns 0 upon success or an implementation specific negative error code.

## 5.9. Extension: System Shutdown (EID #0x08)

```
void sbi_shutdown(void)
```

Puts all the harts to shutdown state from supervisor point of view.

This SBI call doesn't return irrespective whether it succeeds or fails.

## 5.10. Function Listing

Table 5. Legacy Function List

| Function Name       | SBI Version | FID | EID  | Replacement EID |
|---------------------|-------------|-----|------|-----------------|
| sbi_set_timer       | 0.1         | 0   | 0x00 | 0x54494D45      |
| sbi_console_putchar | 0.1         | 0   | 0x01 | N/A             |
| sbi_console_getchar | 0.1         | 0   | 0x02 | N/A             |
| sbi_clear_ipi       | 0.1         | 0   | 0x03 | N/A             |
| sbi_send_ipi        | 0.1         | 0   | 0x04 | 0x735049        |

| Function Name              | SBI Version | FID | EID       | Replacement EID |
|----------------------------|-------------|-----|-----------|-----------------|
| sbi_remote_fence_i         | 0.1         | 0   | 0x05      | 0x52464E43      |
| sbi_remote_sfence_vma      | 0.1         | 0   | 0x06      | 0x52464E43      |
| sbi_remote_sfence_vma_asid | 0.1         | 0   | 0x07      | 0x52464E43      |
| sbi_shutdown               | 0.1         | 0   | 0x08      | 0x53525354      |
| <b>RESERVED</b>            |             |     | 0x09-0x0F |                 |

# Chapter 6. Timer Extension (EID #0x54494D45 "TIME")

This replaces legacy timer extension (EID #0x00). It follows the new calling convention defined in v0.2.

## 6.1. Function: Set Timer (FID #0)

```
struct sbiret sbi_set_timer(uint64_t stime_value)
```

Programs the clock for next event after **stime\_value** time. **stime\_value** is in absolute time. This function must clear the pending timer interrupt bit as well.

If the supervisor wishes to clear the timer interrupt without scheduling the next timer event, it can either request a timer interrupt infinitely far into the future (i.e., (uint64\_t)-1), or it can instead mask the timer interrupt by clearing **sie.STIE** CSR bit.

## 6.2. Function Listing

Table 6. TIME Function List

| Function Name | SBI Version | FID | EID        |
|---------------|-------------|-----|------------|
| sbi_set_timer | 0.2         | 0   | 0x54494D45 |

## Chapter 7. IPI Extension (EID #0x735049 "sPI: s-mode IPI")

This extension replaces the legacy extension (EID #0x04). The other IPI related legacy extension(0x3) is deprecated now. All the functions in this extension follow the `hart_mask` as defined in the binary encoding section.

### 7.1. Function: Send IPI (FID #0)

```
struct sbiret sbi_send_ipi(unsigned long hart_mask,
                          unsigned long hart_mask_base)
```

Send an inter-processor interrupt to all the harts defined in `hart_mask`. Interprocessor interrupts manifest at the receiving harts as the supervisor software interrupts.

The possible error codes returned in `sbiret.error` are shown in the [Table 7](#) below.

Table 7. IPI Send Errors

| Error code  | Description  |
|-------------|--|
| SBI_SUCCESS | IPI was sent to all the targeted harts successfully. |

### 7.2. Function Listing

Table 8. IPI Function List

| Function Name             | SBI Version | FID | EID      |
|---------------------------|-------------|-----|----------|
| <code>sbi_send_ipi</code> | 0.2         | 0   | 0x735049 |

## Chapter 8. RFENCE Extension (EID #0x52464E43 "RFNC")

This extension defines all remote fence related functions and replaces the legacy extensions (EIDs #0x05 - #0x07). All the functions follow the `hart_mask` as defined in binary encoding section. Any function wishes to use range of addresses (i.e. `start_addr` and `size`), have to abide by the below constraints on range parameters.

The remote fence function acts as a full TLB flush if

- `start_addr` and `size` are both 0
- `size` is equal to  $2^{XLEN-1}$

### 8.1. Function: Remote FENCE.I (FID #0)

```
struct sbiret sbi_remote_fence_i(unsigned long hart_mask,
                                unsigned long hart_mask_base)
```

Instructs remote harts to execute `FENCE.I` instruction.

The possible error codes returned in `sbiret.error` are shown in the [Table 9](#) below.

Table 9. RFENCE Remote FENCE.I Errors

| Error code  | Description  |
|-------------|--|
| SBI_SUCCESS | IPI was sent to all the targeted harts successfully. |

### 8.2. Function: Remote SFENCE.VMA (FID #1)

```
struct sbiret sbi_remote_sfence_vma(unsigned long hart_mask,
                                     unsigned long hart_mask_base,
                                     unsigned long start_addr,
                                     unsigned long size)
```

Instructs the remote harts to execute one or more `SFENCE.VMA` instructions, covering the range of virtual addresses between `start` and `size`.

The possible error codes returned in `sbiret.error` are shown in the [Table 10](#) below.

Table 10. RFENCE Remote SFENCE.VMA Errors

| Error code              | Description  |
|-------------------------|--|
| SBI_SUCCESS             | IPI was sent to all the targeted harts successfully.       |
| SBI_ERR_INVALID_ADDRESS | <code>start_addr</code> or <code>size</code> is not valid. |



## 8.3. Function: Remote SFENCE.VMA with ASID (FID #2)

```
struct sbiret sbi_remote_sfence_vma_asid(unsigned long hart_mask,
                                         unsigned long hart_mask_base,
                                         unsigned long start_addr,
                                         unsigned long size,
                                         unsigned long asid)
```

Instruct the remote harts to execute one or more **SFENCE.VMA** instructions, covering the range of virtual addresses between start and size. This covers only the given **ASID**.

The possible error codes returned in **sbiret.error** are shown in the [Table 11](#) below.

Table 11. RFENCE Remote SFENCE.VMA with ASID Errors

| Error code              | Description  |
|-------------------------|--|
| SBI_SUCCESS             | IPI was sent to all the targeted harts successfully. |
| SBI_ERR_INVALID_ADDRESS | <b>start_addr</b> or <b>size</b> is not valid.       |

## 8.4. Function: Remote HFENCE.GVMA with VMID (FID #3)

```
struct sbiret sbi_remote_hfence_gvma_vmid(unsigned long hart_mask,
                                           unsigned long hart_mask_base,
                                           unsigned long start_addr,
                                           unsigned long size,
                                           unsigned long vmid)
```

Instruct the remote harts to execute one or more **HFENCE.GVMA** instructions, covering the range of guest physical addresses between start and size only for the given **VMID**. This function call is only valid for harts implementing hypervisor extension.

The possible error codes returned in **sbiret.error** are shown in the [Table 12](#) below.

Table 12. RFENCE Remote HFENCE.GVMA with VMID Errors

| Error code              | Description   |
|-------------------------|---|
| SBI_SUCCESS             | IPI was sent to all the targeted harts successfully.  |
| SBI_ERR_NOT_SUPPORTED   | This function is not supported as it is not implemented or one of the target hart doesn't support hypervisor extension. |
| SBI_ERR_INVALID_ADDRESS | <b>start_addr</b> or <b>size</b> is not valid.  |

## 8.5. Function: Remote HFENCE.GVMA (FID #4)

```
struct sbiret sbi_remote_hfence_gvma(unsigned long hart_mask,
                                     unsigned long hart_mask_base,
                                     unsigned long start_addr,
                                     unsigned long size)
```

Instruct the remote harts to execute one or more **HFENCE.GVMA** instructions, covering the range of guest physical addresses between start and size for all the guests. This function call is only valid for harts implementing hypervisor extension.

The possible error codes returned in **sbiret.error** are shown in the [Table 13](#) below.

Table 13. RFENCE Remote HFENCE.GVMA Errors

| Error code              | Description   |
|-------------------------|---|
| SBI_SUCCESS             | IPI was sent to all the targeted harts successfully.  |
| SBI_ERR_NOT_SUPPORTED   | This function is not supported as it is not implemented or one of the target hart doesn't support hypervisor extension. |
| SBI_ERR_INVALID_ADDRESS | <b>start_addr</b> or <b>size</b> is not valid.  |

## 8.6. Function: Remote HFENCE.VVMA with ASID (FID #5)

```
struct sbiret sbi_remote_hfence_vvma_asid(unsigned long hart_mask,
                                           unsigned long hart_mask_base,
                                           unsigned long start_addr,
                                           unsigned long size,
                                           unsigned long asid)
```

Instruct the remote harts to execute one or more **HFENCE.VVMA** instructions, covering the range of guest virtual addresses between start and size for the given **ASID** and current **VMID** (in **hgap** CSR) of calling hart. This function call is only valid for harts implementing hypervisor extension.

The possible error codes returned in **sbiret.error** are shown in the [Table 14](#) below.

Table 14. RFENCE Remote HFENCE.VVMA with ASID Errors

| Error code              | Description   |
|-------------------------|---|
| SBI_SUCCESS             | IPI was sent to all the targeted harts successfully.  |
| SBI_ERR_NOT_SUPPORTED   | This function is not supported as it is not implemented or one of the target hart doesn't support hypervisor extension. |
| SBI_ERR_INVALID_ADDRESS | <b>start_addr</b> or <b>size</b> is not valid.  |

## 8.7. Function: Remote HFENCE.VVMA (FID #6)

```
struct sbiret sbi_remote_hfence_vvma(unsigned long hart_mask,
                                     unsigned long hart_mask_base,
                                     unsigned long start_addr,
                                     unsigned long size)
```

Instruct the remote harts to execute one or more **HFENCE.VVMA** instructions, covering the range of guest virtual addresses between start and size for current **VMID** (in **hgatp** CSR) of calling hart. This function call is only valid for harts implementing hypervisor extension.

The possible error codes returned in **sbiret.error** are shown in the [Table 15](#) below.

Table 15. RFENCE Remote HFENCE.VVMA Errors

| Error code              | Description   |
|-------------------------|---|
| SBI_SUCCESS             | IPI was sent to all the targeted harts successfully.  |
| SBI_ERR_NOT_SUPPORTED   | This function is not supported as it is not implemented or one of the target hart doesn't support hypervisor extension. |
| SBI_ERR_INVALID_ADDRESS | <b>start_addr</b> or <b>size</b> is not valid.  |

## 8.8. Function Listing

Table 16. RFENCE Function List

| Function Name               | SBI Version | FID | EID        |
|-----------------------------|-------------|-----|------------|
| sbi_remote_fence_i          | 0.2         | 0   | 0x52464E43 |
| sbi_remote_sfence_vma       | 0.2         | 1   | 0x52464E43 |
| sbi_remote_sfence_vma_asid  | 0.2         | 2   | 0x52464E43 |
| sbi_remote_hfence_gvma_vmid | 0.2         | 3   | 0x52464E43 |
| sbi_remote_hfence_gvma      | 0.2         | 4   | 0x52464E43 |
| sbi_remote_hfence_vvma_asid | 0.2         | 5   | 0x52464E43 |
| sbi_remote_hfence_vvma      | 0.2         | 6   | 0x52464E43 |

# Chapter 9. Hart State Management Extension (EID #0x48534D "HSM")

The Hart State Management (HSM) Extension introduces a set of hart states and a set of functions which allow the supervisor-mode software to request a hart state change.

The [Table 17](#) shown below describes all possible **HSM states** along with a unique **HSM state id** for each state:

Table 17. HSM Hart States

| State ID | State Name      | Description  |
|----------|-----------------|--|
| 0        | STARTED         | The hart is physically powered-up and executing normally.  |
| 1        | STOPPED         | The hart is not executing in supervisor-mode or any lower privilege mode. It is probably powered-down by the SBI implementation if the underlying platform has a mechanism to physically power-down harts.               |
| 2        | START_PENDING   | Some other hart has requested to start (or power-up) the hart from the <b>STOPPED</b> state and the SBI implementation is still working to get the hart in the <b>STARTED</b> state.                                     |
| 3        | STOP_PENDING    | The hart has requested to stop (or power-down) itself from the <b>STARTED</b> state and the SBI implementation is still working to get the hart in the <b>STOPPED</b> state.   |
| 4        | SUSPENDED       | This hart is in a platform specific suspend (or low power) state.  |
| 5        | SUSPEND_PENDING | The hart has requested to put itself in a platform specific low power state from the <b>STARTED</b> state and the SBI implementation is still working to get the hart in the platform specific <b>SUSPENDED</b> state.   |
| 6        | RESUME_PENDING  | An interrupt or platform specific hardware event has caused the hart to resume normal execution from the <b>SUSPENDED</b> state and the SBI implementation is still working to get the hart in the <b>STARTED</b> state. |

At any point in time, a hart should be in one of the above mentioned hart states. The hart state transitions by the SBI implementation should follow the state machine shown below in the [Figure 3](#).




Figure 3. SBI HSM State Machine

A platform can have multiple harts grouped into hierarchical topology groups (namely cores, clusters, nodes, etc.) with separate platform specific low-power states for each hierarchical group. These platform specific low-power states of hierarchical topology groups can be represented as platform specific suspend states of a hart. An SBI implementation can utilize the suspend states of higher topology groups using one of the following approaches:

1. **Platform-coordinated:** In this approach, when a hart becomes idle the supervisor-mode power-managment software will request deepest suspend state for the hart and higher topology groups. An SBI implementation should choose a suspend state at higher topology group which is:
  - a. Not deeper than the specified suspend state
  - b. Wake-up latency is not higher than the wake-up latency of the specified suspend state
2. **OS-initiated:** In this approach, the supervisor-mode power-managment software will directly request a suspend state for higher topology group after the last hart in that group becomes idle. When a hart becomes idle, the supervisor-mode power-managment software will always select suspend state for the hart itself but it will select a suspend state for a higher topology group only if the hart is the last running hart in the group. An SBI implementation should:
  - a. Never choose a suspend state for higher topology group different from the specified suspend state
  - b. Always prefer most recent suspend state requested for higher topology group

## 9.1. Function: HART start (FID #0)

```

struct sbiret sbi_hart_start(unsigned long hartid,
                           unsigned long start_addr,
                           unsigned long opaque)

```

Request the SBI implementation to start executing the target hart in supervisor-mode, at the address specified by `start_addr`, with the specific register values described in [Table 18](#).

Table 18. HSM Hart Start Register State

| Register Name                                     | Register Value   |
|---|------------------|
| satp  | 0                |
| sstatus.SIE                                       | 0                |
| a0  | hartid           |
| a1  | opaque parameter |
| All other registers remain in an undefined state. |                  |



A single **unsigned long** parameter is sufficient as `start_addr`, because the hart will start execution in supervisor-mode with the MMU off, hence `start_addr` must be less than XLEN bits wide.

This call is asynchronous — more specifically, the `sbi_hart_start()` may return before the target hart starts executing as long as the SBI implementation is capable of ensuring the return code is accurate. If the SBI implementation is a platform runtime firmware executing in machine-mode (M-mode), then it **MUST** configure any physical memory protection it supports, such as that defined by PMP, and other M-mode state, before transferring control to supervisor-mode software.

The `hartid` parameter specifies the target hart which is to be started.

The `start_addr` parameter points to a runtime-specified physical address, where the hart can start executing in supervisor-mode.

The `opaque` parameter is an XLEN-bit value which will be set in the `a1` register when the hart starts executing at `start_addr`.

The possible error codes returned in `sbiret.error` are shown in the [Table 19](#) below.

Table 19. HSM Hart Start Errors

| Error code                | Description   |
|---------------------------|---|
| SBI_SUCCESS               | Hart was previously in stopped state. It will start executing from <code>start_addr</code> .  |
| SBI_ERR_INVALID_ADDRESS   | <code>start_addr</code> is not valid, possibly due to the following reasons:<br>* It is not a valid physical address.<br>* Executable access to the address is prohibited by a physical memory protection mechanism or H-extension G-stage for supervisor-mode. |
| SBI_ERR_INVALID_PARAM     | <code>hartid</code> is not a valid hartid as the corresponding hart cannot be started in supervisor mode.   |
| SBI_ERR_ALREADY_AVAILABLE | The given hartid is already started.  |
| SBI_ERR_FAILED            | The start request failed for unspecified or unknown other reasons.  |

## 9.2. Function: HART stop (FID #1)

```
struct sbiret sbi_hart_stop(void)
```

Request the SBI implementation to stop executing the calling hart in supervisor-mode and return its ownership to the SBI implementation. This call is not expected to return under normal conditions. The `sbi_hart_stop()` must be called with supervisor-mode interrupts disabled.

The possible error codes returned in `sbiret.error` are shown in the [Table 20](#) below.

Table 20. HSM Hart Stop Errors

| Error code     | Description                                  |
|----------------|--|
| SBI_ERR_FAILED | Failed to stop execution of the current hart |

## 9.3. Function: HART get status (FID #2)

```
struct sbiret sbi_hart_get_status(unsigned long hartid)
```

Get the current status (or HSM state id) of the given hart in `sbiret.value`, or an error through `sbiret.error`.

The `hartid` parameter specifies the target hart for which status is required.

The possible status (or HSM state id) values returned in `sbiret.value` are described in [Table 17](#).

The possible error codes returned in `sbiret.error` are shown in the [Table 21](#) below.

Table 21. HSM Hart Get Status Errors

| Error code            | Description                                 |
|-----------------------|---|
| SBI_ERR_INVALID_PARAM | The given <code>hartid</code> is not valid. |

The harts may transition HSM states at any time due to any concurrent `sbi_hart_start()` or `sbi_hart_stop()` or `sbi_hart_suspend()` calls, the return value from this function may not represent the actual state of the hart at the time of return value verification.

## 9.4. Function: HART suspend (FID #3)

```
struct sbiret sbi_hart_suspend(uint32_t suspend_type,
                               unsigned long resume_addr,
                               unsigned long opaque)
```

Request the SBI implementation to put the calling hart in a platform specific suspend (or low power) state specified by the `suspend_type` parameter. The hart will automatically come out of suspended state and resume normal execution when it receives an interrupt or platform specific hardware event.

The platform specific suspend states for a hart can be either retentive or non-retentive in nature. A retentive suspend state will preserve hart register and CSR values for all privilege modes whereas a non-retentive suspend state will not preserve hart register and CSR values.

Resuming from a retentive suspend state is straight forward and the supervisor-mode software will see SBI suspend call return without any failures. The `resume_addr` parameter is unused during retentive suspend.

Resuming from a non-retentive suspend state is relatively more involved and requires software to restore various hart registers and CSRs for all privilege modes. Upon resuming from non-retentive suspend state, the hart will jump to supervisor-mode at address specified by `resume_addr` with specific registers values described in the [Table 22](#) below.

Table 22. HSM Hart Resume Register State

| Register Name                                     | Register Value                |
|---|-------------------------------|
| satp  | 0                             |
| sstatus.SIE                                       | 0                             |
| a0  | hartid                        |
| a1  | <code>opaque</code> parameter |
| All other registers remain in an undefined state. |                               |



A single `unsigned long` parameter is sufficient for `resume_addr`, because the hart will resume execution in supervisor-mode with the MMU off, hence `resume_addr` must be less than XLEN bits wide.

The `suspend_type` parameter is 32 bits wide and the possible values are shown in [Table 23](#) below.

Table 23. HSM Hart Suspend Types

| Value                   | Description                             |
|-------------------------|---|
| 0x00000000              | Default retentive suspend               |
| 0x00000001 - 0x0FFFFFFF | Reserved for future use                 |
| 0x10000000 - 0x7FFFFFFF | Platform specific retentive suspend     |
| 0x80000000              | Default non-retentive suspend           |
| 0x80000001 - 0x8FFFFFFF | Reserved for future use                 |
| 0x90000000 - 0xFFFFFFFF | Platform specific non-retentive suspend |
| > 0xFFFFFFFF            | Reserved                                |

The `resume_addr` parameter points to a runtime-specified physical address, where the hart can resume execution in supervisor-mode after a non-retentive suspend.

The `opaque` parameter is an XLEN-bit value which will be set in the `a1` register when the hart resumes execution at `resume_addr` after a non-retentive suspend.

The possible error codes returned in `sbiret.error` are shown in the [Table 24](#) below.



Table 24. HSM Hart Suspend Errors

| Error code              | Description  |
|-------------------------|--|
| SBI_SUCCESS             | Hart has suspended and resumed successfully from a retentive suspend state.  |
| SBI_ERR_INVALID_PARAM   | <code>suspend_type</code> is reserved or is platform-specific and unimplemented.   |
| SBI_ERR_NOT_SUPPORTED   | <code>suspend_type</code> is not reserved and is implemented, but the platform does not support it due to one or more missing dependencies.  |
| SBI_ERR_INVALID_ADDRESS | <code>resume_addr</code> is not valid, possibly due to the following reasons:<br>* It is not a valid physical address.<br>* Executable access to the address is prohibited by a physical memory protection mechanism or H-extension G-stage for supervisor-mode. |
| SBI_ERR_FAILED          | The suspend request failed for unspecified or unknown other reasons.   |

## 9.5. Function Listing

Table 25. HSM Function List

| Function Name                    | SBI Version | FID | EID      |
|----------------------------------|-------------|-----|----------|
| <code>sbi_hart_start</code>      | 0.2         | 0   | 0x48534D |
| <code>sbi_hart_stop</code>       | 0.2         | 1   | 0x48534D |
| <code>sbi_hart_get_status</code> | 0.2         | 2   | 0x48534D |
| <code>sbi_hart_suspend</code>    | 0.3         | 3   | 0x48534D |

# Chapter 10. System Reset Extension (EID #0x53525354 "SRST")

The System Reset Extension provides a function that allow the supervisor software to request system-level reboot or shutdown. The term "system" refers to the world-view of supervisor software and the underlying SBI implementation could be provided by machine mode firmware or a hypervisor.

## 10.1. Function: System reset (FID #0)

```
struct sbiret sbi_system_reset(uint32_t reset_type, uint32_t reset_reason)
```

Reset the system based on provided `reset_type` and `reset_reason`. This is a synchronous call and does not return if it succeeds.

The `reset_type` parameter is 32 bits wide and it's possible values are shown in the [Table 26](#) below.

Table 26. SRST System Reset Types

| Value                   | Description                            |
|-------------------------|--|
| 0x00000000              | Shutdown                               |
| 0x00000001              | Cold reboot                            |
| 0x00000002              | Warm reboot                            |
| 0x00000003 - 0xEFFFFFFF | Reserved for future use                |
| 0xF0000000 - 0xFFFFFFFF | Vendor or platform specific reset type |
| > 0xFFFFFFFF            | Reserved                               |

The `reset_reason` is an optional parameter representing the reason for system reset. This parameter is 32 bits wide with possible values shown in the [Table 27](#) below

Table 27. SRST System Reset Reasons

| Value                   | Description                              |
|-------------------------|--|
| 0x00000000              | No reason                                |
| 0x00000001              | System failure                           |
| 0x00000002 - 0xDFFFFFFF | Reserved for future use                  |
| 0xE0000000 - 0xEFFFFFFF | SBI implementation specific reset reason |
| 0xF0000000 - 0xFFFFFFFF | Vendor or platform specific reset reason |
| > 0xFFFFFFFF            | Reserved                                 |

When supervisor software is running natively, the SBI implementation is provided by machine mode firmware. In this case, shutdown is equivalent to a physical power down of the entire system and cold reboot is equivalent to a physical power cycle of the entire system. Further, warm reboot is equivalent to a power cycle of the main processor and parts of the system, but not the entire system. For example, on a server class system with a BMC (board management controller), a warm reboot will not power cycle the BMC whereas a cold reboot will definitely power cycle the BMC.

When supervisor software is running inside a virtual machine, the SBI implementation is provided by a hypervisor. Shutdown, cold reboot and warm reboot will behave functionally the same as the native case, but might not result in any physical power changes.

The possible error codes returned in `sbiret.error` are shown in the [Table 28](#) below.

Table 28. SRST System Reset Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_ERR_INVALID_PARAM | At least one of <code>reset_type</code> or <code>reset_reason</code> is reserved or is platform-specific and unimplemented.               |
| SBI_ERR_NOT_SUPPORTED | <code>reset_type</code> is not reserved and is implemented, but the platform does not support it due to one or more missing dependencies. |
| SBI_ERR_FAILED        | The reset request failed for unspecified or unknown other reasons.  |

## 10.2. Function Listing

Table 29. SRST Function List

| Function Name                 | SBI Version | FID | EID        |
|-------------------------------|-------------|-----|------------|
| <code>sbi_system_reset</code> | 0.3         | 0   | 0x53525354 |

# Chapter 11. Performance Monitoring Unit Extension (EID #0x504D55 "PMU")

The RISC-V hardware performance counters such as `mcycle`, `minstret`, and `mhpcounterX` CSRs are accessible as read-only from supervisor-mode using `cycle`, `instret`, and `hpmcounterX` CSRs. The SBI performance monitoring unit (PMU) extension is an interface for supervisor-mode to configure and use the RISC-V hardware performance counters with assistance from the machine-mode (or hypervisor-mode). These hardware performance counters can only be started, stopped, or configured from machine-mode using `mcountinhibit` and `mhpmeventX` CSRs. Due to this, a machine-mode SBI implementation may choose to disallow SBI PMU extension if `mcountinhibit` CSR is not implemented by the RISC-V platform.

A RISC-V platform generally supports monitoring of various hardware events using a limited number of hardware performance counters which are up to 64 bits wide. In addition, a SBI implementation can also provide firmware performance counters which can monitor firmware events such as number of misaligned load/store instructions, number of RFENCEs, number of IPIs, etc. All firmware counters must have same number of bits and can be up to 64 bits wide.

The SBI PMU extension provides:

1. An interface for supervisor-mode software to discover and configure per-HART hardware/firmware counters
2. A typical `perf` compatible interface for hardware/firmware performance counters and events
3. Full access to microarchitecture's raw event encodings

To define SBI PMU extension calls, we first define important entities `counter_idx`, `event_idx`, and `event_data`. The `counter_idx` is a logical number assigned to each hardware/firmware counter. The `event_idx` represents a hardware (or firmware) event whereas the `event_data` is 64 bits wide and represents additional configuration (or parameters) for a hardware (or firmware) event.

The `event_idx` is a 20 bits wide number encoded as follows:

```
event_idx[19:16] = type
event_idx[15:0] = code
```

## 11.1. Event: Hardware general events (Type #0)

The `event_idx.type` (i.e. **event type**) should be `0x0` for all hardware general events and each hardware general event is identified by a unique `event_idx.code` (i.e. **event code**) described in the [Table 30](#) below.

Table 30. PMU Hardware Events

| General Event Name  | Code | Description  |
|---------------------|------|--|
| SBI_PMU_HW_NO_EVENT | 0    | Unused event because <code>event_idx</code> cannot be zero |

| General Event Name                 | Code | Description   |
|------------------------------------|------|---|
| SBI_PMU_HW_CPU_CYCLES              | 1    | Event for each CPU cycle                                |
| SBI_PMU_HW_INSTRUCTIONS            | 2    | Event for each completed instruction                    |
| SBI_PMU_HW_CACHE_REFERENCES        | 3    | Event for cache hit                                     |
| SBI_PMU_HW_CACHE_MISSES            | 4    | Event for cache miss                                    |
| SBI_PMU_HW_BRANCH_INSTRUCTIONS     | 5    | Event for a branch instruction                          |
| SBI_PMU_HW_BRANCH_MISSES           | 6    | Event for a branch misprediction                        |
| SBI_PMU_HW_BUS_CYCLES              | 7    | Event for each BUS cycle                                |
| SBI_PMU_HW_STALLED_CYCLES_FRONTEND | 8    | Event for a stalled cycle in microarchitecture frontend |
| SBI_PMU_HW_STALLED_CYCLES_BACKEND  | 9    | Event for a stalled cycle in microarchitecture backend  |
| SBI_PMU_HW_REF_CPU_CYCLES          | 10   | Event for each reference CPU cycle                      |

The **event\_data** (i.e. **event data**) is unused for hardware general events and all non-zero values of **event\_data** are reserved for future use.



A RISC-V platform might halt the CPU clock when it enters WAIT state using the WFI instruction or enters platform specific SUSPEND state using the SBI HSM HART suspend call.



The **SBI\_PMU\_HW\_CPU\_CYCLES** event counts CPU clock cycles as counted by the **cycle** CSR. These may be variable frequency cycles, and are not counted when the CPU clock is halted.



The **SBI\_PMU\_HW\_REF\_CPU\_CYCLES** counts fixed-frequency clock cycles while the CPU clock is not halted. The fixed-frequency of counting might, for example, be the same frequency at which the **time** CSR counts.



The **SBI\_PMU\_HW\_BUS\_CYCLES** counts fixed-frequency clock cycles. The fixed-frequency of counting might be the same frequency at which the **time** CSR counts, or may be the frequency of the clock at the boundary between the HART (and its private caches) and the rest of the system.

## 11.2. Event: Hardware cache events (Type #1)

The **event\_idx.type** (i.e. **event type**) should be **0x1** for all hardware cache events and each hardware cache event is identified by a unique **event\_idx.code** (i.e. **event code**) which is encoded as follows:

```

event_idx.code[15:3] = cache_id
event_idx.code[2:1] = op_id
event_idx.code[0:0] = result_id

```

Below tables show possible values of: `event_idx.code.cache_id` (i.e. cache event id), `event_idx.code.op_id` (i.e. cache operation id) and `event_idx.code.result_id` (i.e. cache result id).

Table 31. PMU Cache Event ID

| Cache Event Name      | Event ID | Description                    |
|-----------------------|----------|--------------------------------|
| SBI_PMU_HW_CACHE_L1D  | 0        | Level1 data cache event        |
| SBI_PMU_HW_CACHE_L1I  | 1        | Level1 instruction cache event |
| SBI_PMU_HW_CACHE_LL   | 2        | Last level cache event         |
| SBI_PMU_HW_CACHE_DTLB | 3        | Data TLB event                 |
| SBI_PMU_HW_CACHE_ITLB | 4        | Instruction TLB event          |
| SBI_PMU_HW_CACHE_BPU  | 5        | Branch predictor unit event    |
| SBI_PMU_HW_CACHE_NODE | 6        | NUMA node cache event          |

Table 32. PMU Cache Operation ID

| Cache Operation Name         | Operation ID | Description         |
|------------------------------|--------------|---------------------|
| SBI_PMU_HW_CACHE_OP_READ     | 0            | Read cache line     |
| SBI_PMU_HW_CACHE_OP_WRITE    | 1            | Write cache line    |
| SBI_PMU_HW_CACHE_OP_PREFETCH | 2            | Prefetch cache line |

Table 33. PMU Cache Operation Result ID

| Cache Result Name              | Result ID | Description  |
|--------------------------------|-----------|--------------|
| SBI_PMU_HW_CACHE_RESULT_ACCESS | 0         | Cache access |
| SBI_PMU_HW_CACHE_RESULT_MISS   | 1         | Cache miss   |

The `event_data` (i.e. event data) is unused for hardware cache events and all non-zero values of `event_data` are reserved for future use.

## 11.3. Event: Hardware raw events (Type #2)

The `event_idx.type` (i.e. event type) should be `0x2` for all hardware raw events and `event_idx.code` (i.e. event code) should be zero.

On RISC-V platform with 32 bits wide `mhpmeventX` CSRs, the `event_data` configuration (or parameter) should have the 32-bit value to to be programmed in the `mhpmeventX` CSR.

On RISC-V platform with 64 bits wide `mhpmeventX` CSRs, the `event_data` configuration (or parameter) should have the 48-bit value to to be programmed in the lower 48-bits of `mhpmeventX` CSR and the SBI implementation shall determine the value to be programmed in the upper 16 bits of

`mhpmeventX` CSR.



The RISC-V platform hardware implementation may choose to define the expected value to be written to `mhpmeventX` CSR for a hardware event. In case of hardware general/cache events, the RISC-V platform hardware implementation may use the zero-extended `event_idx` as the expected value for simplicity.

## 11.4. Event: Firmware events (Type #15)

The `event_idx.type` (i.e. **event type**) should be `0xf` for all firmware events and each firmware event is identified by an unique `event_idx.code` (i.e. **event code**) described in the [Table 34](#) below.

Table 34. PMU Firmware Events

| Firmware Event Name                 | Code | Description   |
|-------------------------------------|------|---|
| SBI_PMU_FW_MISALIGNED_LOAD          | 0    | Misaligned load trap event                                  |
| SBI_PMU_FW_MISALIGNED_STORE         | 1    | Misaligned store trap event                                 |
| SBI_PMU_FW_ACCESS_LOAD              | 2    | Load access trap event                                      |
| SBI_PMU_FW_ACCESS_STORE             | 3    | Store access trap event                                     |
| SBI_PMU_FW_ILLEGAL_INSN             | 4    | Illegal instruction trap event                              |
| SBI_PMU_FW_SET_TIMER                | 5    | Set timer event   |
| SBI_PMU_FW_IPI_SENT                 | 6    | Sent IPI to other HART event                                |
| SBI_PMU_FW_IPI_RECEIVED             | 7    | Received IPI from other HART event                          |
| SBI_PMU_FW_FENCE_I_SENT             | 8    | Sent FENCE.I request to other HART event                    |
| SBI_PMU_FW_FENCE_I_RECEIVED         | 9    | Received FENCE.I request from other HART event              |
| SBI_PMU_FW_SFENCE_VMA_SENT          | 10   | Sent SFENCE.VMA request to other HART event                 |
| SBI_PMU_FW_SFENCE_VMA_RECEIVED      | 11   | Received SFENCE.VMA request from other HART event           |
| SBI_PMU_FW_SFENCE_VMA_ASID_SENT     | 12   | Sent SFENCE.VMA with ASID request to other HART event       |
| SBI_PMU_FW_SFENCE_VMA_ASID_RECEIVED | 13   | Received SFENCE.VMA with ASID request from other HART event |
| SBI_PMU_FW_HFENCE_GVMA_SENT         | 14   | Sent HFENCE.GVMA request to other HART event                |
| SBI_PMU_FW_HFENCE_GVMA_RECEIVED     | 15   | Received HFENCE.GVMA request from other HART event          |

| Firmware Event Name                  | Code        | Description   |
|--------------------------------------|-------------|---|
| SBI_PMU_FW_HFENCE_GVMA_VMID_SENT     | 16          | Sent HFENCE.GVMA with VMID request to other HART event  |
| SBI_PMU_FW_HFENCE_GVMA_VMID_RECEIVED | 17          | Received HFENCE.GVMA with VMID request from other HART event  |
| SBI_PMU_FW_HFENCE_VVMA_SENT          | 18          | Sent HFENCE.VVMA request to other HART event  |
| SBI_PMU_FW_HFENCE_VVMA_RECEIVED      | 19          | Received HFENCE.VVMA request from other HART event  |
| SBI_PMU_FW_HFENCE_VVMA_ASID_SENT     | 20          | Sent HFENCE.VVMA with ASID request to other HART event  |
| SBI_PMU_FW_HFENCE_VVMA_ASID_RECEIVED | 21          | Received HFENCE.VVMA with ASID request from other HART event  |
| Reserved                             | 22 - 255    | Reserved for future use   |
| Implementation specific events       | 256 - 65534 | SBI implementation specific firmware events   |
| SBI_PMU_FW_PLATFORM                  | 65535       | RISC-V platform specific firmware events, where the <code>event_data</code> configuration (or parameter) contains the event encoding. |

For all firmware events except SBI\_PMU\_FW\_PLATFORM, the `event_data` configuration (or parameter) is unused and all non-zero values of `event_data` are reserved for future use.

## 11.5. Function: Get number of counters (FID #0)

```
struct sbiret sbi_pmu_num_counters()
```

**Returns** the number of counters (both hardware and firmware) in `sbiret.value` and always returns `SBI_SUCCESS` in `sbiret.error`.

## 11.6. Function: Get details of a counter (FID #1)

```
struct sbiret sbi_pmu_counter_get_info(unsigned long counter_idx)
```

Get details about the specified counter such as underlying CSR number, width of the counter, type of counter hardware/firmware, etc.



The `counter_info` returned by this SBI call is encoded as follows:

```
counter_info[11:0] = CSR (12bit CSR number)
counter_info[17:12] = Width (One less than number of bits in CSR)
counter_info[XLEN-2:18] = Reserved for future use
counter_info[XLEN-1] = Type (0 = hardware and 1 = firmware)
```

If `counter_info.type == 1` then `counter_info.csr` and `counter_info.width` should be ignored.

Returns the `counter_info` described above in `sbiret.value`.

The possible error codes returned in `sbiret.error` are shown in the [Table 35](#) below.

Table 35. PMU Counter Get Info Errors

| Error code            | Description  |
|-----------------------|--|
| SBI_SUCCESS           | <code>counter_info</code> read successfully.           |
| SBI_ERR_INVALID_PARAM | <code>counter_idx</code> points to an invalid counter. |

## 11.7. Function: Find and configure a matching counter (FID #2)

```
struct sbiret sbi_pmu_counter_config_matching(unsigned long counter_idx_base,
                                             unsigned long counter_idx_mask,
                                             unsigned long config_flags,
                                             unsigned long event_idx,
                                             uint64_t event_data)
```

Find and configure a counter from a set of counters which is not started (or enabled) and can monitor the specified event. The `counter_idx_base` and `counter_idx_mask` parameters represent the set of counters whereas `event_idx` represents the event to be monitored and `event_data` represents any additional event configuration.

The `config_flags` parameter represents additional counter configuration and filter flags. The bit definitions of the `config_flags` parameter are shown in the [Table 36](#) below.

Table 36. PMU Counter Config Match Flags

| Flag Name                    | Bits | Description  |
|------------------------------|------|--|
| SBI_PMU_CFG_FLAG_SKIP_MATCH  | 0:0  | Skip the counter matching                                  |
| SBI_PMU_CFG_FLAG_CLEAR_VALUE | 1:1  | Clear (or zero) the counter value in counter configuration |
| SBI_PMU_CFG_FLAG_AUTO_START  | 2:2  | Start the counter after configuring a matching counter     |

| Flag Name                  | Bits       | Description                                     |
|----------------------------|------------|---|
| SBI_PMU_CFG_FLAG_SET_VUINH | 3:3        | Event counting inhibited in VU-mode             |
| SBI_PMU_CFG_FLAG_SET_VSINH | 4:4        | Event counting inhibited in VS-mode             |
| SBI_PMU_CFG_FLAG_SET_UINH  | 5:5        | Event counting inhibited in U-mode              |
| SBI_PMU_CFG_FLAG_SET_SINH  | 6:6        | Event counting inhibited in S-mode              |
| SBI_PMU_CFG_FLAG_SET_MINH  | 7:7        | Event counting inhibited in M-mode              |
| RESERVED                   | 8:(XLEN-1) | All non-zero values are reserved for future use |



When `SBI_PMU_CFG_FLAG_SKIP_MATCH` is set in `config_flags`, the SBI implementation will unconditionally select the first counter from the set of counters specified by the `counter_idx_base` and `counter_idx_mask`.



The `SBI_PMU_CFG_FLAG_AUTO_START` flag in `config_flags` has no impact on the counter value.



The `config_flags[3:7]` bits are event filtering hints so these can be ignored or overridden by the SBI implementation for security concerns or due to lack of event filtering support in the underlying RISC-V platform.

Returns the `counter_idx` in `sbiret.value` upon success.

In case of failure, the possible error codes returned in `sbiret.error` are shown in the Table 37 below.

Table 37. PMU Counter Config Match Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | counter found and configured successfully.            |
| SBI_ERR_INVALID_PARAM | set of counters has at least one invalid counter.     |
| SBI_ERR_NOT_SUPPORTED | none of the counters can monitor the specified event. |

## 11.8. Function: Start a set of counters (FID #3)

```
struct sbiret sbi_pmu_counter_start(unsigned long counter_idx_base,
                                   unsigned long counter_idx_mask,
                                   unsigned long start_flags,
                                   uint64_t initial_value)
```

Start or enable a set of counters on the calling HART with the specified initial value. The `counter_idx_base` and `counter_idx_mask` parameters represent the set of counters whereas the

`initial_value` parameter specifies the initial value of the counter.

The bit definitions of the `start_flags` parameter are shown in the Table 38 below.

Table 38. PMU Counter Start Flags

| Flag Name                        | Bits       | Description   |
|----------------------------------|------------|---|
| SBI_PMU_START_SET_INIT_VALUE     | 0:0        | Set the value of counters based on the <code>initial_value</code> parameter |
| SBI_PMU_START_FLAG_INIT_SNAPSHOT | 1:1        | Initialize the given counters from shared memory if available.              |
| RESERVED                         | 2:(XLEN-1) | Reserved for future use   |



When `SBI_PMU_START_SET_INIT_VALUE` is not set in `start_flags`, the counter value will not be modified and event counting will start from current counter value.

The shared memory address must be set during boot via `sbi_pmu_snapshot_set_shmem` before the `SBI_PMU_START_FLAG_INIT_SNAPSHOT` flag may be used. The SBI implementation must initialize all the given valid counters (to be started) from the value set in the shared snapshot memory.



`SBI_PMU_START_SET_INIT_VALUE` and `SBI_PMU_START_FLAG_INIT_SNAPSHOT` are mutually exclusive as the former is only valid for a single counter.

The possible error codes returned in `sbiret.error` are shown in the Table 39 below.

Table 39. PMU Counter Start Errors

| Error code              | Description  |
|-------------------------|--|
| SBI_SUCCESS             | counter started successfully.  |
| SBI_ERR_INVALID_PARAM   | set of counters has at least one invalid counter. or the snapshot address is not configured and <code>SBI_PMU_START_FLAG_INIT_SNAPSHOT</code> is set in the flags. |
| SBI_ERR_ALREADY_STARTED | set of counters includes at least one counter which is already started.  |

## 11.9. Function: Stop a set of counters (FID #4)

```
struct sbiret sbi_pmu_counter_stop(unsigned long counter_idx_base,
                                   unsigned long counter_idx_mask,
                                   unsigned long stop_flags)
```

Stop or disable a set of counters on the calling HART. The `counter_idx_base` and `counter_idx_mask` parameters represent the set of counters. The bit definitions of the `stop_flags` parameter are shown in the Table 40 below.

Table 40. PMU Counter Stop Flags

| Flag Name                       | Bits       | Description  |
|---------------------------------|------------|--|
| SBI_PMU_STOP_FLAG_RESET         | 0:0        | Reset the counter to event mapping.  |
| SBI_PMU_STOP_FLAG_TAKE_SNAPSHOT | 1:1        | Save a snapshot of the given counter's values in the shared memory if available. |
| RESERVED                        | 2:(XLEN-1) | Reserved for future use  |

The shared memory address must be set during boot via `sbi_pmu_snapshot_set_shmem` before the `SBI_PMU_STOP_FLAG_TAKE_SNAPSHOT` flag may be used. The SBI implementation must save the current value of all the stopped counters in the shared memory if `SBI_PMU_STOP_FLAG_TAKE_SNAPSHOT` is set. The values corresponding to all other counters must not be modified. The SBI implementation must additionally update the overflown counter bitmap in the shared memory.

The possible error codes returned in `sbiret.error` are shown in the Table 41 below.

Table 41. PMU Counter Stop Errors

| Error code              | Description   |
|-------------------------|---|
| SBI_SUCCESS             | counter stopped successfully.   |
| SBI_ERR_INVALID_PARAM   | set of counters has at least one invalid counter. Or the snapshot address is not configured and <code>SBI_PMU_STOP_FLAG_TAKE_SNAPSHOT</code> is set in the flags. |
| SBI_ERR_ALREADY_STOPPED | set of counters includes at least one counter which is already stopped.   |

## 11.10. Function: Read a firmware counter (FID #5)

```
struct sbiret sbi_pmu_counter_fw_read(unsigned long counter_idx)
```

Provide the current firmware counter value in `sbiret.value`. On RV32 systems, the `sbiret.value` will only contain the lower 32 bits of the current firmware counter value.

The possible error codes returned in `sbiret.error` are shown in the Table 42 below.

Table 42. PMU Counter Firmware Read Errors

| Error code            | Description  |
|-----------------------|--|
| SBI_SUCCESS           | firmware counter read successfully.  |
| SBI_ERR_INVALID_PARAM | <code>counter_idx</code> points to a hardware counter or an invalid counter. |

## 11.11. Function: Read a firmware counter high bits (FID #6)

```
struct sbiret sbi_pmu_counter_fw_read_hi(unsigned long counter_idx)
```

Provide the upper 32 bits of the current firmware counter value in `sbiret.value`. This function always returns zero in `sbiret.value` for RV64 (or higher) systems.

The possible error codes returned in `sbiret.error` are shown in [Table 43](#) below.

Table 43. PMU Counter Firmware Read High Errors

| Error code            | Description  |
|-----------------------|--|
| SBI_SUCCESS           | Firmware counter read successfully.  |
| SBI_ERR_INVALID_PARAM | <code>counter_idx</code> points to a hardware counter or an invalid counter. |

## 11.12. Function: Enable PMU snapshot feature (FID #7)

```
struct sbiret sbi_pmu_snapshot_set_shmem(unsigned long shmem_phys_lo,
                                         unsigned long shmem_phys_hi)
```

Set shared memory area for PMU state snapshot. The `shmem_phys_lo` specifies the lower XLEN bits and `shmem_phys_hi` specifies the upper XLEN bits of the shared memory physical address. The `shmem_phys_lo` MUST be 4096 bytes (i.e. page) aligned. The shared memory size must be 4096 bytes. The layout of the shared memory is described in [Table 44](#).

Table 44. SBI PMU Snapshot shared memory layout

| Name                    | Offset | Size | Description  |
|-------------------------|--------|------|--|
| counter_overflow_bitmap | 0x0000 | 8    | A bitmap of all logical overflowed counters. This is valid only if the <code>Sscofpmf</code> ISA extension is available. Otherwise, it must be zero. |
| counter_values          | 0x0008 | 512  | An array of 64-bit logical counters where each index represents the value of each logical counter associated with hardware/firmware.                 |
| Reserved                | 0x0208 | 3576 | Reserved for future use  |

Any future revisions to this structure should be made in a backward compatible manner and will be

associated with an SBI version.

This function should be invoked only once per hart at boot time. Once configured, the SBI implementation has read/write access to the shared memory when `sbi_pmu_counter_stop` is invoked with the `SBI_PMU_STOP_FLAG_TAKE_SNAPSHOT` flag set. The SBI implementation has read only access when `sbi_pmu_counter_start` is invoked with the `SBI_PMU_START_FLAG_INIT_SNAPSHOT` flag set. The SBI implementation must not access this memory any other time.

The possible error codes returned in `sbiret.error` are shown in [Table 45](#) below.

Table 45. PMU Setup Snapshot Area Errors

| Error code              | Description  |
|-------------------------|--|
| SBI_SUCCESS             | firmware counter read successfully.  |
| SBI_ERR_INVALID_ADDRESS | The shared memory pointed to by the <code>shmem_phys_lo</code> and <code>shmem_phys_hi</code> parameters is not writable or does not satisfy other requirements of <a href="#">Section 3.2</a> . |

## 11.13. Function Listing

Table 46. PMU Function List

| Function Name                                | SBI Version | FID | EID      |
|--|-------------|-----|----------|
| <code>sbi_pmu_num_counters</code>            | 0.3         | 0   | 0x504D55 |
| <code>sbi_pmu_counter_get_info</code>        | 0.3         | 1   | 0x504D55 |
| <code>sbi_pmu_counter_config_matching</code> | 0.3         | 2   | 0x504D55 |
| <code>sbi_pmu_counter_start</code>           | 0.3         | 3   | 0x504D55 |
| <code>sbi_pmu_counter_stop</code>            | 0.3         | 4   | 0x504D55 |
| <code>sbi_pmu_counter_fw_read</code>         | 0.3         | 5   | 0x504D55 |
| <code>sbi_pmu_counter_fw_read_hi</code>      | 2.0         | 6   | 0x504D55 |
| <code>sbi_pmu_snapshot_set_shmem</code>      | 2.0         | 7   | 0x504D55 |

## Chapter 12. Debug Console Extension (EID #0x4442434E "DBCN")

The debug console extension defines a generic mechanism for debugging and boot-time early prints from supervisor-mode software.

This extension replaces the legacy console putchar (EID #0x01) and console getchar (EID #0x02) extensions. The debug console extension allows supervisor-mode software to write or read multiple bytes in a single SBI call.

If the underlying physical console has extra bits for error checking (or correction) then these extra bits should be handled by the SBI implementation.



*It is recommended that bytes sent/received using the debug console extension follow UTF-8 character encoding.*

### 12.1. Function: Console Write (FID #0)

```
struct sbiret sbi_debug_console_write(unsigned long num_bytes,
                                     unsigned long base_addr_lo,
                                     unsigned long base_addr_hi)
```

Write bytes to the debug console from input memory.

The **num\_bytes** parameter specifies the number of bytes in the input memory. The physical base address of the input memory is represented by two XLEN bits wide parameters. The **base\_addr\_lo** parameter specifies the lower XLEN bits and the **base\_addr\_hi** parameter specifies the upper XLEN bits of the input memory physical base address.

This is a non-blocking SBI call and it may do partial/no writes if the debug console is not able to accept more bytes.

The number of bytes written is returned in **sbiret.value** and the possible error codes returned in **sbiret.error** are shown in [Table 47](#) below.

Table 47. Debug Console Write Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | Bytes written successfully.   |
| SBI_ERR_INVALID_PARAM | The memory pointed to by the <b>num_bytes</b> , <b>base_addr_lo</b> , and <b>base_addr_hi</b> parameters does not satisfy the requirements described in the <a href="#">Section 3.2</a> |
| SBI_ERR_FAILED        | Failed to write due to I/O errors.  |

## 12.2. Function: Console Read (FID #1)

```
struct sbiret sbi_debug_console_read(unsigned long num_bytes,
                                     unsigned long base_addr_lo,
                                     unsigned long base_addr_hi)
```

Read bytes from the debug console into an output memory.

The **num\_bytes** parameter specifies the maximum number of bytes which can be written into the output memory. The physical base address of the output memory is represented by two XLEN bits wide parameters. The **base\_addr\_lo** parameter specifies the lower XLEN bits and the **base\_addr\_hi** parameter specifies the upper XLEN bits of the output memory physical base address.

This is a non-blocking SBI call and it will not write anything into the output memory if there are no bytes to be read in the debug console.

The number of bytes read is returned in **sbiret.value** and the possible error codes returned in **sbiret.error** are shown in [Table 48](#) below.

Table 48. Debug Console Read Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | Bytes read successfully.  |
| SBI_ERR_INVALID_PARAM | The memory pointed to by the <b>num_bytes</b> , <b>base_addr_lo</b> , and <b>base_addr_hi</b> parameters does not satisfy the requirements described in the <a href="#">Section 3.2</a> |
| SBI_ERR_FAILED        | Failed to read due to I/O errors.   |

## 12.3. Function: Console Write Byte (FID #2)

```
struct sbiret sbi_debug_console_write_byte(uint8_t byte)
```

Write a single byte to the debug console.

This is a blocking SBI call and it will only return after writing the specified byte to the debug console. It will also return, with SBI\_ERR\_FAILED, if there are I/O errors.

The **sbiret.value** is set to zero and the possible error codes returned in **sbiret.error** are shown in [Table 49](#) below.

Table 49. Debug Console Write Byte Errors

| Error code     | Description                                 |
|----------------|---|
| SBI_SUCCESS    | Byte written successfully.                  |
| SBI_ERR_FAILED | Failed to write the byte due to I/O errors. |



## 12.4. Function Listing

Table 50. DBCN Function List

| Function Name                | SBI Version | FID | EID        |
|------------------------------|-------------|-----|------------|
| sbi_debug_console_write      | 2.0         | 0   | 0x4442434E |
| sbi_debug_console_read       | 2.0         | 1   | 0x4442434E |
| sbi_debug_console_write_byte | 2.0         | 2   | 0x4442434E |

## Chapter 13. System Suspend Extension (EID #0x53555350 "SUSP")

The system suspend extension defines a set of system-level sleep states and a function which allows the supervisor-mode software to request that the system transitions to a sleep state. Sleep states are identified with 32-bit wide identifiers (**sleep\_type**). The possible values for the identifiers are shown in [Table 51](#).

The term "system" refers to the world-view of supervisor software. The underlying SBI implementation may be provided by machine mode firmware or a hypervisor.

The system suspend extension does not provide any way for supported sleep types to be probed. Platforms are expected to specify their supported system sleep types and per-type wake up devices in their hardware descriptions. The **SUSPEND\_TO\_RAM** sleep type is the one exception, and its presence is implied by that of the extension.

Table 51. SUSP System Sleep Types

| Type                    | Name           | Description   |
|-------------------------|----------------|---|
| 0                       | SUSPEND_TO_RAM | This is a “suspend to RAM” sleep type, similar to ACPI’s S2 or S3. Entry requires all but the calling hart be in the HSM <b>STOPPED</b> state and all hart registers and CSRs saved to RAM. |
| 0x00000001 - 0x7fffffff |                | Reserved for future use   |
| 0x80000000 - 0xffffffff |                | Platform-specific system sleep types  |
| > 0xffffffff            |                | Reserved  |

### 13.1. Function: System Suspend (FID #0)

```
struct sbiret sbi_system_suspend(uint32_t sleep_type,
                                unsigned long resume_addr,
                                unsigned long opaque)
```

A return from a **sbi\_system\_suspend()** call implies an error and an error code from [Table 53](#) will be in **sbiret.error**. A successful suspend and wake up, results in the hart which initiated the suspend, resuming from the **STOPPED** state. To resume, the hart will jump to supervisor-mode, at the address specified by **resume\_addr**, with the specific register values described in [Table 52](#).

Table 52. SUSP System Resume Register State

| Register Name | Register Value |
|---------------|----------------|
| satp          | 0              |

| Register Name                                     | Register Value   |
|---|------------------|
| sstatus.SIE                                       | 0                |
| a0  | hartid           |
| a1  | opaque parameter |
| All other registers remain in an undefined state. |                  |



A single **unsigned long** parameter is sufficient for **resume\_addr**, because the hart will resume execution in supervisor-mode with the MMU off, hence **resume\_addr** must be less than XLEN bits wide.

The **resume\_addr** parameter points to a runtime-specified physical address, where the hart can resume execution in supervisor-mode after a system suspend.

The **opaque** parameter is an XLEN-bit value which will be set in the **a1** register when the hart resumes execution at **resume\_addr** after a system suspend.

Besides ensuring all entry criteria for the selected sleep type are met, such as ensuring other harts are in the **STOPPED** state, the caller must ensure all power units and domains are in a state compatible with the selected sleep type. The preparation of the power units, power domains, and wake-up devices used for resumption from the system sleep state is platform specific and beyond the scope of this specification.

When supervisor software is running inside a virtual machine, the SBI implementation is provided by a hypervisor. The system suspend will behave functionally the same as the native case, but might not result in any physical power changes.

The possible error codes returned in **sbiret.error** are shown in [Table 53](#).

Table 53. SUSP System Suspend Errors

| Error code              | Description  |
|-------------------------|--|
| SBI_SUCCESS             | System has suspended and resumed successfully.   |
| SBI_ERR_INVALID_PARAM   | <b>sleep_type</b> is reserved or is platform-specific and unimplemented.   |
| SBI_ERR_NOT_SUPPORTED   | <b>sleep_type</b> is not reserved and is implemented, but the platform does not support it due to one or more missing dependencies.  |
| SBI_ERR_INVALID_ADDRESS | <b>resume_addr</b> is not valid, possibly due to the following reasons:<br>* It is not a valid physical address.<br>* Executable access to the address is prohibited by a physical memory protection mechanism or H-extension G-stage for supervisor mode. |
| SBI_ERR_FAILED          | The suspend request failed for unspecified or unknown other reasons.   |

## 13.2. Function Listing

Table 54. SUSP Function List

| Function Name      | SBI Version | FID | EID        |
|--------------------|-------------|-----|------------|
| sbi_system_suspend | 2.0         | 0   | 0x53555350 |

# Chapter 14. CPPC Extension (EID #0x43505043 "CPPC")

ACPI defines the Collaborative Processor Performance Control (CPPC) mechanism, which is an abstract and flexible mechanism for the supervisor-mode power-management software to collaborate with an entity in the platform to manage the performance of the processors.

The SBI CPPC extension provides an abstraction to access the CPPC registers through SBI calls. The CPPC registers can be memory locations shared with a separate platform entity such as a BMC. Even though CPPC is defined in the ACPI specification, it may be possible to implement a CPPC driver based on Device Tree.

[Table 55](#) defines 32-bit identifiers for all CPPC registers to be used by the SBI CPPC functions. The first half of the 32-bit register space corresponds to the registers as defined by the ACPI specification. The second half provides the information not defined in the ACPI specification, but is additionally required by the supervisor-mode power-management software.

*Table 55. CPPC Registers*

| Register ID | Register                              | Bit Width | Attribute    | Description                |
|-------------|---------------------------------------|-----------|--------------|----------------------------|
| 0x00000000  | HighestPerformance                    | 32        | Read-only    | ACPI Spec 6.5: 8.4.6.1.1.1 |
| 0x00000001  | NominalPerformance                    | 32        | Read-only    | ACPI Spec 6.5: 8.4.6.1.1.2 |
| 0x00000002  | LowestNonlinearPerformance            | 32        | Read-only    | ACPI Spec 6.5: 8.4.6.1.1.4 |
| 0x00000003  | LowestPerformance                     | 32        | Read-only    | ACPI Spec 6.5: 8.4.6.1.1.5 |
| 0x00000004  | GuaranteedPerformanceRegister         | 32        | Read-only    | ACPI Spec 6.5: 8.4.6.1.1.6 |
| 0x00000005  | DesiredPerformanceRegister            | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.2.3 |
| 0x00000006  | MinimumPerformanceRegister            | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.2.2 |
| 0x00000007  | MaximumPerformanceRegister            | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.2.1 |
| 0x00000008  | PerformanceReductionToleranceRegister | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.2.4 |
| 0x00000009  | TimeWindowRegister                    | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.2.5 |
| 0x0000000A  | CounterWraparoundTime                 | 32 / 64   | Read-only    | ACPI Spec 6.5: 8.4.6.1.3.1 |
| 0x0000000B  | ReferencePerformanceCounterRegister   | 32 / 64   | Read-only    | ACPI Spec 6.5: 8.4.6.1.3.1 |

| Register ID             | Register                            | Bit Width | Attribute    | Description  |
|-------------------------|-------------------------------------|-----------|--------------|--|
| 0x0000000C              | DeliveredPerformanceCounterRegister | 32 / 64   | Read-only    | ACPI Spec 6.5: 8.4.6.1.3.1   |
| 0x0000000D              | PerformanceLimitedRegister          | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.3.2   |
| 0x0000000E              | CPPCEnableRegister                  | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.4   |
| 0x0000000F              | AutonomousSelectionEnable           | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.5   |
| 0x00000010              | AutonomousActivityWindowRegister    | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.6   |
| 0x00000011              | EnergyPerformancePreferenceRegister | 32        | Read / Write | ACPI Spec 6.5: 8.4.6.1.7   |
| 0x00000012              | ReferencePerformance                | 32        | Read-only    | ACPI Spec 6.5: 8.4.6.1.1.3   |
| 0x00000013              | LowestFrequency                     | 32        | Read-only    | ACPI Spec 6.5: 8.4.6.1.1.7   |
| 0x00000014              | NominalFrequency                    | 32        | Read-only    | ACPI Spec 6.5: 8.4.6.1.1.7   |
| 0x00000015 - 0x7FFFFFFF |                                     |           |              | Reserved for future use.   |
| 0x80000000              | TransitionLatency                   | 32        | Read-only    | Provides the maximum (worst-case) performance state transition latency in nanoseconds. |
| 0x80000001 - 0xFFFFFFFF |                                     |           |              | Reserved for future use.   |

## 14.1. Function: Probe CPPC register (FID #0)

```
struct sbiret sbi_cppc_probe(uint32_t cppc_reg_id)
```

Probe whether the CPPC register as specified by the `cppc_reg_id` parameter is implemented or not by the platform.

If the register is implemented, `sbiret.value` will contain the register width. If the register is not implemented, `sbiret.value` will be set to 0.

The possible error codes returned in `sbiret.error` are shown in [Table 56](#).

Table 56. CPPC Probe Errors

| Error code            | Description  |
|-----------------------|--|
| SBI_SUCCESS           | Probe completed successfully.                                      |
| SBI_ERR_INVALID_PARAM | <code>cppc_reg_id</code> is reserved.                              |
| SBI_ERR_FAILED        | The probe request failed for unspecified or unknown other reasons. |

## 14.2. Function: Read CPPC register (FID #1)

```
struct sbiret sbi_cppc_read(uint32_t cppc_reg_id)
```

Reads the register as specified in the `cppc_reg_id` parameter and returns the value in `sbiret.value`. When supervisor mode XLEN is 32, the `sbiret.value` will only contain the lower 32 bits of the CPPC register value.

The possible error codes returned in `sbiret.error` are shown in [Table 57](#).

Table 57. CPPC Read Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | Read completed successfully.                                      |
| SBI_ERR_INVALID_PARAM | <code>cppc_reg_id</code> is reserved.                             |
| SBI_ERR_NOT_SUPPORTED | <code>cppc_reg_id</code> is not implemented by the platform.      |
| SBI_ERR_DENIED        | <code>cppc_reg_id</code> is a write-only register.                |
| SBI_ERR_FAILED        | The read request failed for unspecified or unknown other reasons. |

## 14.3. Function: Read CPPC register high bits (FID #2)

```
struct sbiret sbi_cppc_read_hi(uint32_t cppc_reg_id)
```

Reads the upper 32-bit value of the register specified in the `cppc_reg_id` parameter and returns the value in `sbiret.value`. This function always returns zero in `sbiret.value` when supervisor mode XLEN is 64 or higher.

The possible error codes returned in `sbiret.error` are shown in [Table 58](#).

Table 58. CPPC Read Hi Errors

| Error code            | Description  |
|-----------------------|--|
| SBI_SUCCESS           | Read completed successfully.                                 |
| SBI_ERR_INVALID_PARAM | <code>cppc_reg_id</code> is reserved.                        |
| SBI_ERR_NOT_SUPPORTED | <code>cppc_reg_id</code> is not implemented by the platform. |
| SBI_ERR_DENIED        | <code>cppc_reg_id</code> is a write-only register.           |

| Error code     | Description   |
|----------------|---|
| SBI_ERR_FAILED | The read request failed for unspecified or unknown other reasons. |

## 14.4. Function: Write to CPPC register (FID #3)

```
struct sbiret sbi_cppc_write(uint32_t cppc_reg_id, uint64_t val)
```

Writes the value passed in the **val** parameter to the register as specified in the **cppc\_reg\_id** parameter.

The possible error codes returned in **sbiret.error** are shown in [Table 59](#).

Table 59. CPPC Write Errors

| Error code            | Description  |
|-----------------------|--|
| SBI_SUCCESS           | Write completed successfully.                                      |
| SBI_ERR_INVALID_PARAM | <b>cppc_reg_id</b> is reserved.                                    |
| SBI_ERR_NOT_SUPPORTED | <b>cppc_reg_id</b> is not implemented by the platform.             |
| SBI_ERR_DENIED        | <b>cppc_reg_id</b> is a read-only register.                        |
| SBI_ERR_FAILED        | The write request failed for unspecified or unknown other reasons. |

## 14.5. Function Listing

Table 60. CPPC Function List

| Function Name    | SBI Version | FID | EID        |
|------------------|-------------|-----|------------|
| sbi_cppc_probe   | 2.0         | 0   | 0x43505043 |
| sbi_cppc_read    | 2.0         | 1   | 0x43505043 |
| sbi_cppc_read_hi | 2.0         | 2   | 0x43505043 |
| sbi_cppc_write   | 2.0         | 3   | 0x43505043 |



# Chapter 15. Nested Acceleration Extension (EID #0x4E41434C "NACL")

Nested virtualization is the ability of a hypervisor to run another hypervisor as a guest. RISC-V nested virtualization requires an LO hypervisor (running in hypervisor-mode) to trap-and-emulate the RISC-V H-extension [priv\_v1.12] functionality (such as CSR accesses, HFENCE instructions, HLV/HSV instructions, etc.) for the L1 hypervisor (running in virtualized supervisor-mode).

The SBI nested acceleration extension defines a shared memory based interface between the SBI implementation (or LO hypervisor) and the supervisor software (or L1 hypervisor) which allows both to collaboratively reduce traps taken by the LO hypervisor for emulating RISC-V H-extension functionality. The nested acceleration shared memory allows the L1 hypervisor to batch multiple RISC-V H-extension CSR accesses and HFENCE requests which are then emulated by the LO hypervisor upon an explicit synchronization SBI call.



*The M-mode firmware should not implement the SBI nested acceleration extension if the underlying platform has the RISC-V H-extension implemented in hardware.*

This SBI extension defines optional features which MUST be discovered by the supervisor software (or L1 hypervisor) before using the corresponding SBI functions. Each nested acceleration feature is assigned a unique ID which is an unsigned 32-bit integer. The Table 61 below provides a list of all nested acceleration features.

Table 61. Nested acceleration features

| Feature ID   | Feature Name               | Description             |
|--------------|----------------------------|-------------------------|
| 0x00000000   | SBI_NACL_FEAT_SYNC_CSR     | Synchronize CSR         |
| 0x00000001   | SBI_NACL_FEAT_SYNC_HFENCE  | Synchronize HFENCE      |
| 0x00000002   | SBI_NACL_FEAT_SYNC_SRET    | Synchronize SRET        |
| 0x00000003   | SBI_NACL_FEAT_AUTOSWAP_CSR | Autoswap CSR            |
| > 0x00000003 | RESERVED                   | Reserved for future use |

To use the SBI nested acceleration extension, the supervisor software (or L1 hypervisor) MUST set up a nested acceleration shared memory physical address for each virtual hart at boot-time. The physical base address of the nested acceleration shared memory MUST be 4096 bytes (i.e. page) aligned and the size of the nested acceleration shared memory is assumed to be  $4096 + (1024 * (XLEN / 8))$  bytes. The Table 62 below shows the layout of nested acceleration shared memory.

Table 62. Nested acceleration shared memory layout

| Name          | Offset     | Size (bytes) | Description                                |
|---------------|------------|--------------|--|
| Scratch space | 0x00000000 | 4096         | Nested acceleration feature specific data. |

| Name      | Offset     | Size (bytes) | Description   |
|-----------|------------|--------------|---|
| CSR space | 0x00001000 | $XLEN * 128$ | An array of 1024 XLEN-bit words where each word corresponds to a possible RISC-V H-extension CSR defined in the Table 2.1 of the RISC-V privileged specification <a href="#">[priv_v1.12]</a> . |

The contents of the scratch space shown in the [Table 62](#) above is defined separately for each nested acceleration feature.

The contents of the CSR space shown in the [Table 62](#) above is an array of RISC-V H-extension CSR values where CSR  $\langle x \rangle$  is at index  $\langle i \rangle = ((\langle x \rangle \& 0xc00) \gg 2) | (\langle x \rangle \& 0xff)$ . The SBI implementation (or LO hypervisor) MUST update the CSR space whenever the state of any RISC-V H-extension CSR changes unless some nested acceleration feature defines a different behaviour. The [Table 63](#) below shows CSR space index ranges for all possible 1024 RISC-V H-extension CSRs.

Table 63. Nested acceleration H-extension CSR index ranges

| H-extension CSR address |       |       |               | SBI NACL CSR space index |
|-------------------------|-------|-------|---------------|--------------------------|
| [11:10]                 | [9:8] | [7:4] | Hex Range     | Hex Range                |
| 00                      | 10    | xxxx  | 0x200 - 0x2ff | 0x000 - 0x0ff            |
| 01                      | 10    | 0xxx  | 0x600 - 0x67f | 0x100 - 0x17f            |
| 01                      | 10    | 10xx  | 0x680 - 0x6bf | 0x180 - 0x1bf            |
| 01                      | 10    | 11xx  | 0x6c0 - 0x6ff | 0x1c0 - 0x1ff            |
| 10                      | 10    | 0xxx  | 0xa00 - 0xa7f | 0x200 - 0x27f            |
| 10                      | 10    | 10xx  | 0xa80 - 0xabf | 0x280 - 0x2bf            |
| 10                      | 10    | 11xx  | 0xac0 - 0xaff | 0x2c0 - 0x2ff            |
| 11                      | 10    | 0xxx  | 0xe00 - 0xe7f | 0x300 - 0x37f            |
| 11                      | 10    | 10xx  | 0xe80 - 0xebf | 0x380 - 0x3bf            |
| 11                      | 10    | 11xx  | 0xec0 - 0xeff | 0x3c0 - 0x3ff            |

## 15.1. Feature: Synchronize CSR (ID #0)

The synchronize CSR feature describes the ability of the SBI implementation (or LO hypervisor) to allow supervisor software (or L1 hypervisor) to write RISC-V H-extension CSRs using the CSR space.

This nested acceleration feature defines the scratch space offset range **0x0F80 - 0x0FFF** (128 bytes) as nested CSR dirty bitmap. The nested CSR dirty bitmap contains 1-bit for each possible RISC-V H-extension CSR.

To write a CSR  $\langle x \rangle$  in nested acceleration shared memory, the supervisor software (or L1 hypervisor) MUST do the following:

1. Compute  $\langle i \rangle = ((\langle x \rangle \& 0xc00) \gg 2) | (\langle x \rangle \& 0xff)$

2. Write a new CSR value at word with index *<i>* in the CSR space
3. Set the *<i>* bit in the nested CSR dirty bitmap

To synchronize a CSR *<x>*, the SBI implementation (or LO hypervisor) MUST do the following:

1. Compute  $\text{<i> = ((<x> \& 0xc00) \gg 2) | (<x> \& 0xff)}$
2. If bit *<i>* is not set in the nested CSR dirty bitmap then goto step 5
3. Emulate write to CSR *<x>* with the new CSR value taken from the word with index *<i>* in the CSR space
4. Clear the *<i>* bit in the nested CSR dirty bitmap
5. Write back the latest CSR value of CSR *<x>* to the word with index *<i>* in the CSR space

When synchronizing multiple CSRs, if the value of a CSR *<y>* depends on the value of some other CSR *<x>* then the SBI implementation (or LO hypervisor) MUST synchronize CSR *<x>* before CSR *<y>*. For example, the value of CSR *hip* depends on the value of the CSR *hvip*, which means *hvip* is emulated and written first, followed by *hip*.

## 15.2. Feature: Synchronize HFENCE (ID #1)

The synchronize HFENCE feature describes the ability of the SBI implementation (or LO hypervisor) to allow supervisor software (or L1 hypervisor) to issue HFENCE using the scratch space.

This nested acceleration feature defines the scratch space offset range *0x0800 - 0x0F7F* (1920 bytes) as an array of nested HFENCE entries. The total number of nested HFENCE entries are *3840 / XLEN* where each nested HFENCE entry consists of four XLEN-bit words.

A nested HFENCE entry is equivalent to an HFENCE over a range of guest addresses. The [Table 64](#) below shows the nested HFENCE entry format whereas [Table 65](#) below provides a list of nested HFENCE entry types. Upon an explicit synchronize HFENCE request from supervisor software (or L1 hypervisor), the SBI implementation (or LO hypervisor) will process nested HFENCE entries with the *Config.Pending* bit set. After processing pending nested HFENCE entries, the SBI implementation (or LO hypervisor) will clear the *Config.Pending* bit of these entries.

*Table 64. Nested HFENCE entry format*

| Word | Name        | Encoding  |
|------|-------------|---|
| 0    | Config      | <p>Config information about the nested HFENCE entry</p> <p>BIT[XLEN-1:XLEN-1] - Pending<br/> BIT[XLEN-2:XLEN-4] - Reserved and must be zero<br/> BIT[XLEN-5:XLEN-8] - Type<br/> BIT[XLEN-9:XLEN-9] - Reserved and must be zero<br/> BIT[XLEN-10:XLEN-16] - Order<br/> if XLEN == 32 then<br/>     BIT[15:9] - VMID<br/>     BIT[8:0] - ASID<br/> else<br/>     BIT[29:16] - VMID<br/>     BIT[15:0] - ASID</p> <p>The page size for invalidation is assumed to be<br/> <math>1 \ll (\text{Config.Order} + 12)</math> bytes.</p> |
| 1    | Page_Number | Page address right shifted by $\text{Config.Order} + 12$  |
| 2    | Reserved    | Reserved for future use and must be zero  |
| 3    | Page_Count  | Number of pages to invalidate   |

Table 65. Nested HFENCE entry types

| Type | Name          | Description   |
|------|---------------|---|
| 0    | GVMA          | Invalidate a guest physical address range across all VMIDs. The VMID and ASID fields of the Config word are ignored and MUST be zero.   |
| 1    | GVMA_ALL      | Invalidate all guest physical addresses across all VMIDs. The Order, VMID and ASID fields of the Config word are ignored and MUST be zero. The Page_Number and Page_Count words are ignored and MUST be zero. |
| 2    | GVMA_VMID     | Invalidate a guest physical address range for a particular VMID. The ASID field of the Config word is ignored and MUST be zero.   |
| 3    | GVMA_VMID_ALL | Invalidate all guest physical addresses for a particular VMID. The Order and ASID fields of the Config word are ignored and MUST be zero. The Page_Number and Page_Count words are ignored and MUST be zero.  |
| 4    | VVMA          | Invalidate a guest virtual address range for a particular VMID. The ASID field of the Config word is ignored and MUST be zero.  |
| 5    | VVMA_ALL      | Invalidate all guest virtual addresses for a particular VMID. The Order and ASID fields of the Config word are ignored and MUST be zero. The Page_Number and Page_Count words are ignored and MUST be zero.   |

| Type | Name          | Description   |
|------|---------------|---|
| 6    | VVMA_ASID     | Invalidate a guest virtual address range for a particular VMID and ASID.  |
| 7    | VVMA_ASID_ALL | Invalidate all guest virtual addresses for a particular VMID and ASID. The <b>Order</b> field of the <b>Config</b> word is ignored and MUST be zero. The <b>Page_Number</b> and <b>Page_Count</b> words are ignored and MUST be zero. |
| > 7  | Reserved      | Reserved for future use.  |

To add a nested HFENCE entry, the supervisor software (or L1 hypervisor) MUST do the following:

1. Find an unused nested HFENCE entry with **Config.Pending == 0**
2. Update the **Page\_Number** and **Page\_Count** words in the nested HFENCE entry
3. Update the **Config** word in the nested HFENCE entry such that **Config.Pending** bit is set

To synchronize a nested HFENCE entry, the SBI implementation (or LO hypervisor) MUST do the following:

1. If **Config.Pending == 0** then do nothing and skip below steps
2. Process HFENCE based on details in the nested HFENCE entry
3. Clear the **Config.Pending** bit in the nested HFENCE entry

## 15.3. Feature: Synchronize SRET (ID #2)

The synchronize SRET feature describes the ability of the SBI implementation (or LO hypervisor) to do synchronization of CSRs and HFENCES in the nested acceleration shared memory for the supervisor software (or L1 hypervisor) along with SRET emulation.

This nested acceleration feature defines the scratch space offset range **0x0000 - 0x01FF** (512 bytes) as nested SRET context. The [Table 66](#) below shows contents of the nested SRET context.

Table 66. Nested SRET context

| Offset                | Name     | Encoding                                 |
|-----------------------|----------|--|
| <b>0 * (XLEN / 8)</b> | Reserved | Reserved for future use and must be zero |
| <b>1 * (XLEN / 8)</b> | X1       | Value to be restored in GPR X1           |
| <b>2 * (XLEN / 8)</b> | X2       | Value to be restored in GPR X2           |
| <b>3 * (XLEN / 8)</b> | X3       | Value to be restored in GPR X3           |
| <b>4 * (XLEN / 8)</b> | X4       | Value to be restored in GPR X4           |
| <b>5 * (XLEN / 8)</b> | X5       | Value to be restored in GPR X5           |
| <b>6 * (XLEN / 8)</b> | X6       | Value to be restored in GPR X6           |
| <b>7 * (XLEN / 8)</b> | X7       | Value to be restored in GPR X7           |
| <b>8 * (XLEN / 8)</b> | X8       | Value to be restored in GPR X8           |
| <b>9 * (XLEN / 8)</b> | X9       | Value to be restored in GPR X9           |

| Offset                  | Name     | Encoding                        |
|-------------------------|----------|---------------------------------|
| 10 * (XLEN / 8)         | X10      | Value to be restored in GPR X10 |
| 11 * (XLEN / 8)         | X11      | Value to be restored in GPR X11 |
| 12 * (XLEN / 8)         | X12      | Value to be restored in GPR X12 |
| 13 * (XLEN / 8)         | X13      | Value to be restored in GPR X13 |
| 14 * (XLEN / 8)         | X14      | Value to be restored in GPR X14 |
| 15 * (XLEN / 8)         | X15      | Value to be restored in GPR X15 |
| 16 * (XLEN / 8)         | X16      | Value to be restored in GPR X16 |
| 17 * (XLEN / 8)         | X17      | Value to be restored in GPR X17 |
| 18 * (XLEN / 8)         | X18      | Value to be restored in GPR X18 |
| 19 * (XLEN / 8)         | X19      | Value to be restored in GPR X19 |
| 20 * (XLEN / 8)         | X20      | Value to be restored in GPR X20 |
| 21 * (XLEN / 8)         | X21      | Value to be restored in GPR X21 |
| 22 * (XLEN / 8)         | X22      | Value to be restored in GPR X22 |
| 23 * (XLEN / 8)         | X23      | Value to be restored in GPR X23 |
| 24 * (XLEN / 8)         | X24      | Value to be restored in GPR X24 |
| 25 * (XLEN / 8)         | X25      | Value to be restored in GPR X25 |
| 26 * (XLEN / 8)         | X26      | Value to be restored in GPR X26 |
| 27 * (XLEN / 8)         | X27      | Value to be restored in GPR X27 |
| 28 * (XLEN / 8)         | X28      | Value to be restored in GPR X28 |
| 29 * (XLEN / 8)         | X29      | Value to be restored in GPR X29 |
| 30 * (XLEN / 8)         | X30      | Value to be restored in GPR X30 |
| 31 * (XLEN / 8)         | X31      | Value to be restored in GPR X31 |
| 32 * (XLEN / 8) - 0x1FF | Reserved | Reserved for future use         |

Before sending a synchronize SRET request to the SBI implementation (or LO hypervisor), the supervisor software (or L1 hypervisor) MUST write the GPR  $X\langle i \rangle$  values to be restored at offset  $\langle i \rangle * (XLEN / 8)$  of the nested SRET context.

Upon a synchronize SRET request from the supervisor software (or L1 hypervisor), the SBI implementation (or LO hypervisor) MUST do the following:

1. If SBI\_NACL\_FEAT\_SYNC\_CSR feature is available then
  - a. All RISC-V H-extension CSRs implemented by the SBI implementation (or LO hypervisor) are synchronized as described in the [Section 15.1](#). This is equivalent to the SBI call `sbi_nacl_sync_csr(-1UL)`.
2. If SBI\_NACL\_FEAT\_SYNC\_HFENCE feature is available then
  - a. All nested HFENCE entries are synchronized as described in the [Section 15.2](#). This is equivalent to the SBI call `sbi_nacl_sync_hfence(-1UL)`.

3. Restore GPR  $X_{<i>i</i>}$  registers from the nested SRET context.
4. Emulate the SRET instruction as defined by the RISC-V Privilege specification [\[priv\\_v1.12\]](#).

## 15.4. Feature: Autoswap CSR (ID #3)

The autoswap CSR feature describes the ability of the SBI implementation (or LO hypervisor) to automatically swap certain RISC-V H-extension CSR values from the nested acceleration shared memory in the following situations:

- Before emulating the SRET instruction for a synchronized SRET request from the supervisor software (or L1 hypervisor).
- After supervisor (or L1) virtualization state changes from ON to OFF.



*The supervisor software (or L1 hypervisor) should use the autoswap CSR feature in conjunction with the synchronize SRET feature.*

This nested acceleration feature defines the scratch space offset range  $0x0200 - 0x027F$  (128 bytes) as nested autoswap context. The [Table 67](#) below shows contents of the nested autoswap context.

Table 67. Nested autoswap context

| Offset                  | Name           | Encoding   |
|-------------------------|----------------|--|
| $0 * (XLEN / 8)$        | Autoswap_Flags | Autoswap flags<br><br>BIT[XLEN-1:1] - Reserved for future use and must be zero<br>BIT[0:0] - HSTATUS |
| $1 * (XLEN / 8)$        | HSTATUS        | Value to be swapped with HSTATUS CSR   |
| $2 * (XLEN / 8) - 0x7F$ | Reserved       | Reserved for future use.   |

To enable automatic swapping of CSRs from the nested autoswap context, the supervisor software (or L1 hypervisor) MUST do the following:

1. Write the HSTATUS swap value in the nested autoswap context.
2. Set Autoswap\_Flags.HSTATUS bit in the nested autoswap context.

To swap CSRs from the nested autoswap context, the SBI implementation (or LO hypervisor) MUST do the following:

1. If Autoswap\_Flags.HSTATUS bit is set in the nested autoswap context then swap the supervisor HSTATUS CSR value with the HSTATUS value in the nested autoswap context.

## 15.5. Function: Probe nested acceleration feature (FID #0)

```
struct sbiret sbi_nacl_probe_feature(uint32_t feature_id)
```



Probe a nested acceleration feature. This is a mandatory function of the SBI nested acceleration extension. The `feature_id` parameter specifies the nested acceleration feature to probe. [Table 61](#) provides a list of possible feature IDs.

This function always returns `SBI_SUCCESS` in `sbiret.error`. It returns `0` in `sbiret.value` if the given `feature_id` is not available, or `1` in `sbiret.value` if it is available.

## 15.6. Function: Set nested acceleration shared memory (FID #1)

```
struct sbiret sbi_nacl_set_shmem(unsigned long shmem_phys_lo,
                                unsigned long shmem_phys_hi,
                                unsigned long flags)
```

Set and enable the shared memory for nested acceleration on the calling hart. This is a mandatory function of the SBI nested acceleration extension.

If both `shmem_phys_lo` and `shmem_phys_hi` parameters are not all-ones bitwise then `shmem_phys_lo` specifies the lower XLEN bits and `shmem_phys_hi` specifies the upper XLEN bits of the shared memory physical base address. `shmem_phys_lo` MUST be 4096 bytes (i.e. page) aligned and the size of the shared memory is assumed to be  $4096 + (\text{XLEN} * 128)$  bytes.

If both `shmem_phys_lo` and `shmem_phys_hi` parameters are all-ones bitwise then the nested acceleration features are disabled.

The `flags` parameter is reserved for future use and must be zero.

The possible error codes returned in `sbiret.error` are shown in [Table 68](#).

Table 68. NACL Set Shared Memory Errors

| Error code              | Description   |
|-------------------------|---|
| SBI_SUCCESS             | Shared memory was set or cleared successfully.  |
| SBI_ERR_INVALID_PARAM   | The <code>flags</code> parameter is not zero or or the <code>shmem_phys_lo</code> parameter is not 4096 bytes aligned.  |
| SBI_ERR_INVALID_ADDRESS | The shared memory pointed to by the <code>shmem_phys_lo</code> and <code>shmem_phys_hi</code> parameters does not satisfy the requirements described in <a href="#">Section 3.2</a> . |

## 15.7. Function: Synchronize shared memory CSRs (FID #2)

```
struct sbiret sbi_nacl_sync_csr(unsigned long csr_num)
```

Synchronize CSRs in the nested acceleration shared memory. This is an optional function which is



only available if the SBI\_NACL\_FEAT\_SYNC\_CSR feature is available. The parameter `csr_num` specifies the set of RISC-V H-extension CSRs to be synchronized.

If `csr_num` is all-ones bitwise then all RISC-V H-extension CSRs implemented by the SBI implementation (or LO hypervisor) are synchronized as described in the [Section 15.1](#).

If  $(\text{csr\_num} \& 0x300) == 0x200$  and  $\text{csr\_num} < 0x1000$  then only a single RISC-V H-extension CSR specified by the `csr_num` parameter is synchronized as described in the [Section 15.1](#).

The possible error codes returned in `sbiret.error` are shown in [Table 69](#).

Table 69. NACL Synchronize CSR Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | CSRs synchronized successfully.   |
| SBI_ERR_NOT_SUPPORTED | SBI_NACL_FEAT_SYNC_CSR feature is not available.  |
| SBI_ERR_INVALID_PARAM | <code>csr_num</code> is not all-ones bitwise and either:<br>* $(\text{csr\_num} \& 0x300) \neq 0x200$ or<br>* $\text{csr\_num} \geq 0x1000$ or<br>* <code>csr_num</code> is not implemented by the SBI implementation |
| SBI_ERR_NO_SHMEM      | Nested acceleration shared memory not available.  |

## 15.8. Function: Synchronize shared memory HFENCEs (FID #3)

```
struct sbiret sbi_nacl_sync_hfence(unsigned long entry_index)
```

Synchronize HFENCEs in the nested acceleration shared memory. This is an optional function which is only available if the SBI\_NACL\_FEAT\_SYNC\_HFENCE feature is available. The parameter `entry_index` specifies the set of nested HFENCE entries to be synchronized.

If `entry_index` is all-ones bitwise then all nested HFENCE entries are synchronized as described in the [Section 15.2](#).

If  $\text{entry\_index} < (3840 / \text{XLEN})$  then only a single nested HFENCE entry specified by the `entry_index` parameter is synchronized as described in the [Section 15.2](#).

The possible error codes returned in `sbiret.error` are shown in [Table 70](#).

Table 70. NACL Synchronize HFENCE Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | HFENCEs synchronized successfully.  |
| SBI_ERR_NOT_SUPPORTED | SBI_NACL_FEAT_SYNC_HFENCE feature is not available.   |
| SBI_ERR_INVALID_PARAM | <code>entry_index</code> is not all-ones bitwise and<br>$\text{entry\_index} \geq (3840 / \text{XLEN})$ . |
| SBI_ERR_NO_SHMEM      | Nested acceleration shared memory not available.  |

# 15.9. Function: Synchronize shared memory and emulate SRET (FID #4)

```
struct sbiret sbi_nacl_sync_sret(void)
```

Synchronize CSRs and HFENCES in the nested acceleration shared memory and emulate the SRET instruction. This is an optional function which is only available if the SBI\_NACL\_FEAT\_SYNC\_SRET feature is available.

This function is used by supervisor software (or L1 hypervisor) to do a synchronize SRET request and the SBI implementation (or LO hypervisor) MUST handle it as described in the [Section 15.3](#).

This function does not return upon success and the possible error codes returned in `sbiret.error` upon failure are shown in [Table 71](#).

Table 71. NACL Synchronize SRET Errors

| Error code            | Description                                       |
|-----------------------|---|
| SBI_ERR_NOT_SUPPORTED | SBI_NACL_FEAT_SYNC_SRET feature is not available. |
| SBI_ERR_NO_SHMEM      | Nested acceleration shared memory not available.  |

# 15.10. Function Listing

Table 72. NACL Function List

| Function Name          | SBI Version | FID | EID        |
|------------------------|-------------|-----|------------|
| sbi_nacl_probe_feature | 2.0         | 0   | 0x4E41434C |
| sbi_nacl_set_shmem     | 2.0         | 1   | 0x4E41434C |
| sbi_nacl_sync_csr      | 2.0         | 2   | 0x4E41434C |
| sbi_nacl_sync_hfence   | 2.0         | 3   | 0x4E41434C |
| sbi_nacl_sync_sret     | 2.0         | 4   | 0x4E41434C |

## Chapter 16. Steal-time Accounting Extension (EID #0x535441 "STA")

SBI implementations may encounter situations where virtual HARTs are ready to run, but must be withheld from running. These situations may be, for example, when multiple SBI domains share processors or when an SBI implementation is a hypervisor and guest contexts share processors with other guest contexts or host tasks. When virtual HARTs are at times withheld from running, observers within the contexts of the virtual HARTs may need a way to account for less progress than would otherwise be expected. The time a virtual HART was ready, but had to wait, is called "stolen time" and the tracking of it is referred to as steal-time accounting. The Steal-time Accounting (STA) extension defines the mechanism in which an SBI implementation provides steal-time and preemption information, for each virtual HART, to supervisor-mode software.

### 16.1. Function: Set Steal-time Shared Memory Address (FID #0)

```
struct sbiret sbi_steal_time_set_shmem(unsigned long shmem_phys_lo,
                                       unsigned long shmem_phys_hi,
                                       uint32_t flags)
```

Set the shared memory physical base address for steal-time accounting of the calling virtual HART and enable the SBI implementation's steal-time information reporting.

If `shmem_phys_lo` and `shmem_phys_hi` are not all-ones bitwise, then `shmem_phys_lo` specifies the lower XLEN bits and `shmem_phys_hi` specifies the upper XLEN bits of the shared memory physical base address. `shmem_phys_lo` MUST be 64-byte aligned. The size of the shared memory is assumed to be at least 64 bytes. All bytes MUST be set to zero by the SBI implementation before returning from the SBI call.

If `shmem_phys_lo` and `shmem_phys_hi` are all-ones bitwise, the SBI implementation will stop reporting steal-time information for the virtual HART.

`flags` MUST be set to zero.

It is not expected for the shared memory to be written by the supervisor-mode software while it is in use for steal-time accounting. However, the SBI implementation MUST not misbehave if a write from supervisor-mode software occurs, however, in that case, it MAY leave the shared memory filled with inconsistent data.

The SBI implementation MUST stop writing to the shared memory upon system reset.

The shared memory layout is defined in [Table 73](#)

*Table 73. STA Shared Memory Structure*

| Name      | Offset | Size | Description   |
|-----------|--------|------|---|
| sequence  | 0      | 4    | <p>The SBI implementation MUST increment this field to an odd value before writing the <b>steal</b> field, and increment it again to an even value after writing <b>steal</b> (i.e. an odd sequence number indicates an in-progress update). The SBI implementation SHOULD ensure that the sequence field remains odd for only very short periods of time.</p> <p>The supervisor-mode software MUST check this field before and after reading the <b>steal</b> field, and repeat the read if it is different or odd.</p> <p><i>This sequence field enables the value of the steal field to be read by supervisor-mode software executing in a 32-bit environment.</i></p> |
| flags     | 4      | 4    | <p>Always zero.</p> <p><i>Future extensions of the SBI call might allow the supervisor-mode software to write to some of the fields of the shared memory. Such extensions will not be enabled as long as a zero value is used for the flags argument to the SBI call.</i></p>   |
| steal     | 8      | 8    | <p>The amount of time in which this virtual HART was not idle and scheduled out, in nanoseconds. The time during which the virtual HART is idle will not be reported as steal-time.</p>   |
| preempted | 16     | 1    | <p>An advisory flag indicating whether the virtual HART which registered this structure is running or not. A non-zero value MAY be written by the SBI implementation if the virtual HART has been preempted (i.e. while the <b>steal</b> field is increasing), while a zero value MUST be written before the virtual HART starts to run again.</p> <p><i>This preempted field can, for example, be used by the supervisor-mode software to check if a lock holder has been preempted, and, in that case, disable optimistic spinning.</i></p>   |
| pad       | 17     | 47   | Pad with zeros to a 64 byte boundary.   |

**sbiret.value** is set to zero and the possible error codes returned in **sbiret.error** are shown in Table 74 below.

Table 74. STA Set Steal-time Shared Memory Address Errors

| Error code              | Description  |
|-------------------------|--|
| SBI_SUCCESS             | The steal-time shared memory physical base address was set or cleared successfully.  |
| SBI_ERR_INVALID_PARAM   | The <code>flags</code> parameter is not zero or the <code>shmem_phys_lo</code> is not 64-byte aligned.   |
| SBI_ERR_INVALID_ADDRESS | The shared memory pointed to by the <code>shmem_phys_lo</code> and <code>shmem_phys_hi</code> parameters is not writable or does not satisfy other requirements of <a href="#">Section 3.2</a> . |
| SBI_ERR_FAILED          | The request failed for unspecified or unknown other reasons.   |

## 16.2. Function Listing

Table 75. STA Function List

| Function Name                         | SBI Version | FID | EID      |
|---------------------------------------|-------------|-----|----------|
| <code>sbi_steal_time_set_shmem</code> | 2.0         | 0   | 0x535441 |

# Chapter 17. SBI Supervisor Software Events Extension (EID #0x535345 "SSE")

The SBI Supervisor Software Events (SSE) extension provides a mechanism to inject software events from an SBI implementation to supervisor software such that it preempts all other traps and interrupts.

The software events can be of two types: local or global. A local software event is local to a HART and can be handled only on that HART whereas a global software event is a system event. A software event can only be handled by a HART which in the HSM **STARTED** state.

## 17.1. Software Event Identification

Each software event is identified by a unique 32-bit unsigned integer called `event_id` which is encoded as shown in [Table 76](#) below.

Table 76. SSE Event Identification

| Software Event ID       | Description                    |
|-------------------------|--------------------------------|
| 0x00000000              | Local RAS event                |
| 0x00000001              | Local PMU event                |
| 0x00000002              | Local async page fault event   |
| 0x00000003 - 0x3ffffff  | Reserved for future use        |
| 0x40000000 - 0x7ffffffe | Local platform specific event  |
| 0x7ffffff               | Local debug event              |
| 0x80000000              | Global RAS event               |
| 0x80000001 - 0xbffffff  | Reserved for future use        |
| 0xc0000000 - 0xffffffe  | Global platform specific event |
| 0xfffff                 | Global debug event             |

## 17.2. Software Event States

At any point in time, a software event can be in one of the following states:

1. **UNUSED** - Software event is not used by supervisor software
2. **REGISTERED** - Supervisor software has provided an event handler for the software event but it is not ready to handle the events.
3. **ENABLED** - Supervisor software is ready to handle the software event.
4. **RUNNING** - Supervisor software has taken the software event and is busy handling it.

The below diagram shows the state transitions of a software event.

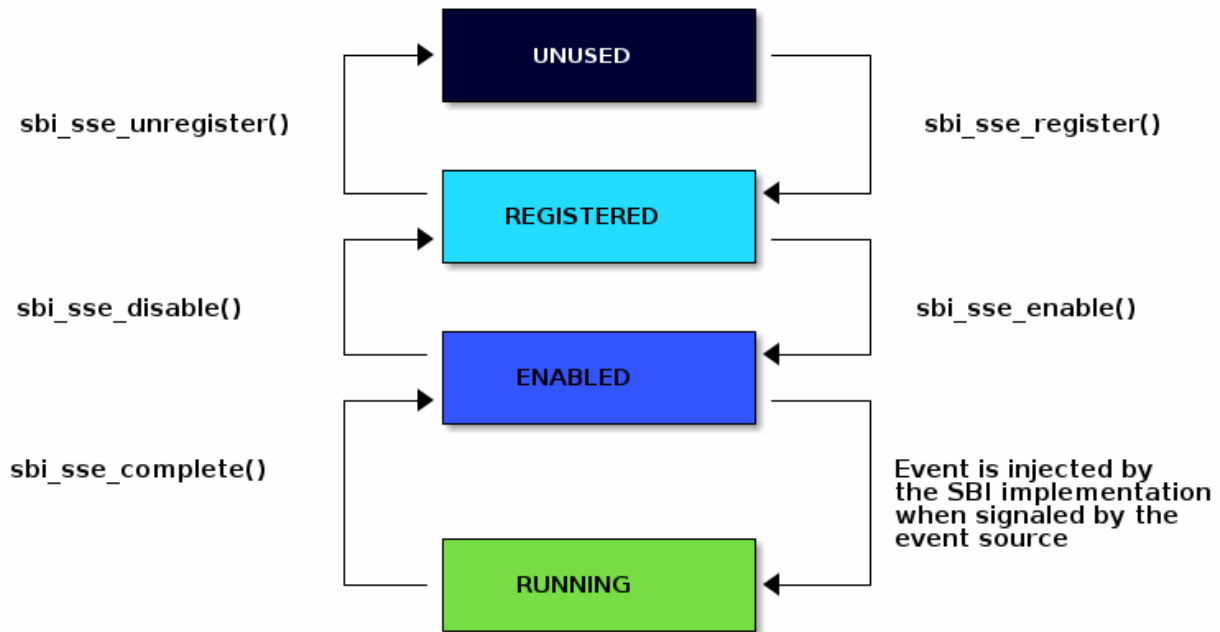


Figure 4. SBI SSE State Machine

Each software event is associated with a target HART (referred to as `target_hart`) which will handle the software event.

A **global** software event **MUST** be registered and enabled only by the `target_hart` of that software event. The default `target_hart` of a global software event is assigned by the SBI implementation which can be overridden by the supervisor software.

A **local** software event **MUST** be registered and enabled by all HARTs which want to handle the software event. The state of a local software event **MUST** be tracked separately for each HART. The `target_hart` for a local software event is fixed to the associated HART.

## 17.3. Software Event Priority

Each software event has an associated event priority (also referred to as `event_priority`) which can be used by the SBI implementation to select a software event for injection when multiple software events are pending on the same HART.

The priority of a software event is a 32-bit unsigned integer where lower value means higher priority. By default, all software events have event priority as zero.

If two or more events have same priority on a given hart then SBI implementation must use `event_id` to break the tie where the lower `event_id` has higher priority.

A higher priority event always preempts a lower priority event on the same HART. Once the higher priority event is marked as completed, the previous handler will be resumed.



*If a software event in the **RUNNING** state is signaled again by the event source, the `pending_status` attribute of the event is set. The event is injected again after the current event completes, provided that supervisor software doesn't disable the event.*

## 17.4. Software Event Attributes

A software event can have various attributes associated with it. A software event attribute is a unique 32-bit unsigned integer called `attr_id`. An attribute can have Read-Only or Read-Write access permissions. The supervisor software can query these event attributes and change the attributes that have Read-Write access permissions. The Table 77 below provides a list event attributes.

Table 77. SSE Event Attributes

| Attribute ID ( <code>attr_id</code> ) | Read-Only  | Description  | Possible values  |
|---------------------------------------|--|--|--|
| 0x00000000                            | Yes  | Software event state   | 0: UNUSED<br>1: REGISTERED<br>2: ENABLED<br>3: RUNNING   |
| 0x00000001                            | No   | Software event priority  | 32-bit unsigned integer  |
| 0x00000002                            | Yes  | Event injection by the supervisor software using <code>sbi_sse_inject</code> call. | 0: Not allowed<br>1: Allowed   |
| 0x00000003                            | No (for global events)<br>Yes (for local events) | HART id of the <code>target_hart</code> .  | unsigned long integer  |
| 0x00000004                            | Yes  | Pending Status   | This is set when the event source signals the event. When the event is injected, it is cleared.<br>0: Not Pending 1: Pending |
| > 0x00000004                          | ---  | Reserved for future use  | ---  |

## 17.5. Software Event Handler

To handle a software event, the supervisor software MUST register an event handler and enable it. Each event handler registered by the supervisor software consists of a handler context (also referred to as `handler_context`).

The `handler_context` contains the following register states:

1. **Entry State** - contains the supervisor register state when handling the software event injected by SBI implementation. It is referred to as `entry_state`. This register state must be initialized by the supervisor software before registering the handler. The handler's entry point is at offset 0 of the `entry_state` as mentioned in Table 78
2. **Interrupted State** - contains interrupted register state and is referred to as `interrupted_state`. The interrupted execution mode is saved at the end of the `interrupted_state`.

The `handler_context` must be contiguous in both virtual and physical address space. The physical



address of the `handler_context` is represented by `handler_context_phys_hi` and `handler_context_phys_lo`.



*It is advisable to use different context for different events. Since a higher priority event can preempt lower priority events, if same context is used, then the interrupted state will be overwritten with register values of the higher priority event. This will make resuming to the previous handler impossible.*

Table 78. SSE Register offsets in entry state

| Register Offsets in Entry State            | Data   |
|--|--|
| <code>entry_state + 0 * (XLEN / 8)</code>  | Entry program counter<br>Must be 2-byte aligned virtual address. |
| <code>entry_state + 1 * (XLEN / 8)</code>  | X1   |
| <code>entry_state + 2 * (XLEN / 8)</code>  | X2   |
| <code>entry_state + 3 * (XLEN / 8)</code>  | X3   |
| <code>entry_state + 4 * (XLEN / 8)</code>  | X4   |
| <code>entry_state + 5 * (XLEN / 8)</code>  | X5   |
| <code>entry_state + 6 * (XLEN / 8)</code>  | X6   |
| <code>entry_state + 7 * (XLEN / 8)</code>  | X7   |
| <code>entry_state + 8 * (XLEN / 8)</code>  | X8   |
| <code>entry_state + 9 * (XLEN / 8)</code>  | X9   |
| <code>entry_state + 10 * (XLEN / 8)</code> | X10  |
| <code>entry_state + 11 * (XLEN / 8)</code> | X11  |
| <code>entry_state + 12 * (XLEN / 8)</code> | X12  |
| <code>entry_state + 13 * (XLEN / 8)</code> | X13  |
| <code>entry_state + 14 * (XLEN / 8)</code> | X14  |
| <code>entry_state + 15 * (XLEN / 8)</code> | X15  |
| <code>entry_state + 16 * (XLEN / 8)</code> | X16  |
| <code>entry_state + 17 * (XLEN / 8)</code> | X17  |
| <code>entry_state + 18 * (XLEN / 8)</code> | X18  |
| <code>entry_state + 19 * (XLEN / 8)</code> | X19  |
| <code>entry_state + 20 * (XLEN / 8)</code> | X20  |
| <code>entry_state + 21 * (XLEN / 8)</code> | X21  |
| <code>entry_state + 22 * (XLEN / 8)</code> | X22  |
| <code>entry_state + 23 * (XLEN / 8)</code> | X23  |
| <code>entry_state + 24 * (XLEN / 8)</code> | X24  |
| <code>entry_state + 25 * (XLEN / 8)</code> | X25  |

| Register Offsets in Entry State            | Data |
|--|------|
| <code>entry_state + 26 * (XLEN / 8)</code> | X26  |
| <code>entry_state + 27 * (XLEN / 8)</code> | X27  |
| <code>entry_state + 28 * (XLEN / 8)</code> | X28  |
| <code>entry_state + 29 * (XLEN / 8)</code> | X29  |
| <code>entry_state + 30 * (XLEN / 8)</code> | X30  |
| <code>entry_state + 31 * (XLEN / 8)</code> | X31  |

Table 79. SSE Register offsets in interrupted state

| Register Offsets in Interrupted State            | Data                        |
|--|-----------------------------|
| <code>interrupted_state + 0 * (XLEN / 8)</code>  | Interrupted program counter |
| <code>interrupted_state + 1 * (XLEN / 8)</code>  | Saved copy of X1            |
| <code>interrupted_state + 2 * (XLEN / 8)</code>  | Saved copy of X2            |
| <code>interrupted_state + 3 * (XLEN / 8)</code>  | Saved copy of X3            |
| <code>interrupted_state + 4 * (XLEN / 8)</code>  | Saved copy of X4            |
| <code>interrupted_state + 5 * (XLEN / 8)</code>  | Saved copy of X5            |
| <code>interrupted_state + 6 * (XLEN / 8)</code>  | Saved copy of X6            |
| <code>interrupted_state + 7 * (XLEN / 8)</code>  | Saved copy of X7            |
| <code>interrupted_state + 8 * (XLEN / 8)</code>  | Saved copy of X8            |
| <code>interrupted_state + 9 * (XLEN / 8)</code>  | Saved copy of X9            |
| <code>interrupted_state + 10 * (XLEN / 8)</code> | Saved copy of X10           |
| <code>interrupted_state + 11 * (XLEN / 8)</code> | Saved copy of X11           |
| <code>interrupted_state + 12 * (XLEN / 8)</code> | Saved copy of X12           |
| <code>interrupted_state + 13 * (XLEN / 8)</code> | Saved copy of X13           |
| <code>interrupted_state + 14 * (XLEN / 8)</code> | Saved copy of X14           |
| <code>interrupted_state + 15 * (XLEN / 8)</code> | Saved copy of X15           |
| <code>interrupted_state + 16 * (XLEN / 8)</code> | Saved copy of X16           |
| <code>interrupted_state + 17 * (XLEN / 8)</code> | Saved copy of X17           |
| <code>interrupted_state + 18 * (XLEN / 8)</code> | Saved copy of X18           |
| <code>interrupted_state + 19 * (XLEN / 8)</code> | Saved copy of X19           |
| <code>interrupted_state + 20 * (XLEN / 8)</code> | Saved copy of X20           |
| <code>interrupted_state + 21 * (XLEN / 8)</code> | Saved copy of X21           |
| <code>interrupted_state + 22 * (XLEN / 8)</code> | Saved copy of X22           |
| <code>interrupted_state + 23 * (XLEN / 8)</code> | Saved copy of X23           |
| <code>interrupted_state + 24 * (XLEN / 8)</code> | Saved copy of X24           |

| Register Offsets in Interrupted State            | Data   |
|--|--|
| <code>interrupted_state + 25 * (XLEN / 8)</code> | Saved copy of X25  |
| <code>interrupted_state + 26 * (XLEN / 8)</code> | Saved copy of X26  |
| <code>interrupted_state + 27 * (XLEN / 8)</code> | Saved copy of X27  |
| <code>interrupted_state + 28 * (XLEN / 8)</code> | Saved copy of X28  |
| <code>interrupted_state + 29 * (XLEN / 8)</code> | Saved copy of X29  |
| <code>interrupted_state + 30 * (XLEN / 8)</code> | Saved copy of X30  |
| <code>interrupted_state + 31 * (XLEN / 8)</code> | Saved copy of X31  |
| <code>interrupted_state + 32 * (XLEN / 8)</code> | <p>Interrupted Execution mode</p> <p><b>bit [0]</b> = Privilege mode which was interrupted<br/>(1 = S-mode, 0 = U-mode)</p> <p><b>bit [1]</b> = Virtualization state which was interrupted<br/>(1 = ON, 0 = OFF)</p> <p><b>bit [2]</b> = Saved copy of sstatus.SPIE</p> <p><b>bit [XLEN-1:3]</b> = Reserved for future use</p> |

## 17.6. Software Event Injection

To inject a software event on a HART, the SBI implementation must do the following:

1. Copy X1 to X31 registers into the `interrupted_state` in `handler_context` from the offsets mentioned in [Table 79](#).
2. Load X1 to X31 registers from `entry_state` in `handler_context` from the offsets mentioned in [Table 78](#).
3. Save the interrupted mode at offset `interrupted_state + 32 * (XLEN / 8)` in `handler_context` as shown in [Table 79](#).
4. Update registers as follows:
  - a. Set `sstatus.SPIE` = `sstatus.SIE`
  - b. Set `sstatus.SIE` = 0
5. Resume execution with:
  - a. Program counter = value at `entry_state + 0 * (XLEN / 8)`
  - b. Privilege mode = S-mode
  - c. Virtualization state = OFF

## 17.7. Software Event Completion

After handling the software event on a HART, the supervisor software must notify the SBI implementation about completion of event handling using `sbi_sse_complete` call. The SBI

implementation must do the following to complete event handling and resume interrupted state:

1. Restore X1 to X31 registers from the `interrupted_state` of `handler_context` from the offsets mentioned in [Table 79](#).
2. Update supervisor CSRs as follows:
  - a. Set `sstatus.SIE` = `sstatus.SPIE`
  - b. Set `sstatus.SPIE` = bit[2] of the value at `interrupted_state` + 32 \* (`XLEN` / 8)
3. Resume execution with:
  - Virtualization state = bit[1] of the value at `interrupted_state` + 32 \* (`XLEN` / 8)
  - Privilege mode = bit[0] of the value at `interrupted_state` + 32 \* (`XLEN` / 8)
  - Program counter = value at `interrupted_state` + 0 \* (`XLEN` / 8)

If the supervisor software wishes to resume from a different location, it can update the `interrupted_state` fields accordingly.

## 17.8. Function: Get a software event attribute (FID #0)

```
struct sbiret sbi_sse_get_attr(uint32_t event_id,
                             uint32_t attr_id)
```

Get an event attribute value of a software event. The `event_id` parameter specifies the software event whereas `attr_id` parameter specifies the software event's attribute.

Upon success, the event attribute value is returned in `sbiret.value`. In case of an error, the possible error codes are shown in the [Table 80](#) below:

Table 80. SSE Event Attribute Read Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | Attribute of given event returned successfully.                   |
| SBI_ERR_INVALID_PARAM | <code>event_id</code> , <code>attr_id</code> or both are invalid. |

## 17.9. Function: Set a software event attribute (FID #1)

```
struct sbiret sbi_sse_set_attr(uint32_t event_id,
                              uint32_t attr_id,
                              unsigned long value)
```

Set an event attribute value of a software event. The `event_id` parameter specifies the software event whereas `attr_id` parameter specifies the event attribute. The new event attribute value is specified by

the `value` parameter.

If the software event specified by the parameter `event_id` is a global software event and is in the `RUNNING` state, the `target_hart` attribute for the software event cannot be set. In such case, supervisor software can retry.



*If the `target_hart` of a global event changes then it doesn't affect the state of the global event.*

An error is returned in `sbiret.error`. In case of error, The possible return values are listed in [Table 81](#) below:

Table 81. SSE Event Attribute Write Errors

| Error code                         | Description   |
|------------------------------------|---|
| <code>SBI_SUCCESS</code>           | Attribute value set successfully.   |
| <code>SBI_ERR_INVALID_PARAM</code> | <code>event_id</code> , <code>attr_id</code> or both are invalid                            |
| <code>SBI_ERR_BAD_RANGE</code>     | <code>value</code> does not match the possible values defined in <a href="#">Table 77</a> . |
| <code>SBI_ERR_BUSY</code>          | Failed to set attribute value. Supervisor software can retry.                               |
| <code>SBI_ERR_DENIED</code>        | The attribute is read-only.   |

## 17.10. Function: Register a software event (FID #2)

```
struct sbiret sbi_sse_register(uint32_t event_id,
                              unsigned long handler_context_phys_hi,
                              unsigned long handler_context_phys_lo)
```

Register a handler for the software event. The `event_id` parameter specifies the event ID for which the handler is being registered. The parameters `handler_context_phys_hi` and `handler_context_phys_lo` contain the upper and the lower XLEN bits, respectively, of the handler's context. The `handler_context_phys_lo` parameter must be  $(XLEN / 8)$  byte aligned.

If the software event specified by the parameter `event_id` is a global event, the registration of the event handler can only be done by the `target_hart` of the software event.

On successful registration, the event state moves from `UNUSED` to `REGISTERED`. In case of an error, possible error codes are listed in [Table 82](#) below:

Table 82. SSE Event Register Errors

| Error code                         | Description   |
|------------------------------------|---|
| <code>SBI_SUCCESS</code>           | Event handler is registered successfully.   |
| <code>SBI_ERR_INVALID_STATE</code> | The event is not in <code>UNUSED</code> state.  |
| <code>SBI_ERR_INVALID_PARAM</code> | <code>event_id</code> is invalid or other parameters not satisfy requirements defined in <a href="#">Section 17.5</a> . |

| Error code              | Description   |
|-------------------------|---|
| SBI_ERR_INVALID_ADDRESS | The memory pointed by <code>handler_context_phys_lo</code> , <code>handler_context_phys_hi</code> , parameters does not satisfy the requirements described in <a href="#">Section 3.2</a> or the <code>handler_context_phys_lo</code> parameter is not $(XLEN / 8)$ byte aligned. |
| SBI_ERR_DENIED          | <code>event_id</code> is a global event and the calling HART id is not the <code>target_hart</code> for the software event.   |

## 17.11. Function: Unregister a software event (FID #3)

```
struct sbiret sbi_sse_unregister(uint32_t event_id)
```

Unregister the handler for the given `event_id`. The event MUST be in the `REGISTERED` state before it can be unregistered.

If the software event specified by the parameter `event_id` is a global event, the event handler can only be unregistered by the `target_hart` of the software event.

On successful unregistration, the event is moved to the `UNREGISTERED` state. In case of an error, possible error codes are listed in [Table 83](#) below:

Table 83. SSE Event Unregister Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | Event handler is unregistered successfully.   |
| SBI_ERR_INVALID_STATE | Event is not in <code>REGISTERED</code> state.  |
| SBI_ERR_INVALID_PARAM | <code>event_id</code> is invalid.   |
| SBI_ERR_DENIED        | <code>event_id</code> is a global event and the calling HART id is not the <code>target_hart</code> for the software event. |

## 17.12. Function: Enable a software event (FID #4)

```
struct sbiret sbi_sse_enable(uint32_t event_id)
```

Enable the software event specified by the `event_id` parameter. For local events, the event is enabled only for the calling HART. For global events, the event can only be enabled by the `target_hart` of the event.

The event MUST be in the `REGISTERED` state otherwise this function will fail.

On success, the event is moved to the `ENABLED` state and the SBI implementation can inject the event when it occurs. In case of an error, possible error codes are listed in [Table 84](#) below:

Table 84. SSE Event Enable Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | Event is successfully enabled.  |
| SBI_ERR_INVALID_PARAM | <code>event_id</code> is not valid.   |
| SBI_ERR_INVALID_STATE | The event is not in the <code>REGISTERED</code> state.  |
| SBI_ERR_DENIED        | <code>event_id</code> is a global event and the calling HART id is not the <code>target_hart</code> for the software event. |

## 17.13. Function: Disable a software event (FID #5)

```
struct sbiret sbi_sse_disable(uint32_t event_id)
```

Disable the software event specified by the `event_id` parameter. For local events, the event is disabled only for the calling HART. For global events, the event can only be disabled by the `target_hart` of the event. The event must be in the `ENABLED` state.

On success, the event is moved to the `REGISTERED` state. In case of an error, possible error codes are listed in [Table 85](#).

Table 85. SSE Event Disable Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | Event is successfully disabled.   |
| SBI_ERR_INVALID_PARAM | <code>event_id</code> is not valid.   |
| SBI_ERR_INVALID_STATE | Event is not in the <code>ENABLED</code> state.   |
| SBI_ERR_DENIED        | <code>event_id</code> is a global event and the calling HART id is not the <code>target_hart</code> for the software event. |

## 17.14. Function: Complete software event handling (FID #6)

```
struct sbiret sbi_sse_complete(uint32_t event_id,
                               uint32_t status,
                               uint32_t flags)
```

Complete the supervisor event handling for the event. The event must be in the `RUNNING` state.

If supervisor software could not handle the event, it must set the `status` parameter to `SBI_SSE_HANDLER_FAILED`. On success, it must set the `status` parameter to `SBI_SSE_HANDLER_SUCCESS`. Other possible status codes are listed in [Table 86](#).

The `flags` parameter represents additional information from supervisor to the SBI implementation and the [Table 87](#) lists the bit-encoding for it.

Table 86. SSE Event Complete Status Values

| Value                           | Enum                    | Description                         |
|---------------------------------|-------------------------|-------------------------------------|
| 0x00000000                      | SBI_SSE_HANDLER_SUCCESS | Supervisor successfully handled the |
| event.                          | 0x00000001              | SBI_SSE_HANDLER_FAILED              |
| Supervisor failed to handle the | event.                  | > 0x00000001                        |

Table 87. SSE Event Complete Flags Values

| Flag Name             | Bits       | Description             |
|-----------------------|------------|-------------------------|
| SBI_SSE_EVENT_DISABLE | 0:0        | Disable the event.      |
| RESERVED              | 1:(XLEN-1) | All non-zero values are |

In case of an error, possible error codes are listed in [Table 88](#).

Table 88. SSE Event Complete Errors

| Error code  | Description   |
|---|---|
| SBI_SUCCESS   | Event is successfully marked completed.                               |
| SBI_ERR_INVALID_PARAM                                       | <b>event_id</b> , <b>status</b> , <b>flags</b> or all of them have an |
| invalid value.  | SBI_ERR_INVALID_STATE   |
| The <b>event_id</b> software event is not in RUNNING state. | SBI_ERR_DENIED  |

## 17.15. Function: Signal a software event (FID #7)

```
struct sbiret sbi_sse_inject(uint32_t event_id,
                             unsigned long hart_id)
```

The supervisor software can inject a software event with the help of this function. The **event\_id** paramater refers to the event to be injected.

For global events, the **hart\_id** parameter is ignored and it is injected on the **target\_hart** for the event. For local events, the **hart\_id** parameter refers to the HART on which the event is to be injected. An event can only be injected if it is allowed by the event attribute as described in [Table 77](#).

In case of an error, possible error codes are listed in [Table 89](#).

Table 89. SSE Event Inject Errors

| Error code            | Description   |
|-----------------------|---|
| SBI_SUCCESS           | Event is successfully injected or <b>pending_status</b> is set on the given hart. |
| SBI_ERR_INVALID_PARAM | <b>event_id</b> or <b>hart_id</b> or both are invalid.                            |



---

| Error code     | Description  |
|----------------|--|
| SBI_ERR_DENIED | The event cannot be injected by the supervisor software. |

# Chapter 18. Experimental SBI Extension Space (EIDs #0x08000000 - #0x08FFFFFF)

No management.

# Chapter 19. Vendor Specific Extension Space (EIDs #0x09000000 - #0x09FFFFFF)

Low bits from `mvendorid`.

## Chapter 20. Firmware Specific Extension Space (EIDs #0x0A000000 - #0x0AFFFFFF)

Low bits is SBI implementation ID. The firmware specific SBI extensions are for SBI implementations. It provides firmware specific SBI functions which are defined in the external firmware specification.

# References

- [priv\_v1.12] The RISC-V Instruction Set Manual, Volume II: Privileged Architecture, Document Version 20211203, URL: [github.com/riscv/riscv-isa-manual/releases/tag/Priv-v1.12](https://github.com/riscv/riscv-isa-manual/releases/tag/Priv-v1.12)