



RISC-V Supervisor Binary Interface Specification

RISC-V Platform Runtime Services Task Group

Version -, 2023-05-31: Draft

Table of Contents

Preamble	1
Copyright and license information	2
Change Log	3
Version 2.0-rc1	3
Version 1.0.0	3
Version 1.0-rc3	3
Version 1.0-rc2	3
Version 1.0-rc1	4
Version 0.3.0	4
Version 0.3-rc1	4
Version 0.2	4
1. Introduction	5
2. Terms and Abbreviations	7
3. Binary Encoding	8
3.1. HART list parameter	9
3.2. Shared memory physical address range parameter	9
4. Base Extension (EID #0x10)	10
4.1. Function: Get SBI specification version (FID #0)	10
4.2. Function: Get SBI implementation ID (FID #1)	10
4.3. Function: Get SBI implementation version (FID #2)	10
4.4. Function: Probe SBI extension (FID #3)	10
4.5. Function: Get machine vendor ID (FID #4)	10
4.6. Function: Get machine architecture ID (FID #5)	11
4.7. Function: Get machine implementation ID (FID #6)	11
4.8. Function Listing	11
4.9. SBI Implementation IDs	11
5. Legacy Extensions (EIDs #0x00 - #0x0F)	13
5.1. Extension: Set Timer (EID #0x00)	13
5.2. Extension: Console Putchar (EID #0x01)	13
5.3. Extension: Console Getchar (EID #0x02)	14
5.4. Extension: Clear IPI (EID #0x03)	14
5.5. Extension: Send IPI (EID #0x04)	14
5.6. Extension: Remote FENCE.I (EID #0x05)	14
5.7. Extension: Remote SFENCE.VMA (EID #0x06)	15
5.8. Extension: Remote SFENCE.VMA with ASID (EID #0x07)	15
5.9. Extension: System Shutdown (EID #0x08)	15
5.10. Function Listing	15
6. Timer Extension (EID #0x54494D45 "TIME")	17
6.1. Function: Set Timer (FID #0)	17
6.2. Function Listing	17

7. IPI Extension (EID #0x735049 "sPI: s-mode IPI")	18
7.1. Function: Send IPI (FID #0)	18
7.2. Function Listing	18
8. RFENCE Extension (EID #0x52464E43 "RFNC")	19
8.1. Function: Remote FENCE.I (FID #0)	19
8.2. Function: Remote SFENCE.VMA (FID #1)	19
8.3. Function: Remote SFENCE.VMA with ASID (FID #2)	20
8.4. Function: Remote HFENCE.GVMA with VMID (FID #3)	20
8.5. Function: Remote HFENCE.GVMA (FID #4)	21
8.6. Function: Remote HFENCE.VVMA with ASID (FID #5)	21
8.7. Function: Remote HFENCE.VVMA (FID #6)	22
8.8. Function Listing	22
9. Hart State Management Extension (EID #0x48534D "HSM")	23
9.1. Function: HART start (FID #0)	24
9.2. Function: HART stop (FID #1)	26
9.3. Function: HART get status (FID #2)	26
9.4. Function: HART suspend (FID #3)	26
9.5. Function Listing	28
10. System Reset Extension (EID #0x53525354 "SRST")	29
10.1. Function: System reset (FID #0)	29
10.2. Function Listing	30
11. Performance Monitoring Unit Extension (EID #0x504D55 "PMU")	31
11.1. Event: Hardware general events (Type #0)	31
11.2. Event: Hardware cache events (Type #1)	32
11.3. Event: Hardware raw events (Type #2)	33
11.4. Event: Firmware events (Type #15)	34
11.5. Function: Get number of counters (FID #0)	35
11.6. Function: Get details of a counter (FID #1)	35
11.7. Function: Find and configure a matching counter (FID #2)	36
11.8. Function: Start a set of counters (FID #3)	37
11.9. Function: Stop a set of counters (FID #4)	38
11.10. Function: Read a firmware counter (FID #5)	39
11.11. Function: Read a firmware counter high bits (FID #6)	39
11.12. Function Listing	39
12. Debug Console Extension (EID #0x4442434E "DBCN")	41
12.1. Function: Console Write (FID #0)	41
12.2. Function: Console Read (FID #1)	42
12.3. Function: Console Write Byte (FID #2)	42
12.4. Function Listing	43
13. System Suspend Extension (EID #0x53555350 "SUSP")	44
13.1. Function: System Suspend (FID #0)	44
13.2. Function Listing	45
14. CPPC Extension (EID #0x43505043 "CPPC")	47

14.1. Function: Probe CPPC register (FID #0)	48
14.2. Function: Read CPPC register (FID #1)	49
14.3. Function: Read CPPC register high bits (FID #2)	49
14.4. Function: Write to CPPC register (FID #3)	50
14.5. Function Listing	50
15. Nested Acceleration Extension (EID #0x4E41434C "NACL")	51
15.1. Feature: Synchronize CSR (ID #0)	52
15.2. Feature: Synchronize HFENCE (ID #1)	53
15.3. Feature: Synchronize SRET (ID #2)	55
15.4. Feature: Autoswap CSR (ID #3)	57
15.5. Function: Probe nested acceleration feature (FID #0)	57
15.6. Function: Set nested acceleration shared memory (FID #1)	58
15.7. Function: Synchronize shared memory CSRs (FID #2)	58
15.8. Function: Synchronize shared memory HFENCES (FID #3)	59
15.9. Function: Synchronize shared memory and emulate SRET (FID #4)	60
15.10. Function Listing	60
16. Steal-time Accounting Extension (EID #0x535441 "STA")	61
16.1. Function: Set Steal-time Shared Memory Address (FID #0)	61
16.2. Function Listing	63
17. SBI Supervisor Software Events Extension (EID #0x535345 "SSE")	64
17.1. Software Event Identification	64
17.2. Software Event States	64
17.3. Software Event Priority	65
17.4. Software Event Attributes	66
17.5. Software Event Handler	66
17.6. Software Event Injection	69
17.7. Software Event Completion	70
17.8. Function: Get a software event attribute (FID #0)	70
17.9. Function: Set a software event attribute (FID #1)	70
17.10. Function: Register a software event (FID #2)	71
17.11. Function: Unregister a software event (FID #3)	72
17.12. Function: Enable a software event (FID #4)	72
17.13. Function: Disable a software event (FID #5)	72
17.14. Function: Complete software event handling (FID #6)	73
17.15. Function: Signal a software event (FID #7)	74
18. Experimental SBI Extension Space (EIDs #0x08000000 - #0x08FFFFFF)	75
19. Vendor Specific Extension Space (EIDs #0x09000000 - #0x09FFFFFF)	76
20. Firmware Specific Extension Space (EIDs #0x0A000000 - #0x0AFFFFFF)	77
References	78

Preamble



This document is in the [Development state](#)

Assume everything can change but backward compatibility with perviously ratified specification will be maintained.

Copyright and license information

This RISC-V SBI specification is © 2022 RISC-V International.

It is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at creativecommons.org/licenses/by/4.0/.

Change Log

Version 2.0-rc1

- Added common description for shared memory physical address range parameter
- Added SBI debug console extension
- Relaxed the counter width requirement on SBI PMU firmware counters
- Added `sbi_pmu_counter_fw_read_hi()` in SBI PMU extension
- Reserved space for SBI implementation specific firmware events
- Added SBI system suspend extension
- Added SBI CPPC extension
- Clarified that an SBI extension can be partially implemented only if it defines a mechanism to discover implemented SBI functions
- Added error code `SBI_ERR_NO_SHMEM`
- Added SBI nested acceleration extension
- Added common description for a virtual HART
- Added SBI steal-time accounting extension
- Added error code to refer to invalid state of a state machine.
- Added error code to refer to a value which is out of its valid range.
- Added SBI supervisor software event extension

Version 1.0.0

- Updated the version for ratification

Version 1.0-rc3

- Updated the calling convention
- Fixed a typo in PMU extension
- Added a abbreviation table

Version 1.0-rc2

- Update to RISC-V formatting
- Improved the introduction
- Removed all references to RV32

Version 1.0-rc1

- A typo fix

Version 0.3.0

- Few typo fixes
- Updated the LICENSE with detailed text instead of a hyperlink

Version 0.3-rc1

- Improved document styling and naming conventions
- Added SBI system reset extension
- Improved SBI introduction section
- Improved documentation of SBI hart state management extension
- Added suspend function to SBI hart state management extension
- Added performance monitoring unit extension
- Clarified that an SBI extension shall not be partially implemented

Version 0.2

- The entire v0.1 SBI has been moved to the legacy extension, which is now an optional extension. This is technically a backwards-incompatible change because the legacy extension is optional and v0.1 of the SBI doesn't allow probing, but it's as good as we can do.

Chapter 1. Introduction

This specification describes the RISC-V Supervisor Binary Interface, known from here on as SBI. The SBI allows supervisor-mode (S-mode or VS-mode) software to be portable across all RISC-V implementations by defining an abstraction for platform (or hypervisor) specific functionality. The design of the SBI follows the general RISC-V philosophy of having a small core along with a set of optional modular extensions.

An SBI extension defines a set of SBI functions which provides a particular functionality to supervisor-mode software. SBI extensions as a whole are optional and cannot be partially implemented unless an SBI extension defines a mechanism to discover implemented SBI functions. If `sbi_probe_extension()` signals that an extension is available, all functions present in the SBI version reported by `sbi_get_spec_version()` must conform to that version of the SBI specification.

The higher privilege software providing SBI interface to the supervisor-mode software is referred as an SBI implementation or Supervisor Execution Environment (SEE). An SBI implementation (or SEE) can be platform runtime firmware executing in machine-mode (M-mode) (see below [Figure 1](#)) or it can be some hypervisor executing in hypervisor-mode (HS-mode) (see below [Figure 2](#)).




Figure 1. RISC-V System without H-extension




Figure 2. RISC-V System with H-extension

HARTs are provisioned by the SBI implementation for supervisor-mode software. Hence, from the perspective of the SBI implementation, the S-mode HART contexts are referred to as virtual HARTs. In the case that the implementation is a hypervisor, virtual HARTs represent the VS-mode guest contexts.

The SBI specification doesn't specify any method for hardware discovery. The supervisor software must rely on the other industry standard hardware discovery methods (i.e. Device Tree or ACPI) for that.

Chapter 2. Terms and Abbreviations

This specification uses the following terms and abbreviations:

Term	Meaning
ACPI	Advanced Configuration and Power Interface
ASID	Address Space Identifier
BMC	Baseboard Management Controller
CPPC	Collaborative Processor Performance Control
EID	Extension ID
FID	Function ID
HSM	Hart State Management
IPI	Inter Processor Interrupt
PMU	Performance Monitoring Unit
SBI	Supervisor Binary Interface
SEE	Supervisor Execution Environment
VMID	Virtual Machine Identifier

Chapter 3. Binary Encoding

All SBI functions share a single binary encoding, which facilitates the mixing of SBI extensions. The SBI specification follows the below calling convention.

- An **ECALL** is used as the control transfer instruction between the supervisor and the SEE.
- **a7** encodes the SBI extension ID (**EID**),
- **a6** encodes the SBI function ID (**FID**) for a given extension ID encoded in **a7** for any SBI extension defined in or after SBI v0.2.
- All registers except **a0** & **a1** must be preserved across an SBI call by the callee.
- SBI functions must return a pair of values in **a0** and **a1**, with **a0** returning an error code. This is analogous to returning the C structure

```
struct sbiret {
    long error;
    long value;
};
```

In the name of compatibility, SBI extension IDs (**EIDs**) and SBI function IDs (**FIDs**) are encoded as signed 32-bit integers. When passed in registers these follow the standard above calling convention rules.

The [Table 1](#) below provides a list of Standard SBI error codes.

Table 1. Standard SBI Errors

Error Type	Value	Description
SBI_SUCCESS	0	Completed successfully
SBI_ERR_FAILED	-1	Failed
SBI_ERR_NOT_SUPPORTED	-2	Not supported
SBI_ERR_INVALID_PARAM	-3	Invalid parameter(s)
SBI_ERR_DENIED	-4	Denied or not allowed
SBI_ERR_INVALID_ADDRESS	-5	Invalid address(s)
SBI_ERR_ALREADY_AVAILABLE	-6	Already available
SBI_ERR_ALREADY_STARTED	-7	Already started
SBI_ERR_ALREADY_STOPPED	-8	Already stopped
SBI_ERR_NO_SHMEM	-9	Shared memory not available
SBI_ERR_INVALID_STATE	-10	SBI_ERR_BAD_RANGE

An **ECALL** with an unsupported SBI extension ID (**EID**) or an unsupported SBI function ID (**FID**) must return the error code **SBI_ERR_NOT_SUPPORTED**.

Every SBI function should prefer **unsigned long** as the data type. It keeps the specification simple and easily adaptable for all RISC-V ISA types. In case the data is defined as 32bit wide, higher privilege

software must ensure that it only uses 32 bit data only.

3.1. HART list parameter

If an SBI function needs to pass a list of harts to the higher privilege mode, it must use a hart mask as defined below. This is applicable to any extensions defined in or after v0.2.

Any function, requiring a hart mask, need to pass following two arguments.

- `unsigned long hart_mask` is a scalar bit-vector containing hartids
- `unsigned long hart_mask_base` is the starting hartid from which bit-vector must be computed.

In a single SBI function call, maximum number harts that can be set is always XLEN. If a lower privilege mode needs to pass information about more than XLEN harts, it should invoke multiple instances of the SBI function call. `hart_mask_base` can be set to `-1` to indicate that `hart_mask` can be ignored and all available harts must be considered.

Any function using hart mask may return error values listed in the [Table 2](#) below which are in addition to function specific error values.

Table 2. HART Mask Errors

Error code	Description
SBI_ERR_INVALID_PARAMETER	Either <code>hart_mask_base</code> , or at least one hartid from <code>hart_mask</code> , is not valid, i.e. either the hartid is not enabled by the platform or is not available to the supervisor.

3.2. Shared memory physical address range parameter

If an SBI function needs to pass a shared memory physical address range to the SBI implementation (or higher privilege mode), then this physical memory address range MUST satisfy the following requirements:

- The SBI implementation MUST check that the supervisor-mode software is allowed to access the specified physical memory range with the access type requested (read and/or write).
- The SBI implementation MUST access the specified physical memory range using the PMA attributes.



If the supervisor-mode software accesses the same physical memory range using a memory type different than the PMA, then a loss of coherence or unexpected memory ordering may occur. The invoking software should follow the rules and sequences defined in the RISC-V Svpbmt specification to prevent the loss of coherence and memory ordering.

- The data in the shared memory MUST follow little-endian byte ordering.

It is recommended that a memory physical address passed to an SBI function should use at least two `unsigned long` parameters to support platforms which have memory physical addresses wider than XLEN bits.

Chapter 4. Base Extension (EID #0x10)

The base extension is designed to be as small as possible. As such, it only contains functionality for probing which SBI extensions are available and for querying the version of the SBI. All functions in the base extension must be supported by all SBI implementations, so there are no error returns defined.

4.1. Function: Get SBI specification version (FID #0)

```
struct sbiret sbi_get_spec_version(void);
```

Returns the current SBI specification version. This function must always succeed. The minor number of the SBI specification is encoded in the low 24 bits, with the major number encoded in the next 7 bits. Bit 31 must be 0 and is reserved for future expansion.

4.2. Function: Get SBI implementation ID (FID #1)

```
struct sbiret sbi_get_impl_id(void);
```

Returns the current SBI implementation ID, which is different for every SBI implementation. It is intended that this implementation ID allows software to probe for SBI implementation quirks.

4.3. Function: Get SBI implementation version (FID #2)

```
struct sbiret sbi_get_impl_version(void);
```

Returns the current SBI implementation version. The encoding of this version number is specific to the SBI implementation.

4.4. Function: Probe SBI extension (FID #3)

```
struct sbiret sbi_probe_extension(long extension_id);
```

Returns 0 if the given SBI extension ID (EID) is not available, or 1 if it is available unless defined as any other non-zero value by the implementation.

4.5. Function: Get machine vendor ID (FID #4)

```
struct sbiret sbi_get_mvendorid(void);
```

Return a value that is legal for the `mvendorid` CSR and 0 is always a legal value for this CSR.

4.6. Function: Get machine architecture ID (FID #5)

```
struct sbiret sbi_get_marchid(void);
```

Return a value that is legal for the `marchid` CSR and 0 is always a legal value for this CSR.

4.7. Function: Get machine implementation ID (FID #6)

```
struct sbiret sbi_get_mimpid(void);
```

Return a value that is legal for the `mimpid` CSR and 0 is always a legal value for this CSR.

4.8. Function Listing

Table 3. Base Function List

Function Name	SBI Version	FID	EID
<code>sbi_get_sbi_spec_version</code>	0.2	0	0x10
<code>sbi_get_sbi_impl_id</code>	0.2	1	0x10
<code>sbi_get_sbi_impl_version</code>	0.2	2	0x10
<code>sbi_probe_extension</code>	0.2	3	0x10
<code>sbi_get_mvendorid</code>	0.2	4	0x10
<code>sbi_get_marchid</code>	0.2	5	0x10
<code>sbi_get_mimpid</code>	0.2	6	0x10

4.9. SBI Implementation IDs

Table 4. SBI Implementation IDs

Implementation ID	Name
0	Berkeley Boot Loader (BBL)
1	OpenSBI
2	Xvisor

Implementation ID	Name
3	KVM
4	RustSBI
5	Diosix
6	Coffer
7	Xen Project

Chapter 5. Legacy Extensions (EIDs #0x00 - #0x0F)

The legacy SBI extensions follow a slightly different calling convention as compared to the SBI v0.2 (or higher) specification where:

- The SBI function ID field in **a6** register is ignored because these are encoded as multiple SBI extension IDs.
- Nothing is returned in **a1** register.
- All registers except **a0** must be preserved across an SBI call by the callee.
- The value returned in **a0** register is SBI legacy extension specific.

The page and access faults taken by the SBI implementation while accessing memory on behalf of the supervisor are redirected back to the supervisor with **sepc** CSR pointing to the faulting **ECALL** instruction.

The legacy SBI extensions is deprecated in favor of the other extensions listed below. The legacy console SBI functions (**sbi_console_getchar()** and **sbi_console_putchar()**) are expected to be deprecated; they have no replacement.

5.1. Extension: Set Timer (EID #0x00)

```
long sbi_set_timer(uint64_t stime_value)
```

Programs the clock for next event after **stime_value** time. This function also clears the pending timer interrupt bit.

If the supervisor wishes to clear the timer interrupt without scheduling the next timer event, it can either request a timer interrupt infinitely far into the future (i.e., (uint64_t)-1), or it can instead mask the timer interrupt by clearing **sie.STIE** CSR bit.

This SBI call returns 0 upon success or an implementation specific negative error code.

5.2. Extension: Console Puchar (EID #0x01)

```
long sbi_console_putchar(int ch)
```

Write data present in **ch** to debug console.

Unlike **sbi_console_getchar()**, this SBI call **will block** if there remain any pending characters to be transmitted or if the receiving terminal is not yet ready to receive the byte. However, if the console doesn't exist at all, then the character is thrown away.

This SBI call returns 0 upon success or an implementation specific negative error code.

5.3. Extension: Console Getchar (EID #0x02)

```
long sbi_console_getchar(void)
```

Read a byte from debug console.

The SBI call returns the byte on success, or -1 for failure.

5.4. Extension: Clear IPI (EID #0x03)

```
long sbi_clear_ipi(void)
```

Clears the pending IPIs if any. The IPI is cleared only in the hart for which this SBI call is invoked. `sbi_clear_ipi()` is deprecated because S-mode code can clear `sip.SSIP` CSR bit directly.

This SBI call returns 0 if no IPI had been pending, or an implementation specific positive value if an IPI had been pending.

5.5. Extension: Send IPI (EID #0x04)

```
long sbi_send_ipi(const unsigned long *hart_mask)
```

Send an inter-processor interrupt to all the harts defined in `hart_mask`. Interprocessor interrupts manifest at the receiving harts as Supervisor Software Interrupts.

`hart_mask` is a virtual address that points to a bit-vector of harts. The bit vector is represented as a sequence of unsigned longs whose length equals the number of harts in the system divided by the number of bits in an unsigned long, rounded up to the next integer.

This SBI call returns 0 upon success or an implementation specific negative error code.

5.6. Extension: Remote FENCE.I (EID #0x05)

```
long sbi_remote_fence_i(const unsigned long *hart_mask)
```

Instructs remote harts to execute `FENCE.I` instruction. The `hart_mask` is same as described in `sbi_send_ipi()`.

This SBI call returns 0 upon success or an implementation specific negative error code.

5.7. Extension: Remote SFENCE.VMA (EID #0x06)

```
long sbi_remote_sfence_vma(const unsigned long *hart_mask,
                           unsigned long start,
                           unsigned long size)
```

Instructs the remote harts to execute one or more **SFENCE.VMA** instructions, covering the range of virtual addresses between start and size.

This SBI call returns 0 upon success or an implementation specific negative error code.

5.8. Extension: Remote SFENCE.VMA with ASID (EID #0x07)

```
long sbi_remote_sfence_vma_asid(const unsigned long *hart_mask,
                                unsigned long start,
                                unsigned long size,
                                unsigned long asid)
```

Instruct the remote harts to execute one or more **SFENCE.VMA** instructions, covering the range of virtual addresses between start and size. This covers only the given **ASID**.

This SBI call returns 0 upon success or an implementation specific negative error code.

5.9. Extension: System Shutdown (EID #0x08)

```
void sbi_shutdown(void)
```

Puts all the harts to shutdown state from supervisor point of view.

This SBI call doesn't return irrespective whether it succeeds or fails.

5.10. Function Listing

Table 5. Legacy Function List

Function Name	SBI Version	FID	EID	Replacement EID
sbi_set_timer	0.1	0	0x00	0x54494D45
sbi_console_putchar	0.1	0	0x01	N/A
sbi_console_getchar	0.1	0	0x02	N/A
sbi_clear_ipi	0.1	0	0x03	N/A
sbi_send_ipi	0.1	0	0x04	0x735049

Function Name	SBI Version	FID	EID	Replacement EID
sbi_remote_fence_i	0.1	0	0x05	0x52464E43
sbi_remote_sfence_vma	0.1	0	0x06	0x52464E43
sbi_remote_sfence_vma_asid	0.1	0	0x07	0x52464E43
sbi_shutdown	0.1	0	0x08	0x53525354
RESERVED			0x09-0x0F	

Chapter 6. Timer Extension (EID #0x54494D45 "TIME")

This replaces legacy timer extension (EID #0x00). It follows the new calling convention defined in v0.2.

6.1. Function: Set Timer (FID #0)

```
struct sbiret sbi_set_timer(uint64_t stime_value)
```

Programs the clock for next event after **stime_value** time. **stime_value** is in absolute time. This function must clear the pending timer interrupt bit as well.

If the supervisor wishes to clear the timer interrupt without scheduling the next timer event, it can either request a timer interrupt infinitely far into the future (i.e., (uint64_t)-1), or it can instead mask the timer interrupt by clearing **sie.STIE** CSR bit.

6.2. Function Listing

Table 6. TIME Function List

Function Name	SBI Version	FID	EID
sbi_set_timer	0.2	0	0x54494D45

Chapter 7. IPI Extension (EID #0x735049 "sPI: s-mode IPI")

This extension replaces the legacy extension (EID #0x04). The other IPI related legacy extension(0x3) is deprecated now. All the functions in this extension follow the `hart_mask` as defined in the binary encoding section.

7.1. Function: Send IPI (FID #0)

```
struct sbiret sbi_send_ipi(unsigned long hart_mask,
                          unsigned long hart_mask_base)
```

Send an inter-processor interrupt to all the harts defined in `hart_mask`. Interprocessor interrupts manifest at the receiving harts as the supervisor software interrupts.

The possible error codes returned in `sbiret.error` are shown in the [Table 7](#) below.

Table 7. IPI Send Errors

Error code	Description
SBI_SUCCESS	IPI was sent to all the targeted harts successfully.

7.2. Function Listing

Table 8. IPI Function List

Function Name	SBI Version	FID	EID
sbi_send_ipi	0.2	0	0x735049

Chapter 8. RFENCE Extension (EID #0x52464E43 "RFNC")

This extension defines all remote fence related functions and replaces the legacy extensions (EIDs #0x05 - #0x07). All the functions follow the `hart_mask` as defined in binary encoding section. Any function wishes to use range of addresses (i.e. `start_addr` and `size`), have to abide by the below constraints on range parameters.

The remote fence function acts as a full TLB flush if

- `start_addr` and `size` are both 0
- `size` is equal to 2^{XLEN-1}

8.1. Function: Remote FENCE.I (FID #0)

```
struct sbiret sbi_remote_fence_i(unsigned long hart_mask,
                                unsigned long hart_mask_base)
```

Instructs remote harts to execute `FENCE.I` instruction.

The possible error codes returned in `sbiret.error` are shown in the [Table 9](#) below.

Table 9. RFENCE Remote FENCE.I Errors

Error code	Description
SBI_SUCCESS	IPI was sent to all the targeted harts successfully.

8.2. Function: Remote SFENCE.VMA (FID #1)

```
struct sbiret sbi_remote_sfence_vma(unsigned long hart_mask,
                                     unsigned long hart_mask_base,
                                     unsigned long start_addr,
                                     unsigned long size)
```

Instructs the remote harts to execute one or more `SFENCE.VMA` instructions, covering the range of virtual addresses between `start` and `size`.

The possible error codes returned in `sbiret.error` are shown in the [Table 10](#) below.

Table 10. RFENCE Remote SFENCE.VMA Errors

Error code	Description
SBI_SUCCESS	IPI was sent to all the targeted harts successfully.
SBI_ERR_INVALID_ADDRESS	<code>start_addr</code> or <code>size</code> is not valid.

8.3. Function: Remote SFENCE.VMA with ASID (FID #2)

```
struct sbiret sbi_remote_sfence_vma_asid(unsigned long hart_mask,
                                         unsigned long hart_mask_base,
                                         unsigned long start_addr,
                                         unsigned long size,
                                         unsigned long asid)
```

Instruct the remote harts to execute one or more **SFENCE.VMA** instructions, covering the range of virtual addresses between start and size. This covers only the given **ASID**.

The possible error codes returned in **sbiret.error** are shown in the [Table 11](#) below.

Table 11. RFENCE Remote SFENCE.VMA with ASID Errors

Error code	Description
SBI_SUCCESS	IPI was sent to all the targeted harts successfully.
SBI_ERR_INVALID_ADDRESS	start_addr or size is not valid.

8.4. Function: Remote HFENCE.GVMA with VMID (FID #3)

```
struct sbiret sbi_remote_hfence_gvma_vmid(unsigned long hart_mask,
                                           unsigned long hart_mask_base,
                                           unsigned long start_addr,
                                           unsigned long size,
                                           unsigned long vmid)
```

Instruct the remote harts to execute one or more **HFENCE.GVMA** instructions, covering the range of guest physical addresses between start and size only for the given **VMID**. This function call is only valid for harts implementing hypervisor extension.

The possible error codes returned in **sbiret.error** are shown in the [Table 12](#) below.

Table 12. RFENCE Remote HFENCE.GVMA with VMID Errors

Error code	Description
SBI_SUCCESS	IPI was sent to all the targeted harts successfully.
SBI_ERR_NOT_SUPPORTED	This function is not supported as it is not implemented or one of the target hart doesn't support hypervisor extension.
SBI_ERR_INVALID_ADDRESS	start_addr or size is not valid.

8.5. Function: Remote HFENCE.GVMA (FID #4)

```
struct sbiret sbi_remote_hfence_gvma(unsigned long hart_mask,
                                     unsigned long hart_mask_base,
                                     unsigned long start_addr,
                                     unsigned long size)
```

Instruct the remote harts to execute one or more **HFENCE.GVMA** instructions, covering the range of guest physical addresses between start and size for all the guests. This function call is only valid for harts implementing hypervisor extension.

The possible error codes returned in **sbiret.error** are shown in the [Table 13](#) below.

Table 13. RFENCE Remote HFENCE.GVMA Errors

Error code	Description
SBI_SUCCESS	IPI was sent to all the targeted harts successfully.
SBI_ERR_NOT_SUPPORTED	This function is not supported as it is not implemented or one of the target hart doesn't support hypervisor extension.
SBI_ERR_INVALID_ADDRESS	start_addr or size is not valid.

8.6. Function: Remote HFENCE.VVMA with ASID (FID #5)

```
struct sbiret sbi_remote_hfence_vvma_asid(unsigned long hart_mask,
                                           unsigned long hart_mask_base,
                                           unsigned long start_addr,
                                           unsigned long size,
                                           unsigned long asid)
```

Instruct the remote harts to execute one or more **HFENCE.VVMA** instructions, covering the range of guest virtual addresses between start and size for the given **ASID** and current **VMID** (in **hgap** CSR) of calling hart. This function call is only valid for harts implementing hypervisor extension.

The possible error codes returned in **sbiret.error** are shown in the [Table 14](#) below.

Table 14. RFENCE Remote HFENCE.VVMA with ASID Errors

Error code	Description
SBI_SUCCESS	IPI was sent to all the targeted harts successfully.
SBI_ERR_NOT_SUPPORTED	This function is not supported as it is not implemented or one of the target hart doesn't support hypervisor extension.
SBI_ERR_INVALID_ADDRESS	start_addr or size is not valid.

8.7. Function: Remote HFENCE.VVMA (FID #6)

```
struct sbiret sbi_remote_hfence_vvma(unsigned long hart_mask,
                                     unsigned long hart_mask_base,
                                     unsigned long start_addr,
                                     unsigned long size)
```

Instruct the remote harts to execute one or more **HFENCE.VVMA** instructions, covering the range of guest virtual addresses between start and size for current **VMID** (in **hgatp** CSR) of calling hart. This function call is only valid for harts implementing hypervisor extension.

The possible error codes returned in **sbiret.error** are shown in the [Table 15](#) below.

Table 15. RFENCE Remote HFENCE.VVMA Errors

Error code	Description
SBI_SUCCESS	IPI was sent to all the targeted harts successfully.
SBI_ERR_NOT_SUPPORTED	This function is not supported as it is not implemented or one of the target hart doesn't support hypervisor extension.
SBI_ERR_INVALID_ADDRESS	start_addr or size is not valid.

8.8. Function Listing

Table 16. RFENCE Function List

Function Name	SBI Version	FID	EID
sbi_remote_fence_i	0.2	0	0x52464E43
sbi_remote_sfence_vma	0.2	1	0x52464E43
sbi_remote_sfence_vma_asid	0.2	2	0x52464E43
sbi_remote_hfence_gvma_vmid	0.2	3	0x52464E43
sbi_remote_hfence_gvma	0.2	4	0x52464E43
sbi_remote_hfence_vvma_asid	0.2	5	0x52464E43
sbi_remote_hfence_vvma	0.2	6	0x52464E43

Chapter 9. Hart State Management Extension (EID #0x48534D "HSM")

The Hart State Management (HSM) Extension introduces a set of hart states and a set of functions which allow the supervisor-mode software to request a hart state change.

The [Table 17](#) shown below describes all possible **HSM states** along with a unique **HSM state id** for each state:

Table 17. HSM Hart States

State ID	State Name	Description
0	STARTED	The hart is physically powered-up and executing normally.
1	STOPPED	The hart is not executing in supervisor-mode or any lower privilege mode. It is probably powered-down by the SBI implementation if the underlying platform has a mechanism to physically power-down harts.
2	START_PENDING	Some other hart has requested to start (or power-up) the hart from the STOPPED state and the SBI implementation is still working to get the hart in the STARTED state.
3	STOP_PENDING	The hart has requested to stop (or power-down) itself from the STARTED state and the SBI implementation is still working to get the hart in the STOPPED state.
4	SUSPENDED	This hart is in a platform specific suspend (or low power) state.
5	SUSPEND_PENDING	The hart has requested to put itself in a platform specific low power state from the STARTED state and the SBI implementation is still working to get the hart in the platform specific SUSPENDED state.
6	RESUME_PENDING	An interrupt or platform specific hardware event has caused the hart to resume normal execution from the SUSPENDED state and the SBI implementation is still working to get the hart in the STARTED state.

At any point in time, a hart should be in one of the above mentioned hart states. The hart state transitions by the SBI implementation should follow the state machine shown below in the [Figure 3](#).




Figure 3. SBI HSM State Machine

A platform can have multiple harts grouped into hierarchical topology groups (namely cores, clusters, nodes, etc.) with separate platform specific low-power states for each hierarchical group. These platform specific low-power states of hierarchical topology groups can be represented as platform specific suspend states of a hart. An SBI implementation can utilize the suspend states of higher topology groups using one of the following approaches:

1. **Platform-coordinated:** In this approach, when a hart becomes idle the supervisor-mode power-managment software will request deepest suspend state for the hart and higher topology groups. An SBI implementation should choose a suspend state at higher topology group which is:
 - a. Not deeper than the specified suspend state
 - b. Wake-up latency is not higher than the wake-up latency of the specified suspend state
2. **OS-initiated:** In this approach, the supervisor-mode power-managment software will directly request a suspend state for higher topology group after the last hart in that group becomes idle. When a hart becomes idle, the supervisor-mode power-managment software will always select suspend state for the hart itself but it will select a suspend state for a higher topology group only if the hart is the last running hart in the group. An SBI implementation should:
 - a. Never choose a suspend state for higher topology group different from the specified suspend state
 - b. Always prefer most recent suspend state requested for higher topology group

9.1. Function: HART start (FID #0)

```

struct sbiret sbi_hart_start(unsigned long hartid,
                           unsigned long start_addr,
                           unsigned long opaque)

```

Request the SBI implementation to start executing the target hart in supervisor-mode, at the address specified by `start_addr`, with the specific register values described in [Table 18](#).

Table 18. HSM Hart Start Register State

Register Name	Register Value
satp	0
sstatus.SIE	0
a0	hartid
a1	opaque parameter
All other registers remain in an undefined state.	



A single **unsigned long** parameter is sufficient as `start_addr`, because the hart will start execution in supervisor-mode with the MMU off, hence `start_addr` must be less than XLEN bits wide.

This call is asynchronous — more specifically, the `sbi_hart_start()` may return before the target hart starts executing as long as the SBI implementation is capable of ensuring the return code is accurate. If the SBI implementation is a platform runtime firmware executing in machine-mode (M-mode), then it **MUST** configure any physical memory protection it supports, such as that defined by PMP, and other M-mode state, before transferring control to supervisor-mode software.

The `hartid` parameter specifies the target hart which is to be started.

The `start_addr` parameter points to a runtime-specified physical address, where the hart can start executing in supervisor-mode.

The `opaque` parameter is an XLEN-bit value which will be set in the `a1` register when the hart starts executing at `start_addr`.

The possible error codes returned in `sbiret.error` are shown in the [Table 19](#) below.

Table 19. HSM Hart Start Errors

Error code	Description
SBI_SUCCESS	Hart was previously in stopped state. It will start executing from <code>start_addr</code> .
SBI_ERR_INVALID_ADDRESS	<code>start_addr</code> is not valid, possibly due to the following reasons: * It is not a valid physical address. * Executable access to the address is prohibited by a physical memory protection mechanism or H-extension G-stage for supervisor-mode.
SBI_ERR_INVALID_PARAM	<code>hartid</code> is not a valid hartid as the corresponding hart cannot be started in supervisor mode.
SBI_ERR_ALREADY_AVAILABLE	The given hartid is already started.
SBI_ERR_FAILED	The start request failed for unspecified or unknown other reasons.

9.2. Function: HART stop (FID #1)

```
struct sbiret sbi_hart_stop(void)
```

Request the SBI implementation to stop executing the calling hart in supervisor-mode and return its ownership to the SBI implementation. This call is not expected to return under normal conditions. The `sbi_hart_stop()` must be called with supervisor-mode interrupts disabled.

The possible error codes returned in `sbiret.error` are shown in the [Table 20](#) below.

Table 20. HSM Hart Stop Errors

Error code	Description
SBI_ERR_FAILED	Failed to stop execution of the current hart

9.3. Function: HART get status (FID #2)

```
struct sbiret sbi_hart_get_status(unsigned long hartid)
```

Get the current status (or HSM state id) of the given hart in `sbiret.value`, or an error through `sbiret.error`.

The `hartid` parameter specifies the target hart for which status is required.

The possible status (or HSM state id) values returned in `sbiret.value` are described in [Table 17](#).

The possible error codes returned in `sbiret.error` are shown in the [Table 21](#) below.

Table 21. HSM Hart Get Status Errors

Error code	Description
SBI_ERR_INVALID_PARAM	The given <code>hartid</code> is not valid.

The harts may transition HSM states at any time due to any concurrent `sbi_hart_start()` or `sbi_hart_stop()` or `sbi_hart_suspend()` calls, the return value from this function may not represent the actual state of the hart at the time of return value verification.

9.4. Function: HART suspend (FID #3)

```
struct sbiret sbi_hart_suspend(uint32_t suspend_type,
                               unsigned long resume_addr,
                               unsigned long opaque)
```

Request the SBI implementation to put the calling hart in a platform specific suspend (or low power) state specified by the `suspend_type` parameter. The hart will automatically come out of suspended state and resume normal execution when it receives an interrupt or platform specific hardware event.

The platform specific suspend states for a hart can be either retentive or non-retentive in nature. A retentive suspend state will preserve hart register and CSR values for all privilege modes whereas a non-retentive suspend state will not preserve hart register and CSR values.

Resuming from a retentive suspend state is straight forward and the supervisor-mode software will see SBI suspend call return without any failures. The `resume_addr` parameter is unused during retentive suspend.

Resuming from a non-retentive suspend state is relatively more involved and requires software to restore various hart registers and CSRs for all privilege modes. Upon resuming from non-retentive suspend state, the hart will jump to supervisor-mode at address specified by `resume_addr` with specific registers values described in the [Table 22](#) below.

Table 22. HSM Hart Resume Register State

Register Name	Register Value
satp	0
sstatus.SIE	0
a0	hartid
a1	<code>opaque</code> parameter
All other registers remain in an undefined state.	



A single `unsigned long` parameter is sufficient for `resume_addr`, because the hart will resume execution in supervisor-mode with the MMU off, hence `resume_addr` must be less than XLEN bits wide.

The `suspend_type` parameter is 32 bits wide and the possible values are shown in [Table 23](#) below.

Table 23. HSM Hart Suspend Types

Value	Description
0x00000000	Default retentive suspend
0x00000001 - 0x0FFFFFFF	Reserved for future use
0x10000000 - 0x7FFFFFFF	Platform specific retentive suspend
0x80000000	Default non-retentive suspend
0x80000001 - 0x8FFFFFFF	Reserved for future use
0x90000000 - 0xFFFFFFFF	Platform specific non-retentive suspend
> 0xFFFFFFFF	Reserved

The `resume_addr` parameter points to a runtime-specified physical address, where the hart can resume execution in supervisor-mode after a non-retentive suspend.

The `opaque` parameter is an XLEN-bit value which will be set in the `a1` register when the hart resumes execution at `resume_addr` after a non-retentive suspend.

The possible error codes returned in `sbiret.error` are shown in the [Table 24](#) below.

Table 24. HSM Hart Suspend Errors

Error code	Description
SBI_SUCCESS	Hart has suspended and resumed successfully from a retentive suspend state.
SBI_ERR_INVALID_PARAM	<code>suspend_type</code> is reserved or is platform-specific and unimplemented.
SBI_ERR_NOT_SUPPORTED	<code>suspend_type</code> is not reserved and is implemented, but the platform does not support it due to one or more missing dependencies.
SBI_ERR_INVALID_ADDRESS	<code>resume_addr</code> is not valid, possibly due to the following reasons: * It is not a valid physical address. * Executable access to the address is prohibited by a physical memory protection mechanism or H-extension G-stage for supervisor-mode.
SBI_ERR_FAILED	The suspend request failed for unspecified or unknown other reasons.

9.5. Function Listing

Table 25. HSM Function List

Function Name	SBI Version	FID	EID
<code>sbi_hart_start</code>	0.2	0	0x48534D
<code>sbi_hart_stop</code>	0.2	1	0x48534D
<code>sbi_hart_get_status</code>	0.2	2	0x48534D
<code>sbi_hart_suspend</code>	0.3	3	0x48534D

Chapter 10. System Reset Extension (EID #0x53525354 "SRST")

The System Reset Extension provides a function that allow the supervisor software to request system-level reboot or shutdown. The term "system" refers to the world-view of supervisor software and the underlying SBI implementation could be provided by machine mode firmware or a hypervisor.

10.1. Function: System reset (FID #0)

```
struct sbiret sbi_system_reset(uint32_t reset_type, uint32_t reset_reason)
```

Reset the system based on provided `reset_type` and `reset_reason`. This is a synchronous call and does not return if it succeeds.

The `reset_type` parameter is 32 bits wide and it's possible values are shown in the [Table 26](#) below.

Table 26. SRST System Reset Types

Value	Description
0x00000000	Shutdown
0x00000001	Cold reboot
0x00000002	Warm reboot
0x00000003 - 0xEFFFFFFF	Reserved for future use
0xF0000000 - 0xFFFFFFFF	Vendor or platform specific reset type
> 0xFFFFFFFF	Reserved

The `reset_reason` is an optional parameter representing the reason for system reset. This parameter is 32 bits wide with possible values shown in the [Table 27](#) below

Table 27. SRST System Reset Reasons

Value	Description
0x00000000	No reason
0x00000001	System failure
0x00000002 - 0xDFFFFFFF	Reserved for future use
0xE0000000 - 0xEFFFFFFF	SBI implementation specific reset reason
0xF0000000 - 0xFFFFFFFF	Vendor or platform specific reset reason
> 0xFFFFFFFF	Reserved

When supervisor software is running natively, the SBI implementation is provided by machine mode firmware. In this case, shutdown is equivalent to a physical power down of the entire system and cold reboot is equivalent to a physical power cycle of the entire system. Further, warm reboot is equivalent to a power cycle of the main processor and parts of the system, but not the entire system. For example, on a server class system with a BMC (board management controller), a warm reboot will not power cycle the BMC whereas a cold reboot will definitely power cycle the BMC.

When supervisor software is running inside a virtual machine, the SBI implementation is provided by a hypervisor. Shutdown, cold reboot and warm reboot will behave functionally the same as the native case, but might not result in any physical power changes.

The possible error codes returned in `sbiret.error` are shown in the [Table 28](#) below.

Table 28. SRST System Reset Errors

Error code	Description
SBI_ERR_INVALID_PARAM	At least one of <code>reset_type</code> or <code>reset_reason</code> is reserved or is platform-specific and unimplemented.
SBI_ERR_NOT_SUPPORTED	<code>reset_type</code> is not reserved and is implemented, but the platform does not support it due to one or more missing dependencies.
SBI_ERR_FAILED	The reset request failed for unspecified or unknown other reasons.

10.2. Function Listing

Table 29. SRST Function List

Function Name	SBI Version	FID	EID
<code>sbi_system_reset</code>	0.3	0	0x53525354

Chapter 11. Performance Monitoring Unit Extension (EID #0x504D55 "PMU")

The RISC-V hardware performance counters such as `mcycle`, `minstret`, and `mhpmpcounterX` CSRs are accessible as read-only from supervisor-mode using `cycle`, `instret`, and `hpmcounterX` CSRs. The SBI performance monitoring unit (PMU) extension is an interface for supervisor-mode to configure and use the RISC-V hardware performance counters with assistance from the machine-mode (or hypervisor-mode). These hardware performance counters can only be started, stopped, or configured from machine-mode using `mcountinhibit` and `mhpmeventX` CSRs. Due to this, a machine-mode SBI implementation may choose to disallow SBI PMU extension if `mcountinhibit` CSR is not implemented by the RISC-V platform.

A RISC-V platform generally supports monitoring of various hardware events using a limited number of hardware performance counters which are up to 64 bits wide. In addition, a SBI implementation can also provide firmware performance counters which can monitor firmware events such as number of misaligned load/store instructions, number of RFENCEs, number of IPIs, etc. All firmware counters must have same number of bits and can be up to 64 bits wide.

The SBI PMU extension provides:

1. An interface for supervisor-mode software to discover and configure per-HART hardware/firmware counters
2. A typical `perf` compatible interface for hardware/firmware performance counters and events
3. Full access to microarchitecture's raw event encodings

To define SBI PMU extension calls, we first define important entities `counter_idx`, `event_idx`, and `event_data`. The `counter_idx` is a logical number assigned to each hardware/firmware counter. The `event_idx` represents a hardware (or firmware) event whereas the `event_data` is 64 bits wide and represents additional configuration (or parameters) for a hardware (or firmware) event.

The `event_idx` is a 20 bits wide number encoded as follows:

```
event_idx[19:16] = type
event_idx[15:0] = code
```

11.1. Event: Hardware general events (Type #0)

The `event_idx.type` (i.e. **event type**) should be `0x0` for all hardware general events and each hardware general event is identified by a unique `event_idx.code` (i.e. **event code**) described in the [Table 30](#) below.

Table 30. PMU Hardware Events

General Event Name	Code	Description
SBI_PMU_HW_NO_EVENT	0	Unused event because <code>event_idx</code> cannot be zero

General Event Name	Code	Description
SBI_PMU_HW_CPU_CYCLES	1	Event for each CPU cycle
SBI_PMU_HW_INSTRUCTIONS	2	Event for each completed instruction
SBI_PMU_HW_CACHE_REFERENCES	3	Event for cache hit
SBI_PMU_HW_CACHE_MISSES	4	Event for cache miss
SBI_PMU_HW_BRANCH_INSTRUCTIONS	5	Event for a branch instruction
SBI_PMU_HW_BRANCH_MISSES	6	Event for a branch misprediction
SBI_PMU_HW_BUS_CYCLES	7	Event for each BUS cycle
SBI_PMU_HW_STALLED_CYCLES_FRONTEND	8	Event for a stalled cycle in microarchitecture frontend
SBI_PMU_HW_STALLED_CYCLES_BACKEND	9	Event for a stalled cycle in microarchitecture backend
SBI_PMU_HW_REF_CPU_CYCLES	10	Event for each reference CPU cycle

The **event_data** (i.e. **event data**) is unused for hardware general events and all non-zero values of **event_data** are reserved for future use.



A RISC-V platform might halt the CPU clock when it enters WAIT state using the WFI instruction or enters platform specific SUSPEND state using the SBI HSM HART suspend call.



The **SBI_PMU_HW_CPU_CYCLES** event counts CPU clock cycles as counted by the **cycle** CSR. These may be variable frequency cycles, and are not counted when the CPU clock is halted.



The **SBI_PMU_HW_REF_CPU_CYCLES** counts fixed-frequency clock cycles while the CPU clock is not halted. The fixed-frequency of counting might, for example, be the same frequency at which the **time** CSR counts.



The **SBI_PMU_HW_BUS_CYCLES** counts fixed-frequency clock cycles. The fixed-frequency of counting might be the same frequency at which the **time** CSR counts, or may be the frequency of the clock at the boundary between the HART (and its private caches) and the rest of the system.

11.2. Event: Hardware cache events (Type #1)

The **event_idx.type** (i.e. **event type**) should be **0x1** for all hardware cache events and each hardware cache event is identified by a unique **event_idx.code** (i.e. **event code**) which is encoded as follows:

```

event_idx.code[15:3] = cache_id
event_idx.code[2:1] = op_id
event_idx.code[0:0] = result_id

```

Below tables show possible values of: `event_idx.code.cache_id` (i.e. cache event id), `event_idx.code.op_id` (i.e. cache operation id) and `event_idx.code.result_id` (i.e. cache result id).

Table 31. PMU Cache Event ID

Cache Event Name	Event ID	Description
SBI_PMU_HW_CACHE_L1D	0	Level1 data cache event
SBI_PMU_HW_CACHE_L1I	1	Level1 instruction cache event
SBI_PMU_HW_CACHE_LL	2	Last level cache event
SBI_PMU_HW_CACHE_DTLB	3	Data TLB event
SBI_PMU_HW_CACHE_ITLB	4	Instruction TLB event
SBI_PMU_HW_CACHE_BPU	5	Branch predictor unit event
SBI_PMU_HW_CACHE_NODE	6	NUMA node cache event

Table 32. PMU Cache Operation ID

Cache Operation Name	Operation ID	Description
SBI_PMU_HW_CACHE_OP_READ	0	Read cache line
SBI_PMU_HW_CACHE_OP_WRITE	1	Write cache line
SBI_PMU_HW_CACHE_OP_PREFETCH	2	Prefetch cache line

Table 33. PMU Cache Operation Result ID

Cache Result Name	Result ID	Description
SBI_PMU_HW_CACHE_RESULT_ACCESS	0	Cache access
SBI_PMU_HW_CACHE_RESULT_MISS	1	Cache miss

The `event_data` (i.e. event data) is unused for hardware cache events and all non-zero values of `event_data` are reserved for future use.

11.3. Event: Hardware raw events (Type #2)

The `event_idx.type` (i.e. event type) should be `0x2` for all hardware raw events and `event_idx.code` (i.e. event code) should be zero.

On RISC-V platform with 32 bits wide `mhpmeventX` CSRs, the `event_data` configuration (or parameter) should have the 32-bit value to be programmed in the `mhpmeventX` CSR.

On RISC-V platform with 64 bits wide `mhpmeventX` CSRs, the `event_data` configuration (or parameter) should have the 48-bit value to be programmed in the lower 48-bits of `mhpmeventX` CSR and the SBI implementation shall determine the value to be programmed in the upper 16 bits of

`mhpmeventX` CSR.



The RISC-V platform hardware implementation may choose to define the expected value to be written to `mhpmeventX` CSR for a hardware event. In case of hardware general/cache events, the RISC-V platform hardware implementation may use the zero-extended `event_idx` as the expected value for simplicity.

11.4. Event: Firmware events (Type #15)

The `event_idx.type` (i.e. **event type**) should be `0xf` for all firmware events and each firmware event is identified by an unique `event_idx.code` (i.e. **event code**) described in the [Table 34](#) below.

Table 34. PMU Firmware Events

Firmware Event Name	Code	Description
SBI_PMU_FW_MISALIGNED_LOAD	0	Misaligned load trap event
SBI_PMU_FW_MISALIGNED_STORE	1	Misaligned store trap event
SBI_PMU_FW_ACCESS_LOAD	2	Load access trap event
SBI_PMU_FW_ACCESS_STORE	3	Store access trap event
SBI_PMU_FW_ILLEGAL_INSN	4	Illegal instruction trap event
SBI_PMU_FW_SET_TIMER	5	Set timer event
SBI_PMU_FW_IPI_SENT	6	Sent IPI to other HART event
SBI_PMU_FW_IPI_RECEIVED	7	Received IPI from other HART event
SBI_PMU_FW_FENCE_I_SENT	8	Sent FENCE.I request to other HART event
SBI_PMU_FW_FENCE_I_RECEIVED	9	Received FENCE.I request from other HART event
SBI_PMU_FW_SFENCE_VMA_SENT	10	Sent SFENCE.VMA request to other HART event
SBI_PMU_FW_SFENCE_VMA_RECEIVED	11	Received SFENCE.VMA request from other HART event
SBI_PMU_FW_SFENCE_VMA_ASID_SENT	12	Sent SFENCE.VMA with ASID request to other HART event
SBI_PMU_FW_SFENCE_VMA_ASID_RECEIVED	13	Received SFENCE.VMA with ASID request from other HART event
SBI_PMU_FW_HFENCE_GVMA_SENT	14	Sent HFENCE.GVMA request to other HART event
SBI_PMU_FW_HFENCE_GVMA_RECEIVED	15	Received HFENCE.GVMA request from other HART event

Firmware Event Name	Code	Description
SBI_PMU_FW_HFENCE_GVMA_VMID_SENT	16	Sent HFENCE.GVMA with VMID request to other HART event
SBI_PMU_FW_HFENCE_GVMA_VMID_RECEIVED	17	Received HFENCE.GVMA with VMID request from other HART event
SBI_PMU_FW_HFENCE_VVMA_SENT	18	Sent HFENCE.VVMA request to other HART event
SBI_PMU_FW_HFENCE_VVMA_RECEIVED	19	Received HFENCE.VVMA request from other HART event
SBI_PMU_FW_HFENCE_VVMA_ASID_SENT	20	Sent HFENCE.VVMA with ASID request to other HART event
SBI_PMU_FW_HFENCE_VVMA_ASID_RECEIVED	21	Received HFENCE.VVMA with ASID request from other HART event
Reserved	22 - 255	Reserved for future use
Implementation specific events	256 - 65534	SBI implementation specific firmware events
SBI_PMU_FW_PLATFORM	65535	RISC-V platform specific firmware events, where the <code>event_data</code> configuration (or parameter) contains the event encoding.

For all firmware events except SBI_PMU_FW_PLATFORM, the `event_data` configuration (or parameter) is unused and all non-zero values of `event_data` are reserved for future use.

11.5. Function: Get number of counters (FID #0)

```
struct sbiret sbi_pmu_num_counters()
```

Returns the number of counters (both hardware and firmware) in `sbiret.value` and always returns `SBI_SUCCESS` in `sbiret.error`.

11.6. Function: Get details of a counter (FID #1)

```
struct sbiret sbi_pmu_counter_get_info(unsigned long counter_idx)
```

Get details about the specified counter such as underlying CSR number, width of the counter, type of counter hardware/firmware, etc.

The `counter_info` returned by this SBI call is encoded as follows:

```
counter_info[11:0] = CSR (12bit CSR number)
counter_info[17:12] = Width (One less than number of bits in CSR)
counter_info[XLEN-2:18] = Reserved for future use
counter_info[XLEN-1] = Type (0 = hardware and 1 = firmware)
```

If `counter_info.type == 1` then `counter_info.csr` and `counter_info.width` should be ignored.

Returns the `counter_info` described above in `sbiret.value`.

The possible error codes returned in `sbiret.error` are shown in the [Table 35](#) below.

Table 35. PMU Counter Get Info Errors

Error code	Description
SBI_SUCCESS	<code>counter_info</code> read successfully.
SBI_ERR_INVALID_PARAM	<code>counter_idx</code> points to an invalid counter.

11.7. Function: Find and configure a matching counter (FID #2)

```
struct sbiret sbi_pmu_counter_config_matching(unsigned long counter_idx_base,
                                             unsigned long counter_idx_mask,
                                             unsigned long config_flags,
                                             unsigned long event_idx,
                                             uint64_t event_data)
```

Find and configure a counter from a set of counters which is not started (or enabled) and can monitor the specified event. The `counter_idx_base` and `counter_idx_mask` parameters represent the set of counters whereas `event_idx` represents the event to be monitored and `event_data` represents any additional event configuration.

The `config_flags` parameter represents additional counter configuration and filter flags. The bit definitions of the `config_flags` parameter are shown in the [Table 36](#) below.

Table 36. PMU Counter Config Match Flags

Flag Name	Bits	Description
SBI_PMU_CFG_FLAG_SKIP_MATCH	0:0	Skip the counter matching
SBI_PMU_CFG_FLAG_CLEAR_VALUE	1:1	Clear (or zero) the counter value in counter configuration
SBI_PMU_CFG_FLAG_AUTO_START	2:2	Start the counter after configuring a matching counter

Flag Name	Bits	Description
SBI_PMU_CFG_FLAG_SET_VUINH	3:3	Event counting inhibited in VU-mode
SBI_PMU_CFG_FLAG_SET_VSINH	4:4	Event counting inhibited in VS-mode
SBI_PMU_CFG_FLAG_SET_UINH	5:5	Event counting inhibited in U-mode
SBI_PMU_CFG_FLAG_SET_SINH	6:6	Event counting inhibited in S-mode
SBI_PMU_CFG_FLAG_SET_MINH	7:7	Event counting inhibited in M-mode
RESERVED	8:(XLEN-1)	All non-zero values are reserved for future use



When `SBI_PMU_CFG_FLAG_SKIP_MATCH` is set in `config_flags`, the SBI implementation will unconditionally select the first counter from the set of counters specified by the `counter_idx_base` and `counter_idx_mask`.



The `SBI_PMU_CFG_FLAG_AUTO_START` flag in `config_flags` has no impact on the counter value.



The `config_flags[3:7]` bits are event filtering hints so these can be ignored or overridden by the SBI implementation for security concerns or due to lack of event filtering support in the underlying RISC-V platform.

Returns the `counter_idx` in `sbiret.value` upon success.

In case of failure, the possible error codes returned in `sbiret.error` are shown in the [Table 37](#) below.

Table 37. PMU Counter Config Match Errors

Error code	Description
SBI_SUCCESS	counter found and configured successfully.
SBI_ERR_INVALID_PARAM	set of counters has at least one invalid counter.
SBI_ERR_NOT_SUPPORTED	none of the counters can monitor the specified event.

11.8. Function: Start a set of counters (FID #3)

```
struct sbiret sbi_pmu_counter_start(unsigned long counter_idx_base,
                                   unsigned long counter_idx_mask,
                                   unsigned long start_flags,
                                   uint64_t initial_value)
```

Start or enable a set of counters on the calling HART with the specified initial value. The `counter_idx_base` and `counter_idx_mask` parameters represent the set of counters whereas the

`initial_value` parameter specifies the initial value of the counter.

The bit definitions of the `start_flags` parameter are shown in the Table 38 below.

Table 38. PMU Counter Start Flags

Flag Name	Bits	Description
SBI_PMU_START_SET_INIT_VALUE	0:0	Set the value of counters based on the <code>initial_value</code> parameter
RESERVED	1:(XLEN-1)	All non-zero values are reserved for future use



When `SBI_PMU_START_SET_INIT_VALUE` is not set in `start_flags`, the counter value will not be modified and event counting will start from current counter value.

The possible error codes returned in `sbiret.error` are shown in the Table 39 below.

Table 39. PMU Counter Start Errors

Error code	Description
SBI_SUCCESS	counter started successfully.
SBI_ERR_INVALID_PARAM	set of counters has at least one invalid counter.
SBI_ERR_ALREADY_STARTED	set of counters includes at least one counter which is already started.

11.9. Function: Stop a set of counters (FID #4)

```
struct sbiret sbi_pmu_counter_stop(unsigned long counter_idx_base,
                                   unsigned long counter_idx_mask,
                                   unsigned long stop_flags)
```

Stop or disable a set of counters on the calling HART. The `counter_idx_base` and `counter_idx_mask` parameters represent the set of counters. The bit definitions of the `stop_flags` parameter are shown in the Table 40 below.

Table 40. PMU Counter Stop Flags

Flag Name	Bits	Description
SBI_PMU_STOP_FLAG_RESET	0:0	Reset the counter to event mapping.
RESERVED	1:(XLEN-1)	All non-zero values are reserved for future use

The possible error codes returned in `sbiret.error` are shown in the Table 41 below.

Table 41. PMU Counter Stop Errors

Error code	Description
SBI_SUCCESS	counter stopped successfully.
SBI_ERR_INVALID_PARAM	set of counters has at least one invalid counter.
SBI_ERR_ALREADY_STOPPED	set of counters includes at least one counter which is already stopped.

11.10. Function: Read a firmware counter (FID #5)

```
struct sbiret sbi_pmu_counter_fw_read(unsigned long counter_idx)
```

Provide the current firmware counter value in `sbiret.value`. On RV32 systems, the `sbiret.value` will only contain the lower 32 bits of the current firmware counter value.

The possible error codes returned in `sbiret.error` are shown in the [Table 42](#) below.

Table 42. PMU Counter Firmware Read Errors

Error code	Description
SBI_SUCCESS	firmware counter read successfully.
SBI_ERR_INVALID_PARAM	<code>counter_idx</code> points to a hardware counter or an invalid counter.

11.11. Function: Read a firmware counter high bits (FID #6)

```
struct sbiret sbi_pmu_counter_fw_read_hi(unsigned long counter_idx)
```

Provide the upper 32 bits of the current firmware counter value in `sbiret.value`. This function always returns zero in `sbiret.value` for RV64 (or higher) systems.

The possible error codes returned in `sbiret.error` are shown in [Table 43](#) below.

Table 43. PMU Counter Firmware Read High Errors

Error code	Description
SBI_SUCCESS	Firmware counter read successfully.
SBI_ERR_INVALID_PARAM	<code>counter_idx</code> points to a hardware counter or an invalid counter.

11.12. Function Listing

Table 44. PMU Function List

Function Name	SBI Version	FID	EID
sbi_pmu_num_counters	0.3	0	0x504D55
sbi_pmu_counter_get_info	0.3	1	0x504D55
sbi_pmu_counter_config_matching	0.3	2	0x504D55
sbi_pmu_counter_start	0.3	3	0x504D55
sbi_pmu_counter_stop	0.3	4	0x504D55
sbi_pmu_counter_fw_read	0.3	5	0x504D55
sbi_pmu_counter_fw_read_hi	2.0	6	0x504D55

Chapter 12. Debug Console Extension (EID #0x4442434E "DBCN")

The debug console extension defines a generic mechanism for debugging and boot-time early prints from supervisor-mode software.

This extension replaces the legacy console putchar (EID #0x01) and console getchar (EID #0x02) extensions. The debug console extension allows supervisor-mode software to write or read multiple bytes in a single SBI call.

If the underlying physical console has extra bits for error checking (or correction) then these extra bits should be handled by the SBI implementation.



It is recommended that bytes sent/received using the debug console extension follow UTF-8 character encoding.

12.1. Function: Console Write (FID #0)

```
struct sbiret sbi_debug_console_write(unsigned long num_bytes,
                                     unsigned long base_addr_lo,
                                     unsigned long base_addr_hi)
```

Write bytes to the debug console from input memory.

The **num_bytes** parameter specifies the number of bytes in the input memory. The physical base address of the input memory is represented by two XLEN bits wide parameters. The **base_addr_lo** parameter specifies the lower XLEN bits and the **base_addr_hi** parameter specifies the upper XLEN bits of the input memory physical base address.

This is a non-blocking SBI call and it may do partial/no writes if the debug console is not able to accept more bytes.

The number of bytes written is returned in **sbiret.value** and the possible error codes returned in **sbiret.error** are shown in [Table 45](#) below.

Table 45. Debug Console Write Errors

Error code	Description
SBI_SUCCESS	Bytes written successfully.
SBI_ERR_INVALID_PARAM	The memory pointed to by the num_bytes , base_addr_lo , and base_addr_hi parameters does not satisfy the requirements described in the Section 3.2
SBI_ERR_FAILED	Failed to write due to I/O errors.

12.2. Function: Console Read (FID #1)

```
struct sbiret sbi_debug_console_read(unsigned long num_bytes,
                                     unsigned long base_addr_lo,
                                     unsigned long base_addr_hi)
```

Read bytes from the debug console into an output memory.

The **num_bytes** parameter specifies the maximum number of bytes which can be written into the output memory. The physical base address of the output memory is represented by two XLEN bits wide parameters. The **base_addr_lo** parameter specifies the lower XLEN bits and the **base_addr_hi** parameter specifies the upper XLEN bits of the output memory physical base address.

This is a non-blocking SBI call and it will not write anything into the output memory if there are no bytes to be read in the debug console.

The number of bytes read is returned in **sbiret.value** and the possible error codes returned in **sbiret.error** are shown in [Table 46](#) below.

Table 46. Debug Console Read Errors

Error code	Description
SBI_SUCCESS	Bytes read successfully.
SBI_ERR_INVALID_PARAM	The memory pointed to by the num_bytes , base_addr_lo , and base_addr_hi parameters does not satisfy the requirements described in the Section 3.2
SBI_ERR_FAILED	Failed to read due to I/O errors.

12.3. Function: Console Write Byte (FID #2)

```
struct sbiret sbi_debug_console_write_byte(uint8_t byte)
```

Write a single byte to the debug console.

This is a blocking SBI call and it will only return after writing the specified byte to the debug console. It will also return, with SBI_ERR_FAILED, if there are I/O errors.

The **sbiret.value** is set to zero and the possible error codes returned in **sbiret.error** are shown in [Table 47](#) below.

Table 47. Debug Console Write Byte Errors

Error code	Description
SBI_SUCCESS	Byte written successfully.
SBI_ERR_FAILED	Failed to write the byte due to I/O errors.

12.4. Function Listing

Table 48. DBCN Function List

Function Name	SBI Version	FID	EID
sbi_debug_console_write	2.0	0	0x4442434E
sbi_debug_console_read	2.0	1	0x4442434E
sbi_debug_console_write_byte	2.0	2	0x4442434E

Chapter 13. System Suspend Extension (EID #0x53555350 "SUSP")

The system suspend extension defines a set of system-level sleep states and a function which allows the supervisor-mode software to request that the system transitions to a sleep state. Sleep states are identified with 32-bit wide identifiers (`sleep_type`). The possible values for the identifiers are shown in [Table 49](#).

The term "system" refers to the world-view of supervisor software. The underlying SBI implementation may be provided by machine mode firmware or a hypervisor.

The system suspend extension does not provide any way for supported sleep types to be probed. Platforms are expected to specify their supported system sleep types and per-type wake up devices in their hardware descriptions. The `SUSPEND_TO_RAM` sleep type is the one exception, and its presence is implied by that of the extension.

Table 49. SUSP System Sleep Types

Type	Name	Description
0	SUSPEND_TO_RAM	This is a "suspend to RAM" sleep type, similar to ACPI's S2 or S3. Entry requires all but the calling hart be in the HSM <code>STOPPED</code> state and all hart registers and CSRs saved to RAM.
0x00000001 - 0x7fffffff		Reserved for future use
0x80000000 - 0xffffffff		Platform-specific system sleep types
> 0xffffffff		Reserved

13.1. Function: System Suspend (FID #0)

```
struct sbiret sbi_system_suspend(uint32_t sleep_type,
                                unsigned long resume_addr,
                                unsigned long opaque)
```

A return from a `sbi_system_suspend()` call implies an error and an error code from [Table 51](#) will be in `sbiret.error`. A successful suspend and wake up, results in the hart which initiated the suspend, resuming from the `STOPPED` state. To resume, the hart will jump to supervisor-mode, at the address specified by `resume_addr`, with the specific register values described in [Table 50](#).

Table 50. SUSP System Resume Register State

Register Name	Register Value
satp	0

Register Name	Register Value
sstatus.SIE	0
a0	hartid
a1	opaque parameter
All other registers remain in an undefined state.	



A single **unsigned long** parameter is sufficient for **resume_addr**, because the hart will resume execution in supervisor-mode with the MMU off, hence **resume_addr** must be less than **XLEN** bits wide.

The **resume_addr** parameter points to a runtime-specified physical address, where the hart can resume execution in supervisor-mode after a system suspend.

The **opaque** parameter is an **XLEN**-bit value which will be set in the **a1** register when the hart resumes execution at **resume_addr** after a system suspend.

Besides ensuring all entry criteria for the selected sleep type are met, such as ensuring other harts are in the **STOPPED** state, the caller must ensure all power units and domains are in a state compatible with the selected sleep type. The preparation of the power units, power domains, and wake-up devices used for resumption from the system sleep state is platform specific and beyond the scope of this specification.

When supervisor software is running inside a virtual machine, the SBI implementation is provided by a hypervisor. The system suspend will behave functionally the same as the native case, but might not result in any physical power changes.

The possible error codes returned in **sbiret.error** are shown in [Table 51](#).

Table 51. SUSP System Suspend Errors

Error code	Description
SBI_SUCCESS	System has suspended and resumed successfully.
SBI_ERR_INVALID_PARAM	sleep_type is reserved or is platform-specific and unimplemented.
SBI_ERR_NOT_SUPPORTED	sleep_type is not reserved and is implemented, but the platform does not support it due to one or more missing dependencies.
SBI_ERR_INVALID_ADDRESS	resume_addr is not valid, possibly due to the following reasons: * It is not a valid physical address. * Executable access to the address is prohibited by a physical memory protection mechanism or H-extension G-stage for supervisor mode.
SBI_ERR_FAILED	The suspend request failed for unspecified or unknown other reasons.

13.2. Function Listing

Table 52. SUSP Function List

Function Name	SBI Version	FID	EID
sbi_system_suspend	2.0	0	0x53555350

Chapter 14. CPPC Extension (EID #0x43505043 "CPPC")

ACPI defines the Collaborative Processor Performance Control (CPPC) mechanism, which is an abstract and flexible mechanism for the supervisor-mode power-management software to collaborate with an entity in the platform to manage the performance of the processors.

The SBI CPPC extension provides an abstraction to access the CPPC registers through SBI calls. The CPPC registers can be memory locations shared with a separate platform entity such as a BMC. Even though CPPC is defined in the ACPI specification, it may be possible to implement a CPPC driver based on Device Tree.

[Table 53](#) defines 32-bit identifiers for all CPPC registers to be used by the SBI CPPC functions. The first half of the 32-bit register space corresponds to the registers as defined by the ACPI specification. The second half provides the information not defined in the ACPI specification, but is additionally required by the supervisor-mode power-management software.

Table 53. CPPC Registers

Register ID	Register	Bit Width	Attribute	Description
0x00000000	HighestPerformance	32	Read-only	ACPI Spec 6.5: 8.4.6.1.1.1
0x00000001	NominalPerformance	32	Read-only	ACPI Spec 6.5: 8.4.6.1.1.2
0x00000002	LowestNonlinearPerformance	32	Read-only	ACPI Spec 6.5: 8.4.6.1.1.4
0x00000003	LowestPerformance	32	Read-only	ACPI Spec 6.5: 8.4.6.1.1.5
0x00000004	GuaranteedPerformanceRegister	32	Read-only	ACPI Spec 6.5: 8.4.6.1.1.6
0x00000005	DesiredPerformanceRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.2.3
0x00000006	MinimumPerformanceRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.2.2
0x00000007	MaximumPerformanceRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.2.1
0x00000008	PerformanceReductionToleranceRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.2.4
0x00000009	TimeWindowRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.2.5
0x0000000A	CounterWraparoundTime	32 / 64	Read-only	ACPI Spec 6.5: 8.4.6.1.3.1
0x0000000B	ReferencePerformanceCounterRegister	32 / 64	Read-only	ACPI Spec 6.5: 8.4.6.1.3.1

Register ID	Register	Bit Width	Attribute	Description
0x0000000C	DeliveredPerformanceCounterRegister	32 / 64	Read-only	ACPI Spec 6.5: 8.4.6.1.3.1
0x0000000D	PerformanceLimitedRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.3.2
0x0000000E	CPPCEnableRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.4
0x0000000F	AutonomousSelectionEnable	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.5
0x00000010	AutonomousActivityWindowRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.6
0x00000011	EnergyPerformancePreferenceRegister	32	Read / Write	ACPI Spec 6.5: 8.4.6.1.7
0x00000012	ReferencePerformance	32	Read-only	ACPI Spec 6.5: 8.4.6.1.1.3
0x00000013	LowestFrequency	32	Read-only	ACPI Spec 6.5: 8.4.6.1.1.7
0x00000014	NominalFrequency	32	Read-only	ACPI Spec 6.5: 8.4.6.1.1.7
0x00000015 - 0x7FFFFFFF				Reserved for future use.
0x80000000	TransitionLatency	32	Read-only	Provides the maximum (worst-case) performance state transition latency in nanoseconds.
0x80000001 - 0xFFFFFFFF				Reserved for future use.

14.1. Function: Probe CPPC register (FID #0)

```
struct sbiret sbi_cppc_probe(uint32_t cppc_reg_id)
```

Probe whether the CPPC register as specified by the `cppc_reg_id` parameter is implemented or not by the platform.

If the register is implemented, `sbiret.value` will contain the register width. If the register is not implemented, `sbiret.value` will be set to 0.

The possible error codes returned in `sbiret.error` are shown in [Table 54](#).

Table 54. CPPC Probe Errors

Error code	Description
SBI_SUCCESS	Probe completed successfully.
SBI_ERR_INVALID_PARAM	<code>cppc_reg_id</code> is reserved.
SBI_ERR_FAILED	The probe request failed for unspecified or unknown other reasons.

14.2. Function: Read CPPC register (FID #1)

```
struct sbiret sbi_cppc_read(uint32_t cppc_reg_id)
```

Reads the register as specified in the `cppc_reg_id` parameter and returns the value in `sbiret.value`. When supervisor mode XLEN is 32, the `sbiret.value` will only contain the lower 32 bits of the CPPC register value.

The possible error codes returned in `sbiret.error` are shown in [Table 55](#).

Table 55. CPPC Read Errors

Error code	Description
SBI_SUCCESS	Read completed successfully.
SBI_ERR_INVALID_PARAM	<code>cppc_reg_id</code> is reserved.
SBI_ERR_NOT_SUPPORTED	<code>cppc_reg_id</code> is not implemented by the platform.
SBI_ERR_DENIED	<code>cppc_reg_id</code> is a write-only register.
SBI_ERR_FAILED	The read request failed for unspecified or unknown other reasons.

14.3. Function: Read CPPC register high bits (FID #2)

```
struct sbiret sbi_cppc_read_hi(uint32_t cppc_reg_id)
```

Reads the upper 32-bit value of the register specified in the `cppc_reg_id` parameter and returns the value in `sbiret.value`. This function always returns zero in `sbiret.value` when supervisor mode XLEN is 64 or higher.

The possible error codes returned in `sbiret.error` are shown in [Table 56](#).

Table 56. CPPC Read Hi Errors

Error code	Description
SBI_SUCCESS	Read completed successfully.
SBI_ERR_INVALID_PARAM	<code>cppc_reg_id</code> is reserved.
SBI_ERR_NOT_SUPPORTED	<code>cppc_reg_id</code> is not implemented by the platform.
SBI_ERR_DENIED	<code>cppc_reg_id</code> is a write-only register.

Error code	Description
SBI_ERR_FAILED	The read request failed for unspecified or unknown other reasons.

14.4. Function: Write to CPPC register (FID #3)

```
struct sbiret sbi_cppc_write(uint32_t cppc_reg_id, uint64_t val)
```

Writes the value passed in the **val** parameter to the register as specified in the **cppc_reg_id** parameter.

The possible error codes returned in **sbiret.error** are shown in [Table 57](#).

Table 57. CPPC Write Errors

Error code	Description
SBI_SUCCESS	Write completed successfully.
SBI_ERR_INVALID_PARAM	cppc_reg_id is reserved.
SBI_ERR_NOT_SUPPORTED	cppc_reg_id is not implemented by the platform.
SBI_ERR_DENIED	cppc_reg_id is a read-only register.
SBI_ERR_FAILED	The write request failed for unspecified or unknown other reasons.

14.5. Function Listing

Table 58. CPPC Function List

Function Name	SBI Version	FID	EID
sbi_cppc_probe	2.0	0	0x43505043
sbi_cppc_read	2.0	1	0x43505043
sbi_cppc_read_hi	2.0	2	0x43505043
sbi_cppc_write	2.0	3	0x43505043

Chapter 15. Nested Acceleration Extension (EID #0x4E41434C "NACL")

Nested virtualization is the ability of a hypervisor to run another hypervisor as a guest. RISC-V nested virtualization requires an LO hypervisor (running in hypervisor-mode) to trap-and-emulate the RISC-V H-extension [priv_v1.12] functionality (such as CSR accesses, HFENCE instructions, HLV/HSV instructions, etc.) for the L1 hypervisor (running in virtualized supervisor-mode).

The SBI nested acceleration extension defines a shared memory based interface between the SBI implementation (or LO hypervisor) and the supervisor software (or L1 hypervisor) which allows both to collaboratively reduce traps taken by the LO hypervisor for emulating RISC-V H-extension functionality. The nested acceleration shared memory allows the L1 hypervisor to batch multiple RISC-V H-extension CSR accesses and HFENCE requests which are then emulated by the LO hypervisor upon an explicit synchronization SBI call.



The M-mode firmware should not implement the SBI nested acceleration extension if the underlying platform has the RISC-V H-extension implemented in hardware.

This SBI extension defines optional features which MUST be discovered by the supervisor software (or L1 hypervisor) before using the corresponding SBI functions. Each nested acceleration feature is assigned a unique ID which is an unsigned 32-bit integer. The Table 59 below provides a list of all nested acceleration features.

Table 59. Nested acceleration features

Feature ID	Feature Name	Description
0x00000000	SBI_NACL_FEAT_SYNC_CSR	Synchronize CSR
0x00000001	SBI_NACL_FEAT_SYNC_HFENCE	Synchronize HFENCE
0x00000002	SBI_NACL_FEAT_SYNC_SRET	Synchronize SRET
0x00000003	SBI_NACL_FEAT_AUTOSWAP_CSR	Autoswap CSR
> 0x00000003	RESERVED	Reserved for future use

To use the SBI nested acceleration extension, the supervisor software (or L1 hypervisor) MUST set up a nested acceleration shared memory physical address for each virtual hart at boot-time. The physical base address of the nested acceleration shared memory MUST be 4096 bytes (i.e. page) aligned and the size of the nested acceleration shared memory is assumed to be $4096 + (1024 * (XLEN / 8))$ bytes. The Table 60 below shows the layout of nested acceleration shared memory.

Table 60. Nested acceleration shared memory layout

Name	Offset	Size (bytes)	Description
Scratch space	0x00000000	4096	Nested acceleration feature specific data.

Name	Offset	Size (bytes)	Description
CSR space	0x00001000	$XLEN * 128$	An array of 1024 XLEN-bit words where each word corresponds to a possible RISC-V H-extension CSR defined in the Table 2.1 of the RISC-V privileged specification [priv_v1.12] .

The contents of the scratch space shown in the Table 60 above is defined separately for each nested acceleration feature.

The contents of the CSR space shown in the Table 60 above is an array of RISC-V H-extension CSR values where CSR $\langle x \rangle$ is at index $\langle i \rangle = ((\langle x \rangle \& 0xc00) \gg 2) | (\langle x \rangle \& 0xff)$. The SBI implementation (or LO hypervisor) MUST update the CSR space whenever the state of any RISC-V H-extension CSR changes unless some nested acceleration feature defines a different behaviour. The Table 61 below shows CSR space index ranges for all possible 1024 RISC-V H-extension CSRs.

Table 61. Nested acceleration H-extension CSR index ranges

H-extension CSR address				SBI NACL CSR space index
[11:10]	[9:8]	[7:4]	Hex Range	Hex Range
00	10	xxxx	0x200 - 0x2ff	0x000 - 0x0ff
01	10	0xxx	0x600 - 0x67f	0x100 - 0x17f
01	10	10xx	0x680 - 0x6bf	0x180 - 0x1bf
01	10	11xx	0x6c0 - 0x6ff	0x1c0 - 0x1ff
10	10	0xxx	0xa00 - 0xa7f	0x200 - 0x27f
10	10	10xx	0xa80 - 0xabf	0x280 - 0x2bf
10	10	11xx	0xac0 - 0xaff	0x2c0 - 0x2ff
11	10	0xxx	0xe00 - 0xe7f	0x300 - 0x37f
11	10	10xx	0xe80 - 0xebf	0x380 - 0x3bf
11	10	11xx	0xec0 - 0xeff	0x3c0 - 0x3ff

15.1. Feature: Synchronize CSR (ID #0)

The synchronize CSR feature describes the ability of the SBI implementation (or LO hypervisor) to allow supervisor software (or L1 hypervisor) to write RISC-V H-extension CSRs using the CSR space.

This nested acceleration feature defines the scratch space offset range **0x0F80 - 0x0FFF** (128 bytes) as nested CSR dirty bitmap. The nested CSR dirty bitmap contains 1-bit for each possible RISC-V H-extension CSR.

To write a CSR $\langle x \rangle$ in nested acceleration shared memory, the supervisor software (or L1 hypervisor) MUST do the following:

1. Compute $\langle i \rangle = ((\langle x \rangle \& 0xc00) \gg 2) | (\langle x \rangle \& 0xff)$

2. Write a new CSR value at word with index *<i>* in the CSR space
3. Set the *<i>* bit in the nested CSR dirty bitmap

To synchronize a CSR *<x>*, the SBI implementation (or LO hypervisor) MUST do the following:

1. Compute $\text{<i> = ((<x> \& 0xc00) \gg 2) | (<x> \& 0xff)}$
2. If bit *<i>* is not set in the nested CSR dirty bitmap then goto step 5
3. Emulate write to CSR *<x>* with the new CSR value taken from the word with index *<i>* in the CSR space
4. Clear the *<i>* bit in the nested CSR dirty bitmap
5. Write back the latest CSR value of CSR *<x>* to the word with index *<i>* in the CSR space

When synchronizing multiple CSRs, if the value of a CSR *<y>* depends on the value of some other CSR *<x>* then the SBI implementation (or LO hypervisor) MUST synchronize CSR *<x>* before CSR *<y>*. For example, the value of CSR *hip* depends on the value of the CSR *hvip*, which means *hvip* is emulated and written first, followed by *hip*.

15.2. Feature: Synchronize HFENCE (ID #1)

The synchronize HFENCE feature describes the ability of the SBI implementation (or LO hypervisor) to allow supervisor software (or L1 hypervisor) to issue HFENCE using the scratch space.

This nested acceleration feature defines the scratch space offset range *0x0800 - 0x0F7F* (1920 bytes) as an array of nested HFENCE entries. The total number of nested HFENCE entries are *3840 / XLEN* where each nested HFENCE entry consists of four XLEN-bit words.

A nested HFENCE entry is equivalent to an HFENCE over a range of guest addresses. The [Table 62](#) below shows the nested HFENCE entry format whereas [Table 63](#) below provides a list of nested HFENCE entry types. Upon an explicit synchronize HFENCE request from supervisor software (or L1 hypervisor), the SBI implementation (or LO hypervisor) will process nested HFENCE entries with the *Config.Pending* bit set. After processing pending nested HFENCE entries, the SBI implementation (or LO hypervisor) will clear the *Config.Pending* bit of these entries.

Table 62. Nested HFENCE entry format

Word	Name	Encoding
0	Config	<p>Config information about the nested HFENCE entry</p> <p>BIT[XLEN-1:XLEN-1] - Pending BIT[XLEN-2:XLEN-4] - Reserved and must be zero BIT[XLEN-5:XLEN-8] - Type BIT[XLEN-9:XLEN-9] - Reserved and must be zero BIT[XLEN-10:XLEN-16] - Order if XLEN == 32 then BIT[15:9] - VMID BIT[8:0] - ASID else BIT[29:16] - VMID BIT[15:0] - ASID</p> <p>The page size for invalidation is assumed to be $1 \ll (\text{Config.Order} + 12)$ bytes.</p>
1	Page_Number	Page address right shifted by $\text{Config.Order} + 12$
2	Reserved	Reserved for future use and must be zero
3	Page_Count	Number of pages to invalidate

Table 63. Nested HFENCE entry types

Type	Name	Description
0	GVMA	Invalidate a guest physical address range across all VMIDs. The VMID and ASID fields of the Config word are ignored and MUST be zero.
1	GVMA_ALL	Invalidate all guest physical addresses across all VMIDs. The Order , VMID and ASID fields of the Config word are ignored and MUST be zero. The Page_Number and Page_Count words are ignored and MUST be zero.
2	GVMA_VMID	Invalidate a guest physical address range for a particular VMID. The ASID field of the Config word is ignored and MUST be zero.
3	GVMA_VMID_ALL	Invalidate all guest physical addresses for a particular VMID. The Order and ASID fields of the Config word are ignored and MUST be zero. The Page_Number and Page_Count words are ignored and MUST be zero.
4	VVMA	Invalidate a guest virtual address range for a particular VMID. The ASID field of the Config word is ignored and MUST be zero.
5	VVMA_ALL	Invalidate all guest virtual addresses for a particular VMID. The Order and ASID fields of the Config word are ignored and MUST be zero. The Page_Number and Page_Count words are ignored and MUST be zero.

Type	Name	Description
6	VVMA_ASID	Invalidate a guest virtual address range for a particular VMID and ASID.
7	VVMA_ASID_ALL	Invalidate all guest virtual addresses for a particular VMID and ASID. The Order field of the Config word is ignored and MUST be zero. The Page_Number and Page_Count words are ignored and MUST be zero.
> 7	Reserved	Reserved for future use.

To add a nested HFENCE entry, the supervisor software (or L1 hypervisor) MUST do the following:

1. Find an unused nested HFENCE entry with **Config.Pending == 0**
2. Update the **Page_Number** and **Page_Count** words in the nested HFENCE entry
3. Update the **Config** word in the nested HFENCE entry such that **Config.Pending** bit is set

To synchronize a nested HFENCE entry, the SBI implementation (or LO hypervisor) MUST do the following:

1. If **Config.Pending == 0** then do nothing and skip below steps
2. Process HFENCE based on details in the nested HFENCE entry
3. Clear the **Config.Pending** bit in the nested HFENCE entry

15.3. Feature: Synchronize SRET (ID #2)

The synchronize SRET feature describes the ability of the SBI implementation (or LO hypervisor) to do synchronization of CSRs and HFENCES in the nested acceleration shared memory for the supervisor software (or L1 hypervisor) along with SRET emulation.

This nested acceleration feature defines the scratch space offset range **0x0000 - 0x01FF** (512 bytes) as nested SRET context. The [Table 64](#) below shows contents of the nested SRET context.

Table 64. Nested SRET context

Offset	Name	Encoding
0 * (XLEN / 8)	Reserved	Reserved for future use and must be zero
1 * (XLEN / 8)	X1	Value to be restored in GPR X1
2 * (XLEN / 8)	X2	Value to be restored in GPR X2
3 * (XLEN / 8)	X3	Value to be restored in GPR X3
4 * (XLEN / 8)	X4	Value to be restored in GPR X4
5 * (XLEN / 8)	X5	Value to be restored in GPR X5
6 * (XLEN / 8)	X6	Value to be restored in GPR X6
7 * (XLEN / 8)	X7	Value to be restored in GPR X7
8 * (XLEN / 8)	X8	Value to be restored in GPR X8
9 * (XLEN / 8)	X9	Value to be restored in GPR X9

Offset	Name	Encoding
10 * (XLEN / 8)	X10	Value to be restored in GPR X10
11 * (XLEN / 8)	X11	Value to be restored in GPR X11
12 * (XLEN / 8)	X12	Value to be restored in GPR X12
13 * (XLEN / 8)	X13	Value to be restored in GPR X13
14 * (XLEN / 8)	X14	Value to be restored in GPR X14
15 * (XLEN / 8)	X15	Value to be restored in GPR X15
16 * (XLEN / 8)	X16	Value to be restored in GPR X16
17 * (XLEN / 8)	X17	Value to be restored in GPR X17
18 * (XLEN / 8)	X18	Value to be restored in GPR X18
19 * (XLEN / 8)	X19	Value to be restored in GPR X19
20 * (XLEN / 8)	X20	Value to be restored in GPR X20
21 * (XLEN / 8)	X21	Value to be restored in GPR X21
22 * (XLEN / 8)	X22	Value to be restored in GPR X22
23 * (XLEN / 8)	X23	Value to be restored in GPR X23
24 * (XLEN / 8)	X24	Value to be restored in GPR X24
25 * (XLEN / 8)	X25	Value to be restored in GPR X25
26 * (XLEN / 8)	X26	Value to be restored in GPR X26
27 * (XLEN / 8)	X27	Value to be restored in GPR X27
28 * (XLEN / 8)	X28	Value to be restored in GPR X28
29 * (XLEN / 8)	X29	Value to be restored in GPR X29
30 * (XLEN / 8)	X30	Value to be restored in GPR X30
31 * (XLEN / 8)	X31	Value to be restored in GPR X31
32 * (XLEN / 8) - 0x1FF	Reserved	Reserved for future use

Before sending a synchronize SRET request to the SBI implementation (or LO hypervisor), the supervisor software (or L1 hypervisor) MUST write the GPR $X\langle i \rangle$ values to be restored at offset $\langle i \rangle * (XLEN / 8)$ of the nested SRET context.

Upon a synchronize SRET request from the supervisor software (or L1 hypervisor), the SBI implementation (or LO hypervisor) MUST do the following:

1. If SBI_NACL_FEAT_SYNC_CSR feature is available then
 - a. All RISC-V H-extension CSRs implemented by the SBI implementation (or LO hypervisor) are synchronized as described in the [Section 15.1](#). This is equivalent to the SBI call `sbi_nacl_sync_csr(-1UL)`.
2. If SBI_NACL_FEAT_SYNC_HFENCE feature is available then
 - a. All nested HFENCE entries are synchronized as described in the [Section 15.2](#). This is equivalent to the SBI call `sbi_nacl_sync_hfence(-1UL)`.

3. Restore GPR **X<i>i</i>** registers from the nested SRET context.
4. Emulate the SRET instruction as defined by the RISC-V Privilege specification [\[priv_v1.12\]](#).

15.4. Feature: Autoswap CSR (ID #3)

The autoswap CSR feature describes the ability of the SBI implementation (or LO hypervisor) to automatically swap certain RISC-V H-extension CSR values from the nested acceleration shared memory in the following situations:

- Before emulating the SRET instruction for a synchronized SRET request from the supervisor software (or L1 hypervisor).
- After supervisor (or L1) virtualization state changes from ON to OFF.



The supervisor software (or L1 hypervisor) should use the autoswap CSR feature in conjunction with the synchronize SRET feature.

This nested acceleration feature defines the scratch space offset range **0x0200 - 0x027F** (128 bytes) as nested autoswap context. The [Table 65](#) below shows contents of the nested autoswap context.

Table 65. Nested autoswap context

Offset	Name	Encoding
0 * (XLEN / 8)	Autoswap_Flags	Autoswap flags BIT[XLEN-1:1] - Reserved for future use and must be zero BIT[0:0] - HSTATUS
1 * (XLEN / 8)	HSTATUS	Value to be swapped with HSTATUS CSR
2 * (XLEN / 8) - 0x7F	Reserved	Reserved for future use.

To enable automatic swapping of CSRs from the nested autoswap context, the supervisor software (or L1 hypervisor) MUST do the following:

1. Write the **HSTATUS** swap value in the nested autoswap context.
2. Set **Autoswap_Flags.HSTATUS** bit in the nested autoswap context.

To swap CSRs from the nested autoswap context, the SBI implementation (or LO hypervisor) MUST do the following:

1. If **Autoswap_Flags.HSTATUS** bit is set in the nested autoswap context then swap the supervisor **HSTATUS** CSR value with the **HSTATUS** value in the nested autoswap context.

15.5. Function: Probe nested acceleration feature (FID #0)

```
struct sbiret sbi_nacl_probe_feature(uint32_t feature_id)
```

Probe a nested acceleration feature. This is a mandatory function of the SBI nested acceleration extension. The `feature_id` parameter specifies the nested acceleration feature to probe. [Table 59](#) provides a list of possible feature IDs.

This function always returns `SBI_SUCCESS` in `sbiret.error`. It returns `0` in `sbiret.value` if the given `feature_id` is not available, or `1` in `sbiret.value` if it is available.

15.6. Function: Set nested acceleration shared memory (FID #1)

```
struct sbiret sbi_nacl_set_shmem(unsigned long shmem_phys_lo,
                                unsigned long shmem_phys_hi,
                                unsigned long flags)
```

Set and enable the shared memory for nested acceleration on the calling hart. This is a mandatory function of the SBI nested acceleration extension.

If both `shmem_phys_lo` and `shmem_phys_hi` parameters are not all-ones bitwise then `shmem_phys_lo` specifies the lower XLEN bits and `shmem_phys_hi` specifies the upper XLEN bits of the shared memory physical base address. `shmem_phys_lo` MUST be 4096 bytes (i.e. page) aligned and the size of the shared memory is assumed to be $4096 + (\text{XLEN} * 128)$ bytes.

If both `shmem_phys_lo` and `shmem_phys_hi` parameters are all-ones bitwise then the nested acceleration features are disabled.

The `flags` parameter is reserved for future use and must be zero.

The possible error codes returned in `sbiret.error` are shown in [Table 66](#).

Table 66. NACL Set Shared Memory Errors

Error code	Description
<code>SBI_SUCCESS</code>	Shared memory was set or cleared successfully.
<code>SBI_ERR_INVALID_PARAM</code>	The <code>flags</code> parameter is not zero or the <code>shmem_phys_lo</code> parameter is not 4096 bytes aligned.
<code>SBI_ERR_INVALID_ADDRESS</code>	The shared memory pointed to by the <code>shmem_phys_lo</code> and <code>shmem_phys_hi</code> parameters does not satisfy the requirements described in Section 3.2 .

15.7. Function: Synchronize shared memory CSRs (FID #2)

```
struct sbiret sbi_nacl_sync_csr(unsigned long csr_num)
```

Synchronize CSRs in the nested acceleration shared memory. This is an optional function which is

only available if the SBI_NACL_FEAT_SYNC_CSR feature is available. The parameter `csr_num` specifies the set of RISC-V H-extension CSRs to be synchronized.

If `csr_num` is all-ones bitwise then all RISC-V H-extension CSRs implemented by the SBI implementation (or LO hypervisor) are synchronized as described in the [Section 15.1](#).

If $(\text{csr_num} \& 0x300) == 0x200$ and $\text{csr_num} < 0x1000$ then only a single RISC-V H-extension CSR specified by the `csr_num` parameter is synchronized as described in the [Section 15.1](#).

The possible error codes returned in `sbiret.error` are shown in [Table 67](#).

Table 67. NACL Synchronize CSR Errors

Error code	Description
SBI_SUCCESS	CSRs synchronized successfully.
SBI_ERR_NOT_SUPPORTED	SBI_NACL_FEAT_SYNC_CSR feature is not available.
SBI_ERR_INVALID_PARAM	<code>csr_num</code> is not all-ones bitwise and either: * $(\text{csr_num} \& 0x300) \neq 0x200$ or * $\text{csr_num} \geq 0x1000$ or * <code>csr_num</code> is not implemented by the SBI implementation
SBI_ERR_NO_SHMEM	Nested acceleration shared memory not available.

15.8. Function: Synchronize shared memory HFENCES (FID #3)

```
struct sbiret sbi_nacl_sync_hfence(unsigned long entry_index)
```

Synchronize HFENCES in the nested acceleration shared memory. This is an optional function which is only available if the SBI_NACL_FEAT_SYNC_HFENCE feature is available. The parameter `entry_index` specifies the set of nested HFENCE entries to be synchronized.

If `entry_index` is all-ones bitwise then all nested HFENCE entries are synchronized as described in the [Section 15.2](#).

If $\text{entry_index} < (3840 / \text{XLEN})$ then only a single nested HFENCE entry specified by the `entry_index` parameter is synchronized as described in the [Section 15.2](#).

The possible error codes returned in `sbiret.error` are shown in [Table 68](#).

Table 68. NACL Synchronize HFENCE Errors

Error code	Description
SBI_SUCCESS	HFENCES synchronized successfully.
SBI_ERR_NOT_SUPPORTED	SBI_NACL_FEAT_SYNC_HFENCE feature is not available.
SBI_ERR_INVALID_PARAM	<code>entry_index</code> is not all-ones bitwise and $\text{entry_index} \geq (3840 / \text{XLEN})$.
SBI_ERR_NO_SHMEM	Nested acceleration shared memory not available.

15.9. Function: Synchronize shared memory and emulate SRET (FID #4)

```
struct sbiret sbi_nacl_sync_sret(void)
```

Synchronize CSRs and HFENCES in the nested acceleration shared memory and emulate the SRET instruction. This is an optional function which is only available if the SBI_NACL_FEAT_SYNC_SRET feature is available.

This function is used by supervisor software (or L1 hypervisor) to do a synchronize SRET request and the SBI implementation (or LO hypervisor) MUST handle it as described in the [Section 15.3](#).

This function does not return upon success and the possible error codes returned in `sbiret.error` upon failure are shown in [Table 69](#).

Table 69. NACL Synchronize SRET Errors

Error code	Description
SBI_ERR_NOT_SUPPORTED	SBI_NACL_FEAT_SYNC_SRET feature is not available.
SBI_ERR_NO_SHMEM	Nested acceleration shared memory not available.

15.10. Function Listing

Table 70. NACL Function List

Function Name	SBI Version	FID	EID
sbi_nacl_probe_feature	2.0	0	0x4E41434C
sbi_nacl_set_shmem	2.0	1	0x4E41434C
sbi_nacl_sync_csr	2.0	2	0x4E41434C
sbi_nacl_sync_hfence	2.0	3	0x4E41434C
sbi_nacl_sync_sret	2.0	4	0x4E41434C

Chapter 16. Steal-time Accounting Extension (EID #0x535441 "STA")

SBI implementations may encounter situations where virtual HARTs are ready to run, but must be withheld from running. These situations may be, for example, when multiple SBI domains share processors or when an SBI implementation is a hypervisor and guest contexts share processors with other guest contexts or host tasks. When virtual HARTs are at times withheld from running, observers within the contexts of the virtual HARTs may need a way to account for less progress than would otherwise be expected. The time a virtual HART was ready, but had to wait, is called "stolen time" and the tracking of it is referred to as steal-time accounting. The Steal-time Accounting (STA) extension defines the mechanism in which an SBI implementation provides steal-time and preemption information, for each virtual HART, to supervisor-mode software.

16.1. Function: Set Steal-time Shared Memory Address (FID #0)

```
struct sbiret sbi_steal_time_set_shmem(unsigned long shmem_phys_lo,
                                       unsigned long shmem_phys_hi,
                                       uint32_t flags)
```

Set the shared memory physical base address for steal-time accounting of the calling virtual HART and enable the SBI implementation's steal-time information reporting.

If `shmem_phys_lo` and `shmem_phys_hi` are not all-ones bitwise, then `shmem_phys_lo` specifies the lower XLEN bits and `shmem_phys_hi` specifies the upper XLEN bits of the shared memory physical base address. `shmem_phys_lo` MUST be 64-byte aligned. The size of the shared memory is assumed to be at least 64 bytes. All bytes MUST be set to zero by the SBI implementation before returning from the SBI call.

If `shmem_phys_lo` and `shmem_phys_hi` are all-ones bitwise, the SBI implementation will stop reporting steal-time information for the virtual HART.

`flags` MUST be set to zero.

It is not expected for the shared memory to be written by the supervisor-mode software while it is in use for steal-time accounting. However, the SBI implementation MUST not misbehave if a write from supervisor-mode software occurs, however, in that case, it MAY leave the shared memory filled with inconsistent data.

The SBI implementation MUST stop writing to the shared memory upon system reset.

The shared memory layout is defined in [Table 71](#)

Table 71. STA Shared Memory Structure

Name	Offset	Size	Description
sequence	0	4	<p>The SBI implementation MUST increment this field to an odd value before writing the steal field, and increment it again to an even value after writing steal (i.e. an odd sequence number indicates an in-progress update). The SBI implementation SHOULD ensure that the sequence field remains odd for only very short periods of time.</p> <p>The supervisor-mode software MUST check this field before and after reading the steal field, and repeat the read if it is different or odd.</p> <p><i>This sequence field enables the value of the steal field to be read by supervisor-mode software executing in a 32-bit environment.</i></p>
flags	4	4	<p>Always zero.</p> <p><i>Future extensions of the SBI call might allow the supervisor-mode software to write to some of the fields of the shared memory. Such extensions will not be enabled as long as a zero value is used for the flags argument to the SBI call.</i></p>
steal	8	8	<p>The amount of time in which this virtual HART was not idle and scheduled out, in nanoseconds. The time during which the virtual HART is idle will not be reported as steal-time.</p>
preempted	16	1	<p>An advisory flag indicating whether the virtual HART which registered this structure is running or not. A non-zero value MAY be written by the SBI implementation if the virtual HART has been preempted (i.e. while the steal field is increasing), while a zero value MUST be written before the virtual HART starts to run again.</p> <p><i>This preempted field can, for example, be used by the supervisor-mode software to check if a lock holder has been preempted, and, in that case, disable optimistic spinning.</i></p>
pad	17	47	Pad with zeros to a 64 byte boundary.

sbiret.value is set to zero and the possible error codes returned in **sbiret.error** are shown in Table 72 below.

Table 72. STA Set Steal-time Shared Memory Address Errors

Error code	Description
SBI_SUCCESS	The steal-time shared memory physical base address was set or cleared successfully.
SBI_ERR_INVALID_PARAM	The <code>flags</code> parameter is not zero or the <code>shmem_phys_lo</code> is not 64-byte aligned.
SBI_ERR_INVALID_ADDRESS	The shared memory pointed to by the <code>shmem_phys_lo</code> and <code>shmem_phys_hi</code> parameters is not writable or does not satisfy other requirements of Section 3.2 .
SBI_ERR_FAILED	The request failed for unspecified or unknown other reasons.

16.2. Function Listing

Table 73. STA Function List

Function Name	SBI Version	FID	EID
<code>sbi_steal_time_set_shmem</code>	2.0	0	0x535441

Chapter 17. SBI Supervisor Software Events Extension (EID #0x535345 "SSE")

The SBI Supervisor Software Events (SSE) extension provides a mechanism to inject software events from an SBI implementation to supervisor software such that it preempts all other traps and interrupts.

The software events can be of two types: local or global. A local software event is local to a HART and can be handled only on that HART whereas a global software event is a system event and can be handled by any HART.

17.1. Software Event Identification

Each software event is identified by a unique 32-bit unsigned integer called `event_id` which is encoded as shown in [Table 74](#) below.

Table 74. SSE Event Identification

Software Event ID	Description
0x00000000	Local RAS event
0x00000001	Local PMU event
0x00000002	Local async page fault event
0x00000003 - 0x3fffffff	Reserved for future use
0x40000000 - 0x7ffffffe	Local platform specific event
0x7fffffff	Local debug event
0x80000000	Global RAS event
0x80000001 - 0xbfffffff	Reserved for future use
0xc0000000 - 0xffffffff	Global platform specific event
0xffffffff	Global debug event

17.2. Software Event States

At any point in time, a software event can be in one of the following states:

1. **UNUSED** - Software event is not used by supervisor software
2. **REGISTERED** - Supervisor software has provided an event handler for the software event but it is not ready to handle the events.
3. **ENABLED** - Supervisor software is ready to handle the software event.
4. **PENDING** - Software event is pending and not yet delivered to the supervisor software.
5. **RUNNING** - Supervisor software has taken the software event and is busy handling it.

The below diagram shows the state transitions of a software event.

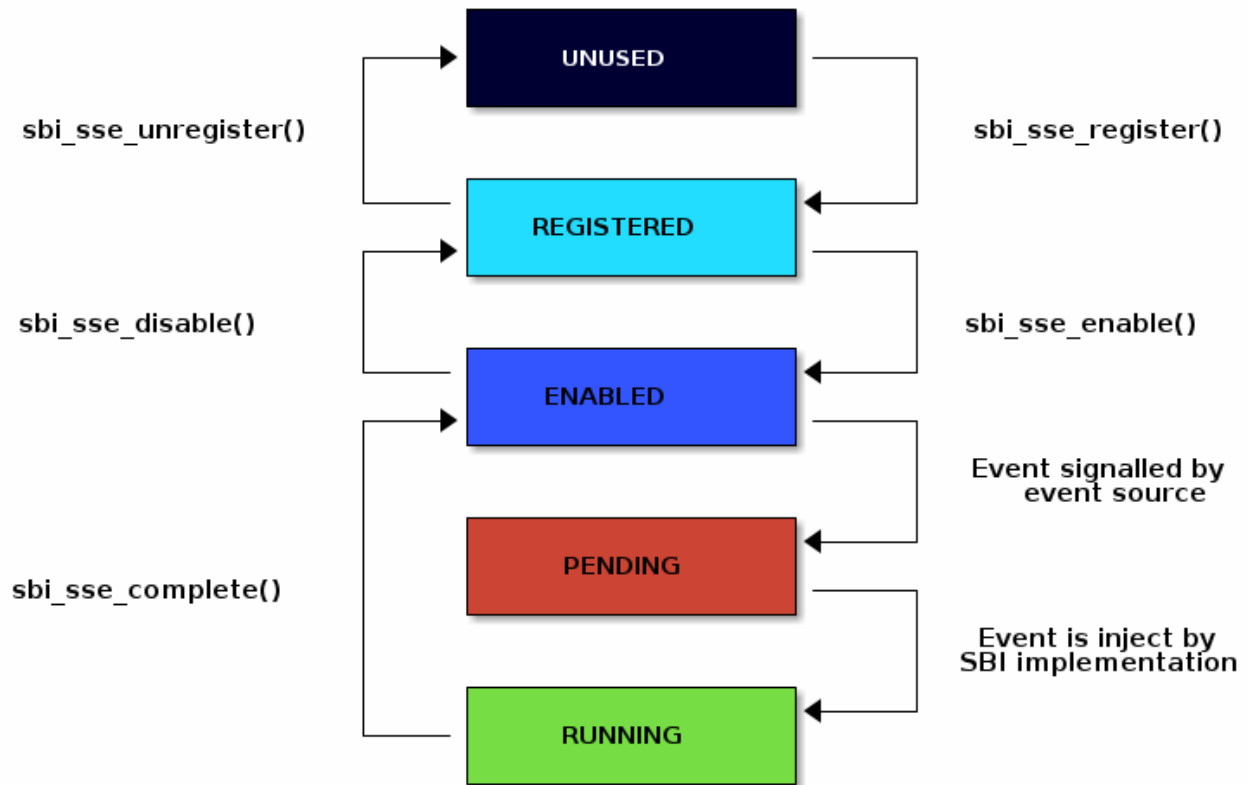


Figure 4. SBI SSE State Machine

A **global** software event **MUST** be registered and enabled only once by any HART. By default, a global software event will be routed to any HART but supervisor software can select a specific HART to handle this event. The state of a global software event **MUST** be common to all HARTs.

A **local** software event **MUST** be registered and enabled by all HARTs who want to handle the event. The state of a local software event **MUST** be tracked separately for each HART.

17.3. Software Event Priority

Each software event has an associated event priority (also referred as `event_priority`) which can be used by an SBI implementation to select a software event for injection when multiple software events are pending on the same HART.

The priority of a software event is a 32-bit unsigned integer where lower value means higher priority. By default, all software events have event priority as zero.

If two or more event have same priority on a given hart then SBI implementation must use `event_id` to break the tie where lower `event_id` has higher priority.

A higher priority event, unless disabled by supervisor software, **always** preempts a lower priority event on the same HART. Once the higher priority event is marked as completed, the previous handler will be resumed.



If a pending or running event is signalled by the event source again, the event will become pending only after the current event completes provided that supervisor software doesn't disable the event on completion.

17.4. Software Event Attributes

A software event can have various attributes associated to it. A software event attribute is a unique 32-bit unsigned integer called `attr_id`. An attribute can have a Read-Only or Read-Write access permissions. The supervisor software can query these event attributes and change the attributes that have Read-Write access permissions. The [Table 75](#) below provides a list event attributes.

Table 75. SSE Event Attributes

Attribute ID (<code>attr_id</code>)	Read-Only	Description	Possible values
0x00000000	Yes	Software event state	0: UNUSED 1: REGISTERED 2: ENABLED 3: PENDING 4: RUNNING
0x00000001	No	Software event priority	32-bit unsigned integer
0x00000002	Yes	Event injection by the supervisor software using <code>sbi_sse_inject</code> call.	0: Not allowed 1: Allowed
0x00000003	No (global) Yes (local)	The HART id of HART that should be preferred to handle the global software event	unsigned long integer
0x00000004	Yes	Raw Pending Status	This is set when the event source signals the event. When the event is injected, it is cleared. 0: Not Pending 1: Pending
> 0x00000004	---	Reserved for future use	---

17.5. Software Event Handler

To handle a software event, the supervisor software MUST register an event handler and enable it. Each event handler registered by the supervisor software consists of a handler context (also referred to as `handler_context`).

The `handler_context` contains the following register states:

1. **Entry State** - It contains the supervisor register state when handling the software event injected by SBI implementation. It is referred as `entry_state`. This register state must be initialized by the supervisor software before registering the handler. The handler's entry point is at the offset 0 of

the `entry_state` as mentioned in Table 76

2. **Interrupted State** - It contains interrupted register state and is referred as `interrupted_state`. The interrupted execution mode is saved at the end of the `interrupted_state`.

The `handler_context` must be contiguous in both virtual and physical address space. The physical address of the `handler_context` is represented by `handler_context_phys`.



It is advisable to use different context for different events. Since a higher priority event can preempt the lower priority event, if same context is used then the interrupted state will be overwritten with register values of the higher priority event. This will make resuming to the previous handler impossible.

Table 76. SSE Register offsets in entry state

Register Offsets in Entry State	Data
<code>entry_state + 0 * (XLEN / 8)</code>	Entry program counter Must be 2-byte aligned virtual address.
<code>entry_state + 1 * (XLEN / 8)</code>	X1
<code>entry_state + 2 * (XLEN / 8)</code>	X2
<code>entry_state + 3 * (XLEN / 8)</code>	X3
<code>entry_state + 4 * (XLEN / 8)</code>	X4
<code>entry_state + 5 * (XLEN / 8)</code>	X5
<code>entry_state + 6 * (XLEN / 8)</code>	X6
<code>entry_state + 7 * (XLEN / 8)</code>	X7
<code>entry_state + 8 * (XLEN / 8)</code>	X8
<code>entry_state + 9 * (XLEN / 8)</code>	X9
<code>entry_state + 10 * (XLEN / 8)</code>	X10
<code>entry_state + 11 * (XLEN / 8)</code>	X11
<code>entry_state + 12 * (XLEN / 8)</code>	X12
<code>entry_state + 13 * (XLEN / 8)</code>	X13
<code>entry_state + 14 * (XLEN / 8)</code>	X14
<code>entry_state + 15 * (XLEN / 8)</code>	X15
<code>entry_state + 16 * (XLEN / 8)</code>	X16
<code>entry_state + 17 * (XLEN / 8)</code>	X17
<code>entry_state + 18 * (XLEN / 8)</code>	X18
<code>entry_state + 19 * (XLEN / 8)</code>	X19
<code>entry_state + 20 * (XLEN / 8)</code>	X20
<code>entry_state + 21 * (XLEN / 8)</code>	X21
<code>entry_state + 22 * (XLEN / 8)</code>	X22

Register Offsets in Entry State	Data
<code>entry_state + 23 * (XLEN / 8)</code>	X23
<code>entry_state + 24 * (XLEN / 8)</code>	X24
<code>entry_state + 25 * (XLEN / 8)</code>	X25
<code>entry_state + 26 * (XLEN / 8)</code>	X26
<code>entry_state + 27 * (XLEN / 8)</code>	X27
<code>entry_state + 28 * (XLEN / 8)</code>	X28
<code>entry_state + 29 * (XLEN / 8)</code>	X29
<code>entry_state + 30 * (XLEN / 8)</code>	X30
<code>entry_state + 31 * (XLEN / 8)</code>	X31

Table 77. SSE Register offsets in interrupted state

Register Offsets in Interrupted State	Data
<code>interrupted_state + 0 * (XLEN / 8)</code>	Interrupted program counter
<code>interrupted_state + 1 * (XLEN / 8)</code>	Saved copy of X1
<code>interrupted_state + 2 * (XLEN / 8)</code>	Saved copy of X2
<code>interrupted_state + 3 * (XLEN / 8)</code>	Saved copy of X3
<code>interrupted_state + 4 * (XLEN / 8)</code>	Saved copy of X4
<code>interrupted_state + 5 * (XLEN / 8)</code>	Saved copy of X5
<code>interrupted_state + 6 * (XLEN / 8)</code>	Saved copy of X6
<code>interrupted_state + 7 * (XLEN / 8)</code>	Saved copy of X7
<code>interrupted_state + 8 * (XLEN / 8)</code>	Saved copy of X8
<code>interrupted_state + 9 * (XLEN / 8)</code>	Saved copy of X9
<code>interrupted_state + 10 * (XLEN / 8)</code>	Saved copy of X10
<code>interrupted_state + 11 * (XLEN / 8)</code>	Saved copy of X11
<code>interrupted_state + 12 * (XLEN / 8)</code>	Saved copy of X12
<code>interrupted_state + 13 * (XLEN / 8)</code>	Saved copy of X13
<code>interrupted_state + 14 * (XLEN / 8)</code>	Saved copy of X14
<code>interrupted_state + 15 * (XLEN / 8)</code>	Saved copy of X15
<code>interrupted_state + 16 * (XLEN / 8)</code>	Saved copy of X16
<code>interrupted_state + 17 * (XLEN / 8)</code>	Saved copy of X17
<code>interrupted_state + 18 * (XLEN / 8)</code>	Saved copy of X18
<code>interrupted_state + 19 * (XLEN / 8)</code>	Saved copy of X19
<code>interrupted_state + 20 * (XLEN / 8)</code>	Saved copy of X20
<code>interrupted_state + 21 * (XLEN / 8)</code>	Saved copy of X21

Register Offsets in Interrupted State	Data
<code>interrupted_state + 22 * (XLEN / 8)</code>	Saved copy of X22
<code>interrupted_state + 23 * (XLEN / 8)</code>	Saved copy of X23
<code>interrupted_state + 24 * (XLEN / 8)</code>	Saved copy of X24
<code>interrupted_state + 25 * (XLEN / 8)</code>	Saved copy of X25
<code>interrupted_state + 26 * (XLEN / 8)</code>	Saved copy of X26
<code>interrupted_state + 27 * (XLEN / 8)</code>	Saved copy of X27
<code>interrupted_state + 28 * (XLEN / 8)</code>	Saved copy of X28
<code>interrupted_state + 29 * (XLEN / 8)</code>	Saved copy of X29
<code>interrupted_state + 30 * (XLEN / 8)</code>	Saved copy of X30
<code>interrupted_state + 31 * (XLEN / 8)</code>	Saved copy of X31
<code>interrupted_state + 32 * (XLEN / 8)</code>	<p>Interrupted Execution mode</p> <p>bit [0] = Privilege mode which was interrupted (1 = S-mode, 0 = U-mode)</p> <p>bit [1] = Virtualization state which was interrupted (1 = ON, 0 = OFF)</p> <p>bit [2] = Saved copy of sstatus.SPIE</p> <p>bit [XLEN-1:3] = Reserved for future use</p>

17.6. Software Event Injection

To inject a software event on a HART, the SBI implementation must do the following:

1. Copy X1 to X31 registers into the `interrupted_state` in `handler_context` from the offsets mentioned in [Table 77](#).
2. Load X1 to X31 registers from `entry_state` in `handler_context` from the offsets mentioned in [Table 76](#).
3. Save the interrupted mode at offset `interrupted_state + 32 * (XLEN / 8)` in `handler_context` as shown in [Table 77](#).
4. Update registers as follows:
 - a. Set `sstatus.SPIE` = `sstatus.SIE`
 - b. Set `sstatus.SIE` = 0
5. Resume execution with:
 - a. Program counter = value at `entry_state + 0 * (XLEN / 8)`
 - b. Privilege mode = S-mode
 - c. Virtualization state = OFF

17.7. Software Event Completion

After handling the software event on a HART, the supervisor software must notify the SBI implementation about completion of event handling using `sbi_sse_complete` call. The SBI implementation must do the following to complete event handling and resume interrupted state:

1. Restore X1 to X31 registers from the `interrupted_state` of `handler_context` from the offsets mentioned in [Table 77](#).
2. Update supervisor CSRs as follows:
 - a. Set `sstatus.SIE` = `sstatus.SPIE`
 - b. Set `sstatus.SPIE` = bit[2] of the value at `interrupted_state` + 32 * (`XLEN` / 8)
3. Resume execution with:
 - Virtualization state = bit[1] of the value at `interrupted_state` + 32 * (`XLEN` / 8)
 - Privilege mode = bit[0] of the value at `interrupted_state` + 32 * (`XLEN` / 8)
 - Program counter = value at `interrupted_state` + 0 * (`XLEN` / 8)

If the supervisor software wishes to resume from a different location, it can update the `interrupted_state` fields accordingly.

17.8. Function: Get a software event attribute (FID #0)

```
struct sbiret sbi_sse_get_attr(uint32_t event_id,
                             uint32_t attr_id)
```

Get an event attribute value of software event. The `event_id` parameter specifies the software event whereas `attr_id` parameter specifies the event attribute.

Upon success the event attribute value is returned in `sbiret.value`. In case of an error, the possible error codes are shown in the [Table 78](#) below:

Table 78. SSE Event Attribute Read Errors

Error code	Description
SBI_SUCCESS	Attribute of given event returned successfully.
SBI_ERR_INVALID_PARAM	<code>event_id</code> or <code>attr_id</code> or both are invalid.

17.9. Function: Set a software event attribute (FID #1)

```
struct sbiret sbi_sse_set_attr(uint32_t event_id,
                             uint32_t attr_id,
                             unsigned long value)
```

Set an event attribute value of software event. The `event_id` parameter specifies the software event whereas `attr_id` parameter specifies the event attribute. The new event attribute value is specified by `value` parameter.

Any error is returned in `sbiret.error`. The possible return values are listed in [Table 79](#) below:

Table 79. SSE Event Attribute Write Errors

Error code	Description
SBI_SUCCESS	Attribute value set successfully.
SBI_ERR_INVALID_PARAM	<code>event_id</code> or <code>attr_id</code> or both are invalid
SBI_ERR_BAD_RANGE	<code>value</code> does not match the possible values defined in Table 75 .

17.10. Function: Register a software event (FID #2)

```
struct sbiret sbi_sse_register(uint32_t event_id,
                              unsigned long handler_context_phys_hi,
                              unsigned long handler_context_phys_lo)
```

Register a handler for the software event. The `event_id` parameter specifies the event ID for which handler is being registered. The parameters `handler_context_phys_hi` and `handler_context_phys_lo` contain the upper and lower XLEN bits, respectively, of the handler's context. The `handler_context_phys_lo` parameter must be $(XLEN / 8)$ byte aligned.

On successful registration, the event state moves from `UNUSED` to `REGISTERED`. In case of an error, possible error codes are listed in [Table 80](#) below:

Table 80. SSE Event Register Errors

Error code	Description
SBI_SUCCESS	Event handler is registered successfully.
SBI_ERR_INVALID_STATE	The event is not in <code>UNUSED</code> state.
SBI_ERR_INVALID_PARAM	<code>event_id</code> is invalid or other parameters not satisfy requirements defined in Section 17.5 .
SBI_ERR_INVALID_ADDRESS	The memory pointed by <code>handler_context_phys_lo</code> , <code>handler_context_phys_hi</code> , parameters does not satisfy the requirements described in Section 3.2 or The <code>handler_context_phys_lo</code> parameter is not $(XLEN / 8)$ byte aligned.

17.11. Function: Unregister a software event (FID #3)

```
struct sbiret sbi_sse_unregister(uint32_t event_id)
```

Unregister the handler for given `event_id`. The event MUST be in `REGISTERED` state before it can be unregistered.

On successful unregistration, the event is moved to `UNREGISTERED` state. In case of an error, possible error codes are listed in [Table 81](#) below:

Table 81. SSE Event Unregister Errors

Error code	Description
SBI_SUCCESS	Event handler is unregistered successfully.
SBI_ERR_INVALID_STATE	Event is not in <code>REGISTERED</code> state.
SBI_ERR_INVALID_PARAM	<code>event_id</code> is invalid.

17.12. Function: Enable a software event (FID #4)

```
struct sbiret sbi_sse_enable(uint32_t event_id)
```

Enable the software event specified by the `event_id` parameter. For local events, the event is enabled only for the calling HART. For global events, the event is enabled for all the harts of supervisor software.

The event MUST be in `REGISTERED` state otherwise this function will fail.

On success, the event is moved to `ENABLED` state and SBI implementation can inject the event when it occurs. In case of an error, possible error codes are listed in [Table 82](#) below:

Table 82. SSE Event Enable Errors

Error code	Description
SBI_SUCCESS	Event is successfully enabled.
SBI_ERR_INVALID_PARAM	<code>event_id</code> is not valid.
SBI_ERR_INVALID_STATE	The event is not in <code>REGISTERED</code> state.

17.13. Function: Disable a software event (FID #5)

```
struct sbiret sbi_sse_disable(uint32_t event_id)
```

Disable the software event specified by the `event_id` parameter. For local events, the event is disabled

only for the calling HART. For global events, the event is disabled for all the harts of supervisor software. The event must be in **ENABLED** state.

On success, the event is moved to **REGISTERED** state. In case of an error, possible error codes are listed in [Table 83](#).

Table 83. SSE Event Disable Errors

Error code	Description
SBI_SUCCESS	Event is successfully disabled.
SBI_ERR_INVALID_PARAM	event_id is not valid.
SBI_ERR_INVALID_STATE	Event is not in ENABLED state.

17.14. Function: Complete software event handling (FID #6)

```
struct sbiret sbi_sse_complete(uint32_t event_id,
                              uint32_t status,
                              uint32_t flags)
```

Complete the supervisor event handling for the event. The event must be in **RUNNING** state.

If supervisor software could not handle the event, it must set the **status** parameter to **SBI_SSE_HANDLER_FAILED**. On success, it must set the **status** parameter to **SBI_SSE_HANDLER_SUCCESS**. Other possible status codes are listed in [Table 84](#). Any other value of **status** field is ignored.

The **flags** parameter represents additional information from supervisor to the SBI implementation and the [Table 85](#) lists the bit-encoding for it.

Table 84. SSE Event Complete Status Values

Value	Enum Name	Description
0x00000000	SBI_SSE_HANDLER_SUCCESS	Supervisor successfully handled the event.
0x00000001	SBI_SSE_HANDLER_FAILED	Supervisor failed to handle the event.
> 0x00000001	-	Reserved

Table 85. SSE Event Complete Flags Values

Flag Name	Bits	Description
SBI_SSE_EVENT_DISABLE	0:0	Disable the event.
RESERVED	1:(XLEN-1)	All non-zero values are

In case of an error, possible error codes are listed in [Table 86](#).

Table 86. SSE Event Complete Errors

Error code	Description
SBI_SUCCESS	Event is successfully marked completed.
SBI_ERR_INVALID_PARAM	<code>event_id</code> is invalid or <code>status</code> has invalid value.
SBI_ERR_INVALID_STATE	The <code>event_id</code> event is not in RUNNING state.

17.15. Function: Signal a software event (FID #7)

```
struct sbiret sbi_sse_inject(uint32_t event_id,
                             unsigned long hart_id)
```

The supervisor software can inject a software event with the help of this function. The `event_id` paramater refers to the event to be injected.

For global events, the `hart_id` parameter is ignored. For local events, the `hart_id` parameter refers to the HART on which the event is to be injected. An event can only be injected if it is allowed by the event attribute as described in [Table 75](#).

In case of an error, possible error codes are listed in [Table 87](#).

Table 87. SSE Event Inject Errors

Error code	Description
SBI_SUCCESS	Event is successfully injected or marked PENDING on given HART
SBI_ERR_INVALID_PARAM	<code>event_id</code> or <code>hart_id</code> is invalid.

Chapter 18. Experimental SBI Extension Space (EIDs #0x08000000 - #0x08FFFFFF)

No management.

Chapter 19. Vendor Specific Extension Space (EIDs #0x09000000 - #0x09FFFFFF)

Low bits from `mvendorid`.

Chapter 20. Firmware Specific Extension Space (EIDs #0x0A000000 - #0x0AFFFFFF)

Low bits is SBI implementation ID. The firmware specific SBI extensions are for SBI implementations. It provides firmware specific SBI functions which are defined in the external firmware specification.

References

- [priv_v1.12] The RISC-V Instruction Set Manual, Volume II: Privileged Architecture, Document Version 20211203, URL: github.com/riscv/riscv-isa-manual/releases/tag/Priv-v1.12